

Correlated Synthetic Controls



Tzvetan Moev
Nuffield College
Word Count: 16470

Supervisors: Prof Francis DiTraglia; Prof Maximilian Kasy
Field: Econometrics

*Submitted in partial fulfilment of the requirements for the degree of the Master
of Philosophy in Economics*

Correlated Synthetic Controls

Tzvetan Moev

September 19, 2021

Abstract

Synthetic Control methods have recently gained considerable attention in applications with only one treated unit. Their popularity is partly based on the key insight that we can *predict* good synthetic counterfactuals for our treated unit. However, this insight of predicting counterfactuals is generalisable to microeconomic settings where we often observe many treated units. We propose the Correlated Synthetic Controls (CSC) estimator for such situations: intuitively, it creates synthetic controls that are correlated across individuals with similar observables. When treatment assignment is correlated with unobservables, we show that the CSC estimator has more desirable theoretical properties than the difference-in-difference estimator. We also utilise CSC in practice to obtain heterogeneous treatment effects in the well-known Mariel Boatlift study, leveraging additional information from the PSID.

Keywords: Synthetic Controls, Correlated Random Coefficients, Mariel Boatlift

Acknowledgements

My MPhil thesis would not have been possible without the support of many people across Oxford. Firstly, I would like to thank my supervisors: Frank DiTraglia for encouraging me to pursue my own ideas and for stressing the importance of writing clearly and Max Kasy for motivating me to seek good empirical applications and for helping me clarify my thoughts. Next, I acknowledge the support of James Duffy, Bent Nielsen and Ian Jewitt who convinced me to write a thesis in Econometrics in the first place despite my non-technical background. I am grateful for the invaluable comments of Anders Kock, Barbara Petrongolo, Frank Windmeijer, Yi Ying Tan, Xiyu Jiao, Stanislav Slavov and Otso Hao. Last but not least, I owe this thesis to Tom O’Grady’s extremely clear and intuitive lectures on the Synthetic Control Method and Multilevel Modelling at UCL. Combining these two methods forms the main motivation for this thesis. However, all errors in the text are my own.

I am also grateful to my family - Elitsa, Ioan, Ivan, Mimi, Raya, Stoyko, Teodora and Tsonka - and I also apologise to them for missing out on so many big events due to writing this thesis (especially the recent premature birth of my nephew Ioan and the surrounding circumstances). Moreover, I could not have survived this lockdown without my Oxford friends: Max, Micol, Mihai, Ed and Yi Ying! Also, I would have never gotten into social science without having met Charles Gill, Nadejda Germanova and Lora Kiosseva, so a huge thanks goes to them. Lastly, I humbly acknowledge that this work was supported by ESRC Project ES/P000649/1 and in particular grant No. 2260944.

Data & Repository

The data used in the thesis comes from the publicly available release of the Panel Study of Income and Dynamics. For details on how the data was extracted, please check the GitHub repository for this paper by clicking [here](#).

Contents

1	Introduction	1
2	Motivating Example	4
3	Estimator Construction	8
3.a	Related work	8
3.b	Set-up	9
3.c	The Synthetic Control Method	12
3.d	Many treated units	14
3.e	Correlated Synthetic Controls	17
3.f	Estimation	20
4	Theoretical properties	22
4.a	Data Generating Process	23
4.b	CSC Bound	26
4.c	Estimation Error of DiD	30
5	Simulations	32
5.a	Bias of DiD	33
5.b	Infeasible DiD	35
5.c	Distributional Assumptions on DGP	36
5.d	Results	40
6	Empirical Application	43
6.a	Mariel Boatlift	44
6.b	PSID	46
6.c	Penalised vs Correlated Synthetic Controls	48
6.d	Heterogeneous Treatment Effects	49
7	Discussion, Limitations and Conclusions	51
	Appendices	59
A	Supplementary results and examples	59
B	Proofs of main results	76

1 Introduction

The Fundamental Problem of Causal Inference states that we cannot directly infer the causal effect of some intervention for a single individual because we do not simultaneously observe what happens to them with and without treatment. In other words, we cannot do a *within-person* comparison of the two potential outcomes. One way around this problem is Randomized Control Trials, in which researchers randomly allocate treatment among a set of individuals, thereby forming a control group and a treatment group which are similar in terms of observables (and hopefully unobservables). This strategy ensures that in expectation the two groups have the same potential outcome with treatment and the same potential outcome without treatment. The difference of the average outcomes in the treatment group and the control group would identify the Average Treatment Effect (ATE), suggesting that we have managed to escape the Fundamental Problem via a *between-person* comparison.

However, what if we can construct a surrogate for the potential outcome without treatment for every treated observation? If we can achieve this, then we can solve the Fundamental Problem of Causal Inference by approximating the ideal *within-person* comparison. In the case of one treated unit, the Synthetic Control (SC) Method ([Abadie et al., 2010](#)) method constructs such a synthetic counterfactual by taking a weighted average of the control units. The idea is as follows: if the SC is a good match for the treated unit's outcome in the *pre-treatment* period, we can extrapolate its time series to the *post-treatment* period to find what would have happened to the treated unit, had it not received treatment. Taking the difference between the SC and the treated unit will give us the causal effect.

This insight of constructing counterfactuals can also be generalised to the case of many treated units. Similarly to [Abadie and L'Hour \(2020\)](#), this paper develops an estimator that provides SCs for all treated individuals. We call the estimator Correlated Synthetic Controls (CSC), because it builds counterfactuals that are similar across treated individuals with similar observable characteristics.

To explain how CSC works, we should note that generalising SC to panels with

many treated units that are common in applied microeconomics is not straightforward because the original SC method was designed for long T settings whereas panel data in such cases often has short T . Subsequently, there are two extreme approaches that we can take to achieve this. One is to construct a separate SC for every treated individual. The other is to aggregate the time-series for all treated individuals and for all control individuals, because often they belong well-defined groups such as cities. For instance, in the Mariel Boatlift study of [Card \(1990\)](#), we can group all treated individuals to Miami and all control individuals to other US cities. Then, we can construct a SC for Miami by combining other cities' time series. The CSC estimator is a compromise between the two extremes. Similarly, a very recent paper by [Ben-Michael et al. \(2021\)](#) that appeared while this thesis was being written builds an estimator that balances a related but different tradeoff between two extremes.¹

As we shall see, the two extremes that we consider are related to the distinction between pooled OLS and fixed coefficients in the panel data literature. However, there is a middle ground: correlated random coefficients models ([Wooldridge, 2003](#); [Hsiao and Peseran, 2008](#)). The CSC approach postulates such a model for the weights allocated to different donors for every treated individual. In particular, the weights are allowed to differ across treated individuals in a deterministic way based on their observables, similarly to coefficients in correlated random coefficients models. Thus, individuals with similar observables will have similar (or correlated) synthetic controls. This is an attractive approach in microeconomic approaches with many treated and control individuals but few time periods. Thus, the Correlated Synthetic Control approach overcomes the challenge of short T by using information from comparable treated observations when constructing the SCs.

Beyond proposing a novel estimator, we make three further contributions. Firstly, we will compare the theoretical properties of CSC to those of the Difference-in-Difference estimator (DiD). When treatment is strongly correlated with unobservable characteristics (e.g. we suspect selection on unobservables), CSC should be used because its estimate of the treatment effect has a smaller estimation error than DiD. This pro-

¹In Section 3.d, we show how the two extremes that we consider are different from their distinction between *pooled* and *separate* SCs.

vides one good reason for empirical researchers to choose CSC in the context of a panel with many treated units, even though DiD is often considered the default option in such cases (Arkhangelsky et al., 2020). Importantly, this theoretical result does not depend on assumptions specific to CSC but holds more generally for many estimators from the SC family when used in a setting with many treated units.

Secondly, via a simulation study we can compare CSC not only to DiD, but also to Penalised Synthetic Control (Abadie and L’Hour, 2020), our estimator’s closest sibling from the family of SC estimators. We provide an infeasible estimator that estimates consistently the parameter of interest (the Average Treatment Effect on the Treated, or ATT) and study the conditions, under which the three feasible estimators approximate its behaviour. In particular, CSC outperforms other feasible estimators such as DiD in terms of calculating the ATT with a lower estimation error, when our data is generated from an interactive fixed effects model and treatment assignment is correlated with unobservables. This confirms our main theoretical result.

Thirdly, we illustrate how researchers can use CSC in practice via studying the effect of immigration on wages and labour supply in the context of the Mariel Boatlift (Card, 1990). Intuitively, what we do is construct a synthetic doppelganger for every treated worker based on workers in other states. Next, we show that CSC performs slightly better than Penalised Synthetic Control with real data in terms of predicting good counterfactuals in the pre-treatment period. In contrast to previous studies, our empirical application uses an alternative data source, namely the Panel Study of Income and Dynamics, because the credibility of typical data sources used for evaluations of the Mariel Boatlift² have been recently questioned (Clemens and Hunt, 2019). Lastly, we show how our estimator can be used to obtain heterogeneous treatment effects of immigration on low-skilled and high-skilled native workers.

So, the Correlated Synthetic Controls provide a useful addition to the toolkit of empirical researchers for conducting causal inference and the rest of this paper attempts to illustrate this point. Section 2 presents a motivating example, in which the estimator seems more appropriate than standard causal inference techniques. The formal setup

²Mostly variants of CPS. See Peri and Yasenov (2018, p.6-12).

of the problem and the construction of the CSC estimator is detailed in Section 3. Next, its theoretical properties are explored in Section 4 whereas Section 5 provides the result from our simulation exercise. Section 6 applies the estimator to the Mariel Boatlift. Some avenues for further research and limitations are discussed in Section 7. Some supplementary examples and further clarifications are contained in Appendix A whereas Appendix B contains the proofs of the main results in the thesis. The coding for the simulation and the empirical application can be found in [this GitHub repository](#).

2 Motivating Example

The key purpose of this section is to motivate the Correlated Synthetic Control (CSC) estimator via an example, in which it seems more appropriate than other causal inference techniques. A key implication of models in economic geography postulates that market access is an important determinant of the spatial distribution of economic activity (Krugman, 1991; Davis and Weinstein, 2002). Redding and Sturm (2008) provide one of the first empirical tests of this causal link. They exploit the German division after the Second World War as an exogenous shock to the market access enjoyed by cities. In particular, they speculate that West German cities close to the East-West border experienced a bigger decline in market access relative to other West German cities. The decline should additionally be more pronounced for small cities close to the border relative to big cities. In such cases, economic geography models imply a decline in economic development.

So, Redding and Sturm (2008) would like to test this mechanism by measuring the possibly heterogeneous treatment effect (depending on city size) of the East-West border on the population of cities close to the border.³ A key decision problem which they are facing is which causal inference technique to use. For the purpose of illustrating the decision problem, suppose that we only observe four cities close to the border, two of which are small and two are big.⁴ Note that we also observe only one covariate -

³In the paper, they argue that population growth is a good proxy for economic development.

⁴In reality, they observe many treated cities but here we abstract from this in order to focus on

whether a city is small or big. See Table 1 for more details on the made-up data for the treated cities.

Table 1: *Made-up data on the population of some German cities close to the border (in 000s) inspired by Redding and Sturm (2008)’s real dataset.*

	Covariates		Population		
	Size	Close to border	1922	1937	1952
Bayreuth	Small	Yes	38	48	62
Erlangen	Small	Yes	32	40	50
Hamburg	Big	Yes	1350	1590	1680
Hanover	Big	Yes	420	480	560

Given this dataset, there are two causal inference techniques that come to mind immediately: difference-in-difference (DiD) and synthetic control (SC). Let us now explain why each of them is not completely adequate in this case. Firstly, we can consider a conventional DiD approach (Angrist and Pischke, 2008), similarly to what Redding and Sturm (2008) did in the original study:

$$Population_{it} = \rho + \gamma_i + \delta_t + \tau D_{it} + e_{it} \quad (1)$$

where i is city, t is time (1922, 1937 or 1952), γ_i are city-level fixed effects, δ_t are time fixed effects, e_{it} are idiosyncratic shocks and D_{it} is a treatment indicator which equals to 1 only for the four cities in Table 1 in 1952. The parameter of interest is τ , which can be interpreted as identifying the Average Treatment Effect on the Treated (ATT). However, we may be concerned that with just 2 pre-treatment periods we cannot evaluate if the parallel trends assumption which is necessary for identifying ATT via DiD holds. Moreover, even if the lack of a pre-treatment period is not such a big concern, it is possible that the parallel trends assumption fails when treatment assignment is not as good as random.⁵ More importantly, as suggested by Redding and Sturm (2008), the treatment effects may be heterogeneous, i.e. they are bigger for smaller cities. However, recent work on DiD with heterogeneous effects has shown the decision problem.

⁵An example of this is provided in Section 5.a where our data is generated from an interactive fixed effects and treatment assignment is not as good as random.

that in such cases τ from (1) will not identify the ATT, even when the parallel trends assumption holds (De Chaisemartin and d’Haultfoeuille, 2020). With few discrete covariates, we can consider doing DiD separately for each group of cities (e.g. small and big cities). However, this approach becomes infeasible if we want to add continuous covariates or if we have many discrete covariates. Thus, overall, using DiD in this particular example does not seem optimal.

Alternatively, we can consider creating a separate SC for every treated city in our sample. The benefit of this approach is that we can get the individual treatment effects for every city with small estimation error under certain conditions (Abadie et al., 2010). This will allow us to capture the heterogeneity among different cities without having to rely on parallel trends.⁶ SC postulates that the population of a treated city is a weighted average of populations of untreated cities:

$$Population_{it} = \sum_{j=1}^{n_0} w_{ij} Population_{jt} + \tau_i D_{it} + e_{it}$$

where w_{ij} is the weight of untreated city j for treated city i , n_0 is the total number of untreated cities and D_{it} and e_{it} are defined as above. The parameter of interest is the individual treatment effect τ_i . Note that w_{ij} are constrained to be non-negative and to sum up to 1, so that they can be understood as probabilities. So, for example, it could be that Hannover’s SC is 20% Munich and 70% Frankfurt. There are two related problems with using SC in this case, especially with a lot of control cities and short T as in our case. There might be more than one combination of donors that perfectly matches the time series of a certain treated city pre-treatment, i.e. a multiplicity of solutions for w_{ij} . For example, consider Hanover in Table 1 and suppose that there are three other cities with the following populations in the pre-treatment period (1922

⁶This points out to a very interesting conceptual differences between how DiD and SC are thought about in the literature. DiD is used as an estimator when there are many treated cities where SC is the preferred method in the case of few treated cities. So, when their theoretical properties are studied, the quantity of interest in SC is usually an individual treatment effect whereas in the case of DiD it is some average quantity such as the ATT. However, under certain assumptions on the structure of our panel data, both approaches can be used in the same set of applications, independently of how many treated units there are.

and 1937):

$$A = \{400, 450\} \quad B = \{440, 510\} \quad C = \{500, 600\}$$

where using control cities A and B with weights 0.5 and 0.5 or using control cities A and C with weights 0.8 and 0.2 both match exactly the time series for Hannover. Moreover, if we try to estimate separate SC for every treated city with just 2 pre-treatment periods, we risk over-fitting a lot, even when multiplicity of solutions is not an issue. So, doing SC is not ideal in this case.

So, if we have decided that neither the standard DiD or SC are appropriate, then we have to look out for another techniques such as the Correlated Synthetic Controls (CSC). Essentially, CSC modifies the weights in such a way that it tackles overfitting and multiplicity of solutions simultaneously. We let the weights depend on observable characteristics such as the size of cities and allow for an individual fixed effect γ_i (as in DiD):

$$Population_{it} = \gamma_i + \sum_{j=1}^{n_0} w_{ij} Population_{jt} + \tau_i D_{it} + e_{it}$$

$$w_{ij} = w_j^{small} Small_i + w_j^{big} Big_i$$

where w_{ij} is the weight of donor j on treated unit i , D_{it} is treatment indicator and τ_i captures the fact that we get an individual estimate of the treatment effect. $Small_i$ and Big_i indicator functions for being a small or big city. While the weights w_{ij} are still constrained to be non-negative and sum up to 1, the most significant difference is that now we constrain them to be the same for our two small cities (Erlangen and Bayreuth) and the same for our two big cities (Hamburg and Hannover). As a result, we do not get multiplicity of solutions because for small cities the weights are balancing both the time-series of Bayreuth and Erlangen. However, the two time trends will not be perfectly matched simultaneously by the same set of weights, irrespectively of the individual fixed effect. Analogically, over-fitting is less of a concern because we have fewer free parameters and each set of weights is exploiting information from different

treated cities belonging to the same group.

So, we conclude that CSCs should be preferred to standard SC and DiD in this case. In contrast to DiD it does not rely on assumptions like parallel trends and can naturally accommodate heterogeneity of treatment effects. On the other hand, CSC does not suffer from multiplicity of solutions and overfitting as opposed to SC. Therefore, there are empirical situations as in [Redding and Sturm \(2008\)](#) where researchers may want to use our estimator.

3 Estimator Construction

3.a Related work

In this section we will briefly review the literature and place the Correlated Synthetic Control (CSC) in a wider context. In a seminal paper, [Abadie et al. \(2010\)](#) introduced the Synthetic Control (SC) method for constructing synthetic counterfactuals for a single treated unit. This was followed by a series of empirical papers, using SC to estimate the treatment effects of various macro interventions such as the effect of Brexit on the British economy ([Born et al., 2017](#)) or the effect of German unification on Germany’s economy ([Abadie et al., 2015](#)). Despite the wave of applied papers using the method, little was known about its theoretical properties until [Doudchenko and Imbens \(2018\)](#) and the work of Bruno Ferman and coauthors ([Botosaru and Ferman, 2019](#); [Ferman, 2020](#); [Ferman and Pinto, 2021](#)). [Doudchenko and Imbens \(2018\)](#) present a general framework for estimators which nests difference-in-difference and SC as special cases and illustrate the connections between the two estimators. Regarding whether difference-in-difference or SC are preferable in applications, [Ferman and Pinto \(2021\)](#) provide conditions under which SC has better theoretical properties. In a related article, [Ferman \(2020\)](#) show that when our data is generated from an interactive fixed effects models, the SC can yield an asymptotically unbiased estimate of the treatment effect under certain conditions. Lastly, [Botosaru and Ferman \(2019\)](#) show that even if our SC does not match perfectly the true SC in the pretreatment period, the approach can still yield a meaningful estimate of the treatment effect.

The papers discussed so far have considered the case of one treated unit. However, the idea of constructing counterfactuals is generalisable to settings with many treated individuals: we can construct SC for every individual that has been treated in our dataset. Several recent papers have proposed estimators that work in this context such as the synthetic difference-in-difference (Arkhangelsky et al., 2020) and matrix completion methods (?). However, the two papers that are most relevant for our contribution, namely the Correlated Synthetic Control, are by Abadie and L’Hour (2020) and Ben-Michael et al. (2021). Firstly, Abadie and L’Hour (2020) propose the Penalised Synthetic Control which tackles a key problem when trying to generalise the original SC to a setting with many treated units, namely the multiplicity of solutions. Secondly, while primarily aimed at building an estimator inspired by SC for staggered treatment adoption, the estimator that Ben-Michael et al. (2021) propose also works for the case of many treated units and has a similar motivation to CSC. The CSC estimator that we propose is explicitly aimed at the many-treated unit setting without staggered adoption and so its closest sibling is Abadie and L’Hour (2020). However, it attempts to improve on existing techniques by bringing insights from the panel data literature in order to exploit the information from treated individuals that are similar in terms of observables.

3.b Set-up

We introduce in this subsection a formal framework for thinking about the family of SC estimators with many treated units based on Doudchenko and Imbens (2018) and more generally the matrix completion literature (?). This formulation of the problem is useful because it allows us to see how the different estimators in the literature relate to each other and to compare the optimisation problem that they solve. Moreover, it brings home the main insight from the different SC methods: we can translate the causal inference problem into a prediction problem (under certain assumptions).

Suppose that we observe many treated units for a relatively short time period. We have available a panel dataset with the outcome variable of N individuals who are

observed for T periods. A policy is implemented *once*⁷ at time $T_0 + 1$ and we are interested in estimating its Average Treatment Effect on the Treated (ATT, denoted by τ_t throughout the paper) and the individual treatment effects (denoted by τ_{it}). The intervention affects a total of n_1 people which forms our treatment group, whereas the other $n_0 \equiv N - n_1$ people form our control group, which we refer to interchangeably as the donors following [Abadie et al. \(2010\)](#). So, we can observe the outcome variable $y_{it}(D)$ for both donors $i \in \{1, 2, \dots, n_0\}$ and treated individuals $i \in \{n_0 + 1, n_0 + 2, \dots, n_0 + n_1\}$ in the pre-treatment period $t \in \{1, 2, \dots, T_0\}$ and the post-treatment period $t \in \{T_0 + 1, \dots, T\}$. D indicates if an individual is treated, so that we only have $D = 1$ for treated units after T_0 . In addition to the outcomes y_{it}^D , there is data on K time-invariant covariates in the $(K \times 1)$ vector $\mathbf{x}_i = (x_i^{(1)}, x_i^{(2)}, \dots, x_i^{(k)})'$. We can now define formally the unobserved individual treatment effects as $\tau_{it} = y_{it}(D = 1) - y_{it}(D = 0)$ and so the Average Treatment Effect on the Treated (ATT) at time t is:

$$\tau_t = \frac{\sum_{i=n_0+1}^N \tau_{it}}{n_1} = \frac{\sum_{i=n_0+1}^N [y_{it}(1) - y_{it}(0)]}{n_1}$$

where we use ATT at time t and τ_t interchangeably below.

In the empirical application to the 1980 Mariel Boatlift with PSID data, for instance, the outcome variable is the wages of individuals for the period 1974 – 1984, i.e. $T = 11$. On the other hand, the treated individuals are people who live in Miami whereas the donors are people who live elsewhere in the US. More generally, we can neatly illustrate this set-up via the matrix of outcomes $y_{it}(D)$, called Θ , as:

⁷Thereby ruling out staggered adoption.

$$\Theta \equiv \left(\begin{array}{cccc|cccc} y_{1,1}(0) & y_{1,2}(0) & \dots & y_{1,T_0}(0) & y_{1,T_0+1}(0) & \dots & y_{1,T}(0) & \\ y_{2,1}(0) & y_{1,2}(0) & \dots & y_{2,T_0}(0) & y_{2,T_0+1}(0) & \dots & y_{2,T}(0) & \\ \vdots & \ddots & \ddots & \vdots & \vdots & \dots & \vdots & \\ y_{n_0,1}(0) & \dots & \dots & y_{n_0,T_0}(0) & y_{n_0,T_0+1}(0) & \dots & y_{n_0,T}(0) & \\ \hline y_{n_0+1,1}(0) & y_{n_0+1,2}(0) & \dots & y_{n_0+1,T_0}(0) & y_{n_0+1,T_0+1}(1) & \dots & y_{n_0+1,T}(1) & \\ y_{n_0+2,1}(0) & y_{n_0+2,2}(0) & \dots & y_{n_0+2,T_0}(0) & y_{n_0+2,T_0+1}(1) & \dots & y_{n_0+2,T}(1) & \\ \vdots & \ddots & \ddots & \vdots & \vdots & \dots & \vdots & \\ y_{n_0+n_1,1}(0) & y_{n_0+n_1,2}(0) & \dots & y_{n_0+n_1,T_0}(0) & y_{n_0+n_1,T_0+1}(1) & \dots & y_{n_0+n_1,T}(1) & \end{array} \right) \quad (2)$$

Note that Θ has four panels. The top-left panel contains the outcome variables for the donors in the pre-treatment period ($\mathbf{Y}_{n_0}^{pre}$) whereas the top-right panel is filled with the same values for the post-treatment period ($\mathbf{Y}_{n_0}^{post}$). The pre-treatment outcome variables for the treated observations are in the bottom-left panel ($\mathbf{Y}_{n_1}^{pre}$) and the post-treatment values for the treated group are in the bottom-right panel ($\mathbf{Y}_{n_1}^{post}$). Thus, we can rewrite Θ as a block matrix:

$$\Theta = \left(\begin{array}{c|c} \mathbf{Y}_{n_0}^{pre}(0) & \mathbf{Y}_{n_0}^{post}(0) \\ \hline \mathbf{Y}_{n_1}^{pre}(0) & \mathbf{Y}_{n_1}^{post}(1) \end{array} \right) \quad (3)$$

Our objective is to estimate the ATT of the policy occurring at time T_0+1 . In the post-treatment period $t > T_0$, we only observe the treated outcomes for the treatment group $\mathbf{Y}_{n_1}^{post}(1)$ and not the untreated outcomes for the treatment group $\mathbf{Y}_{n_1}^{post}(0)$ which is the Fundamental Problem of Causal Inference that prevents the within-person comparison. However, in order to calculate τ_i , we need to observe both. The SC method tackles this data limitation by estimating the unobserved $\hat{\mathbf{Y}}_{n_1}^{post}(0)$ using the rest of the information in Θ and then utilises these estimates to calculate individual treatment effect τ_{it} for

the post-treatment period. In this sense, it turns the causal inference problem into a prediction problem as we are trying to predict the missing values of $\mathbf{Y}_{n_1}^{post}(0)$. This is an extremely powerful insight, as it allows us to augment causal inference with cutting-edge techniques from statistics and machine learning that are great at predicting out-of-sample values.

Thus, what we are actually trying to estimate is the unobserved matrix with untreated values for all individuals:

$$\widehat{\Theta}(0) = \left(\begin{array}{c|c} \mathbf{Y}_{n_0}^{pre}(0) & \mathbf{Y}_{n_0}^{post}(0) \\ \hline \mathbf{Y}_{n_1}^{pre}(0) & \widehat{\mathbf{Y}_{n_1}^{post}(0)} \end{array} \right) \quad (4)$$

and the key question is how to use the information from the observed panels, namely $\mathbf{Y}_{n_0}^{pre}(0)$, $\mathbf{Y}_{n_0}^{post}(0)$ and $\mathbf{Y}_{n_1}^{pre}(0)$ to predict $\widehat{\mathbf{Y}_{n_1}^{post}(0)}$. At this point, one may wonder why we are bothering with SC in the first place, given that we are faced with a problem that requires good prediction of missing values in a matrix and we have an abundance of techniques for such situations in the matrix completion literature (?). One answer is that SC methods are focused on predicting values in a particular block of the matrix whereas many matrix completion techniques do not assume that data is missing in such a systematic way and so aim at fitting the whole matrix. So, SC methods can be more efficient, as they exploit the systematic ways in which data is missing. We discuss this question further in Section 4.a after we introduce a potential DGP for Θ .

3.c The Synthetic Control Method

The set-up in the previous subsection raises the question how we can impute the missing counterfactuals and so tackle the Fundamental Problem of Causal Inference via approximating the within-person comparison. In the context of $n_1 = 1$ and long T_0 , [Abadie et al. \(2010\)](#) introduced the SC method for cases with a single treated unit. As discussed in Section 2, their method assumes that the outcome vector of the treated

unit $y_{n_0+1,t}$ can be represented as a weighted average of n_0 donors' outcomes:

$$y_{n_0+1,t} = \sum_{j=1}^{n_0} \hat{w}_j y_{jt} + \epsilon_{n_0+1,t} \quad (5)$$

where n_0+1 indicates the single treated observation as it would be shown in Θ and the weight on donor j is given by \hat{w}_j . In addition, $\epsilon_{n_0+1,t}$ are idiosyncratic shocks that are independent and have mean 0. Since we are taking a weighted average, the weights are constrained to be between 0 and 1 and to sum up to 1. As such, we can interpret the weights as probabilities. To give an example, [Born et al. \(2019\)](#) construct a synthetic Britain before 2016 to study the effect of Brexit. Their chosen synthetic Britain is made up of 51% US, 17% Italy, 14% New Zealand, 11% Hungary, 5% Germany and other countries with a weight of 1% or less and matches pretty closely the true Britain.⁸

We may then wonder how the weights on donors are estimated. Equation (5) points towards the idea that we are essentially the time-series for the treated unit on the time-series for the donors in the pre-treatment period, except that the coefficients are constrained to be non-negative and sum up to 1. Thus, one possibility would be to rewrite the SC model in (5) as a constrained optimisation problem:

$$\min_{w_j} \sum_{t=1}^{T_0} \left(y_{n_0+1,t} - \sum_{j=1}^{n_0} w_j y_{jt} \right)^2 \quad s.t. \quad \sum_{j=1}^{n_0} w_j = 1 \quad w_j \geq 0 \quad (6)$$

where the two constraints ensure that the weights can be interpreted as probabilities. So, the formulation amounts to a constrained regression problem. Note that the literature has disagreed on whether there are better ways to estimate w_j which also take account of covariates ([Botosaru and Ferman, 2019](#)). See Appendix A.a for details. In any case, (6) shows that estimating the weights can be done in relatively simple way in the case of 1 treated unit.

⁸See [Figure 2](#) in [Borjas \(2017\)](#)

3.d Many treated units

While the previous section considered the case of $n_1 = 1$ and implicitly assumed T_0 is not short, many interesting applications in empirical microeconomics involve having $n_1 > 1$ and short T_0 . This section illustrates why implementing SC is not straightforward in this case. Similarly to Section 2, the problems can be made clear via a particular example:⁹ Card (1990) studies the effect of the Mariel Boatlift, a massive wave of Cuban immigration to Miami in 1980, on natives' wages. However, the same type of problems hold more generally when using SC method in applied microeconomics settings.

Suppose that we have panel data on many treated workers in Miami and on many control workers in other cities that did not experience the treatment of immigration. In such a case, many applied researchers will proceed by using difference-in-difference (DiD) estimator:

$$y_{it} = \rho + \gamma_i + \delta_t + \tau D_{it} + e_{it}$$

where i denotes an individual, ρ is an intercept, γ_i is an individual fixed effect, δ_t is a time fixed effect, D_{it} is the treatment indicator and τ is the ATT. Let us now assume that instead of DiD we would like to estimate the treatment effect via SC. One motivation for this could be that we believe that the parallel trends assumption fails or if we suspect heterogeneous treatment effects, as discussed in Section 2. Given the panel matrix of outcomes Θ , we are faced with two possibilities: either create a *separate synthetic control* for every treated individual in Miami or construct a single *pooled synthetic control* for all treated individuals in Miami that matches well the average wage in Miami. Let us consider each option.

Firstly, we can create a *separate SC* for every treated individual. The benefit of this approach is that we can obtain an individual treatment effect which can allow us to explore how the causal effects vary across different groups, e.g. for low-skilled versus high-skilled workers. However, as discussed in Section 2, we can quickly run into two problems: multiplicity of solutions and overfitting which we can now define

⁹See Section 6 for more details

more formally. The first issue renders SC method infeasible (Abadie and L’Hour, 2020), as in applications we often have a short T_0 , small K and many donor units n_0 . To see why SC will fail in this case, note that the SC for treated unit i is weighted average of donors. With many donors n_0 , it is likely that more than one weighted average of donors matches exactly the short vector of outcomes and covariates. For simple numerical examples of this issue, consider the example in Appendix A.c or the Motivating Example in Section 2.

Moreover, we may also be concerned with overfitting with short T_0 and small K . While SC constraints that require coefficients to sum up to 1 and to be non-negative ensure sparsity of the solution similarly to Lasso (Chernozhukov et al., 2020), the SC can be interpreted as a constrained regression of a vector of outcomes for treated unit i on the same vectors for the donors. However, when T_0 and K are small, we will have very little observations in our constrained regression. As a result, especially with big n_0 , overfitting is likely to occur.

Fortunately, one can still overcome the problem of multiplicity of solutions via changing the objective function in such a way that it picks the optimal SC out of all perfect SCs based on some criteria. This is the approach taken by Abadie and L’Hour (2020) who consider a similar setting as in this paper. Their estimator selects the SC which has the closest values of the matched variable to the actual treated observation. The example in Appendix A.c provides an illustration. However, they are still solving a complicated quadratic problem for every treated observation separately. Thus, their estimator can still suffer from overfitting, depending on what variables are included. In the simulation exercise and the empirical application, we further explore the properties of their proposal.

Secondly, we may consider calculating a (single) *pooled SC* for all treated individuals in Miami, i.e. calculate a single set of weights. This approach has the benefit of not running into multiplicity of solutions and overfitting unless n_0 is large relative to $n_1 \times T$, as we will be solving a constrained quadratic optimisation problem with $n_1 \times T$ observations rather than just n_1 as in *separate SC*. However, with many treated individuals, the single SC will not be matching too well the time series of some outlying

observations. For example, individual who earns a very low wage will get the same SC as an individual who earns a very high wage. While this is a serious limitation, it illustrates that if we can reduce the heterogeneity across individuals by classifying them into groups, then we can create SCs within each group and tackle multiplicity of solutions and overfitting.

Moreover, under certain conditions, the *pooled SC* method nests as a special case a common approach for policy evaluation of the Mariel Boatlift (Peri and Yasenov, 2018). We can call this approach *city-level synthetic control*: aggregate the individual level data on a city-level and then run SC on the city-level. So, we are creating SCs for the average wage in Miami by combining other US cities. Appendix A.d provides a set of conditions, under which the *pooled SC* reduces to *city-level SC*. Intuitively, we show in Appendix A.d that the weights which *city-level SC* can select is a subset of the weight which *pooled SC* can select. Unfortunately, the *city-level SC* also has certain limitations. In particular, inference remains a challenge in this case, as we observe just one estimate of the treatment effect. More broadly, inference with SCs is still work in progress (Chernozhukov et al., 2020). On the other hand, when we estimate separate SCs, an estimate of the treatment effect for every treated individual n_1 is available and so inference can be straightforward under mild assumptions.¹⁰ Furthermore, when T_0 is short, using SC methods on aggregate units is not recommended by Abadie et al. (2010), as the estimate of the ATT can be very biased.¹¹

So, it seems that neither the *pooled SC*, nor the *separate SC* are without problems. However, they motivate our proposed estimator presented in the next subsection, namely the Correlated Synthetic Control (CSC), as it balances between the two extremes. This insight for balancing *pooled SC* and *separate SC* has also been exploited in the *partially pooled SC* proposed very recently by Ben-Michael et al. (2021)

¹⁰Albeit our empirical application in Section 6.d show that there might be some complications.

¹¹In addition, when aggregating, we may be losing the heterogeneity of treatment effects. In contrast to *separate SC* and *pooled SC*, we cannot explore how the effect varies for different groups. One solution to this problem is to estimate different SCs for every group of interest, e.g. low-skilled vs high-skilled workers. I thank Barbara Petrongolo for pointing this out to me. However, when we suspect that there is considerable heterogeneity across many categories or even across continuous covariates, then simply dividing our dataset into groups may quickly become infeasible.

who made their paper public while we were working on CSC. However, the distinction which the authors draw between *pooled SC* and *separate SC* is slightly different from the one we consider here. In particular, we and Ben-Michael et al. (2021) understand *separate SC* similarly: create a separate synthetic counterfactual for every treated individual. However, we differ in how we understand *pooled SC*. For us, *pooled SC* refers to a SC-type of estimator that gives every treated individuals the same set of weights. In contrast, Ben-Michael et al. (2021) understand *pooled SC* to be creating separate synthetic controls for every treated individual. Crucially, instead of matching the individual time-series of a person, their *pooled SC* is picking weights that match the average time-series for all treated individuals. In a certain sense, we are pooling the *weights* across individuals whereas they are pooling the *outcomes* of treated individuals. As a result of this and other differences¹², the two estimators are related but remain different in important ways.

3.e Correlated Synthetic Controls

The key distinction between *separate SC* for each unit and *pooled SC* is analogous to a key distinction in the panel data literature between fixed coefficients models and pooled OLS. If we want to allow for individual-specific coefficients β_i in a linear regression model $y_{it} = \beta_i x_{it} + \epsilon_{it}$ without an intercept, we can run separate regressions for every observation i . This is equivalent to adding an interaction between the covariates and dummies for every observation as $y_{it} = \sum_{j=1}^N \mathbb{1}\{j = i\} \beta_j x_{it} + \epsilon_{it}$. This model is sometimes called fixed *coefficients* model in analogy to the fixed *effects* models (Balestra and Krishnakumar, 2008). It is also equivalent to running a separate regression for every unit i , i.e. we run N regressions of the type $y_{it} = \beta x_{it} + \epsilon_{it}$ for every i separately. On the other extreme, we can run a pooled OLS regression where the coefficients are restricted to be the same across $\beta = \beta_1 = \dots = \beta_N$ as $y_{it} = \beta x_{it} + \epsilon_{it}$.

However, there is a middle ground between the two extremes of fixed coefficients

¹²Another key conceptual difference with Ben-Michael et al. (2021) is that they create an estimator for the staggered adoption case with multiple but not too many treated unit whereas we are focused on an estimator with non-staggered adoption with many treated unit. As a result, Ben-Michael et al. (2021) do not consider how issues such as overfitting and multiplicity of solutions affect the properties of their estimator. So, empirical researchers should choose between CSC and *partially pooled SC* estimator based on whether they are facing staggered adoption or many treated units.

and pooled OLS: (correlated) random coefficients models (Wooldridge, 2003; Suri, 2011; Hsiao and Peseran, 2008). Basically, random coefficient models allow us to capture the heterogeneity in β_i (in contrast to pooled OLS) without requiring a lot of data¹³ for consistency (in contrast to fixed coefficients). The idea is that we allow β_i to differ across i but in a deterministic way based on some time-invariant z_i : we can do this by specifying $\beta_i = \beta + \psi z_i$. These type of models are a generalisation of Chamberlain's (correlated) random *effects* models that allow only the intercept to depend on observables (Chamberlain, 1982; Crépon and Mairesse, 2008).

The main contribution of this MPhil thesis is to apply this idea to the weights that postulated treated group's potential outcome without treatment as weighted average of donors' outcomes. So, we shall assume that the weights for treated unit i follow such a (correlated) random coefficient model:

$$y_{it}(0) = \eta_i + \sum_{j=1}^{n_0} w_{ij} y_{jt}(0) + e_{it} \quad s.t. \quad (7)$$

$$w_{ij} = \underbrace{\omega_j}_{ind.-invariant} + \underbrace{\mathbf{x}_i \boldsymbol{\alpha}^j}_{ind.-specific} \quad (Random\ Coef.)$$

$$\sum_{j=1}^{n_0} w_{ij} = 1 \quad \forall j : w_{ij} \geq 0 \quad (SC\ Constraints)$$

where we also allow for an intercept η_i , following suggestions in Ferman and Pinto (2021) and Doudchenko and Imbens (2018). In the (*Random Coef.*) constraint, each weight w_{ij} has two parts: an individual invariant part ω_j and individual specific part $\mathbf{x}_i \boldsymbol{\alpha}^j$ where \mathbf{x}_i is a $(1 \times K)$ vector of covariates and $\boldsymbol{\alpha}^j$ is a $(K \times 1)$ vector of coefficients. The individual-invariant part ensures we are close to the pooled SC approach discussed above whereas the individual-specific allows us to introduce heterogeneity across weights as in separate SC. On the other hand, the (*SC Constraints*) constraint simply ensures that the solution \mathbf{w} contains weights that sum up to one and are greater than zero.

¹³This is a result of the incidental parameter problem: in $y_{it} = \beta_i x_{it} + \epsilon_{it}$ we need $T \rightarrow \infty$ as well, if we would like to estimate β_i consistently. However, in correlated random coefficients if we let $\beta_i = \beta + \psi z_i$ and z_i is a discrete variable with categories $\{1, 2, \dots, K\}$, we only need $T * n_k \rightarrow \infty$ for consistency of β_i where n_k is the number of people for which $z_i = k$. Or in other words we can get consistency only with $n_k \rightarrow \infty$.

We call this estimator Correlated Synthetic Controls (CSC). Intuitively, the synthetic controls of two individuals are similar (or correlated) if they are similar in terms of observables, as it was the case for small and big German cities in the Motivating Example. It is useful to consider another example:

Example 1. *Suppose we have three treated individuals $i \in \{Ed, Mihai, Yi\ Ying\}$ and two covariates: i) years of education and ii) marital status for being married, single or other. Firstly, holding education the same across them, CSC will yield separate SCs with different weights on donor j if our individuals have three different values of marital status. For instance, suppose Yi Ying is single, Ed is married and Mihai is other:*

$$w_{Yi\ Ying,j} = \omega_j + \alpha^{single} \neq w_{Ed,j} = \omega_j + \alpha^{married} \neq w_{Mihai,j} = \omega_j + \alpha^{other}$$

However, if they have the same marital status, CSC will result in a single set of weights for all of them (pooled SC). Secondly, let education differ across the treated individuals: say, Mihai has 11 years, Ed - 12 and Yi Ying - 17. Suppose also that Ed and Mihai share the same marital status that is different from Yi Ying. Then, Ed and Mihai will get similar weights on donors, albeit not exactly the same, as they are similar in terms of observables. In that sense, their SCs will be correlated. In contrast, Yi Ying will get a set of weights which is considerably different, given her covariates.

A couple of things should be remarked in light Example 1 and the discussion above. Firstly, the estimator balances between estimating one set of weights for all treated individuals and separate sets of weights for treated individuals. The reason is that if *Mihai* and *Ed* have exactly the same covariates, we will give them the same weights and so the same SC that balances between fitting well both of their time series simultaneously. Note that we implicitly assume that their time-series will not be too different, if they are similar in terms of observables. On the other hand, if *Yi Ying* has completely different values of the covariates, she will receive another set of weights which may be substantially different (especially, if we only use discrete covariates). If her covariates are similar to *Mihai* and *Ed*, although not the same, she will receive

similar donor weights.

Secondly, we can see how CSC solves the problem of multiplicity of solutions. Suppose that *Ed* and *Mihai* both have multiple solutions and have the same covariates but their time series are not exactly the same. Then in general the set of potential donor combinations that perfectly match Mihai’s time series will not be intersecting with the set of potential donor combinations that perfectly match Ed’s time series. Since we can only pick one set of weights for them, we will pick the one set of weights that creates a donor which matches simultaneously both of their time series as closely as possible but not perfectly. In that way, we tackle multiplicity of solutions, as by construction there does not exist a synthetic control which perfectly matches both of their time series at the same time.¹⁴

3.f Estimation

Next, one may wonder how the w_{ij} are estimated. Proposition 1 shows one way,¹⁵ in which we can rewrite the (correlated) random coefficients model in (7) for the outcomes y_{it}^0 as a constrained quadratic optimisation problem in a scalar form. One intuition for estimating the SC model with random coefficients (weights) is that we are essentially regressing the vector of outcomes for the treatment group on the outcomes for the donors and an interactions between donors’ outcomes and treatment group’s covariates, given some constraints on the coefficients. In (1) below, this can be seen as a constrained regression of y_{it} for the treated group on donors’ y_{jt} and an interaction between donors’ y_{jt} and treated group’s $x_i^{(k)}$. See Appendix A.b for a formulation in terms of the block matrices of Θ and for details on how we can use the package CVXR (Fu et al., 2019) to code the estimator in R.

Proposition 1. *The parameters $\alpha_j^{(k)}$ and ω_j can be found by solving the constrained*

¹⁴Recall that we are using the same set of donors for both Mihai and Ed.

¹⁵Given the abundance of methods for solving random coefficients models (Hsiao and Peseran, 2008), there are probably more efficient ways to estimate the model

optimisation problem for the pre-treatment period:

$$\begin{aligned} \max_{\alpha_j^{(k)}, \omega_j} \sum_{i=1}^{n_1} \sum_{t=1}^{T_0} \left(y_{it} - \eta_i - \sum_{j=1}^{n_0} y_{jt} \omega_j - \sum_{j=1}^{n_0} \sum_{k=1}^K \alpha_j^{(k)} y_{jt} x_i^{(k)} \right)^2 \quad s.t. \\ \forall i : \sum_{j=1}^{n_0} \left(\omega_j + \sum_{k=1}^K x_i^{(k)} \alpha_j^{(k)} \right) = 1 \quad \forall (i, j) : \omega_j + \sum_{k=1}^K x_i^{(k)} \alpha_j^{(k)} \geq 0 \end{aligned} \quad (8)$$

where y_{it} are the potential outcomes without treatment, η_i is an individual fixed effect, ω_j is the individual-invariant part of the weight on donor j , $x_i^{(k)}$ is the value of the k -th covariate for the treated observation i and $\alpha_j^{(k)}$ is the coefficient of covariate k in determining the weight on donor j .

Proof. All proofs are relegated to Appendix B. See Appendices A.b and B.a for the proof of this Proposition. \square

The main intuition for the proof is that we are substituting the (*Random Coef.*) constraint into a quadratic loss function and into (*SC Constraints*), as defined by (7).

Before proceeding to CSC's theoretical properties, it is important to discuss a seemingly big limitation of the estimator: it does *not* allow for continuous covariates.¹⁶ The reason is a restriction imposed by the fact that the random coefficients weights should sum up to 1. To gain some intuition, consider the example of $K = 3$ covariates with weights:

$$\sum_{j=1}^{n_0} (\omega_j + x_i^1 \alpha_j^1 + x_i^2 \alpha_j^2 + x_i^3 \alpha_j^3) = 1$$

which can be rewritten as:

$$x_i^{(1)} \sum_{j=1}^{n_0} \alpha_j^{(1)} + x_i^{(2)} \sum_{j=1}^{n_0} \alpha_j^{(2)} + x_i^{(3)} \sum_{j=1}^{n_0} \alpha_j^{(3)} = 1 - \sum_{j=1}^{n_0} \omega_j \quad (9)$$

The last expression holds for every treated i and the right-hand side is the same independent of i . If we have continuous $x_i^{(k)}$ and $i > 2$, then (9) will not hold in general for all i and the optimisation problem will be infeasible, as it is impossible to satisfy the constraint exactly. Nevertheless, suppose that x_i^k are dummies for a single categorical

¹⁶I would like to thank Anders Kock for encouraging me to pursue this line of thought.

variable with three mutually exclusive categories (e.g. married, divorced, other), then for the restriction to hold it is sufficient to have: $\sum_{j=1}^{n_0} \alpha_j^{(1)} = \sum_{j=1}^{n_0} \alpha_j^{(2)} = \sum_{j=1}^{n_0} \alpha_j^{(3)}$. Therefore, while the *sum* of the α^k will be constrained to be the same across the three categories k , the coefficient on the same donor j across different k and q groups, respectively $\alpha_j^{(k)}$ and $\alpha_j^{(q)}$, can be different. This is good news, as we can then achieve our objective of having heterogeneity in the weights on different donors.¹⁷

Regarding continuous predictors, it is still possible to integrate them by recoding such a variable (e.g. wages) as a discrete predictor (e.g. income brackets). However, there are approaches, allowing us to handle continuous covariates in a more systematic way. For example, we can pre-process continuous covariates prior to estimating CSC, as done in coarse exact matching (Iacus et al., 2012). The idea is that there is an extra step which allow us to balance the donors and the treatment groups in terms of some continuous observables. As a result, we do not need to worry about controlling for this particular predictor in CSC.

4 Theoretical properties

This section compares the mean absolute deviation (or estimation error) of the true ATT τ_s at time s from the estimated $\hat{\tau}_s^{did}$ from Difference-in-Difference (DiD) and from the estimated $\hat{\tau}_s^{csc}$ from Correlated Synthetic Controls (CSC) when the data is generated via an interacted fixed effects model.¹⁸ This is important, because it illustrates in what circumstances CSC should be preferred to DiD. Specifically, in the case of many treated units, Synthetic Control (SC) methods are not the default method in empirical work. Establishing when (if ever) they perform better than DiD is vital, if applied researchers are going to use them.

Thus, the main takeaway from this section is that CSC should be preferred in the cases when we believe that the treatment is correlated with unobservable characteristics, i.e. selection on unobservables, under the data generation process we consider.

¹⁷A similar restriction on the sums applies if $x_i^{(k)}$ are *not* mutually exclusive categories

¹⁸We use sometimes interactive fixed effects models and factor models interchangeably, although we try to prioritise the former term.

For this condition to hold, we also need the SCs to do a good job at predicting the pretreatment outcomes of the treated group. On the other hand, if the treatment is not strongly correlated with unobserved characteristics, DiD might be a better choice. We assume that our data is generated from an interactive fixed effects models. More generally, to the best of our knowledge there have been just a few papers exploring the theory behind SC techniques in the case of many treated units (Abadie and L'Hour, 2020; Ben-Michael et al., 2021). Unfortunately, these studies do not provide much guidance on when SC methods should be preferred over DiD. This section fills this gap by making a small step towards providing such conditions.

4.a Data Generating Process

As common in the literature on SC (Abadie et al., 2010; Ferman and Pinto, 2021), we shall assume that each y_{it} in Θ follows an interacted fixed effects model (Bai, 2009; Moon and Weidner, 2015; Hsiao, 2018):

$$y_{it} = \theta_t \mathbf{x}_i' + D_{it}\tau + v_{it} \quad (10)$$

$$= \theta_t \mathbf{x}_i' + D_{it}\tau + \boldsymbol{\lambda}_t \boldsymbol{\mu}_i + \epsilon_{it} \quad (11)$$

where \mathbf{x}_i is a $(1 \times K)$ vector of time-invariant covariates, θ_t are $(1 \times K)$ time-varying coefficients on these covariates, D_{it} is the treatment assignment, v_{it} is a composite error term such that $v_{it} = \boldsymbol{\lambda}_t \boldsymbol{\mu}_i + \epsilon_{it}$, $\boldsymbol{\lambda}_t$ is a $(1 \times F)$ vector of common factors, $\boldsymbol{\mu}_i$ is a $(F \times 1)$ vector of factor loadings and ϵ_{it} are idiosyncratic shocks. The main quantity of interest is τ of the policy.¹⁹ The *interactive* fixed effect structure $\boldsymbol{\lambda}_t \boldsymbol{\mu}_i$ is a generalisation of the traditional *additive* fixed effects $\lambda_t + \mu_i$.²⁰ We can interpret the interactive factor structure as saying that each individual i has some unobserved characteristics $\boldsymbol{\mu}_i$ such as ability or motivation that determine their outcome y_{it} . However, at different points in time, different unobserved factors from $\boldsymbol{\mu}_i$ matter, implying that their effects are time-varying which justifies including time-varying common factors $\boldsymbol{\lambda}_t$.

¹⁹The reason why τ is not the ATE is because we are not going to assume that treatment assignment is randomly assigned.

²⁰In fact, for $F = 2$, using the interactive fixed effects with $\boldsymbol{\mu}_i = (1, \mu_i)'$ and $\boldsymbol{\lambda}_t = (\lambda_t, 1)$ reduces to the additive fixed effects.

We will also impose some further structure on the DGP:

Assumption 1 (DGP restrictions). *We assume that:*

- (i) *Treatment-not-at-random: $\text{Cov}(D_{it}, \boldsymbol{\mu}_i) \neq \mathbf{0}_F \implies \text{Cov}(D_{it}, v_{it}) \neq 0$*
- (ii) *Errors are iid with $E[\epsilon_{it}] = 0$ and $E[\epsilon_{it}^2] \leq \infty$. They are independent from all other random variables.*
- (iii) *DGP of \mathbf{x}_i : \mathbf{x}_i is independent of ϵ_{it} and D_{it} but $\text{Cov}(\boldsymbol{\mu}'_i, \mathbf{x}_i) \neq \mathbf{0}_{F \times K}$*
- (iv) *$\boldsymbol{\mu}_i$ are stochastic with $E[\boldsymbol{\mu}_i] = \boldsymbol{\mu} \neq \mathbf{0}_F$*
- (v) *$\boldsymbol{\lambda}_t$ are fixed parameters with $\frac{1}{T} \sum_{t=1}^T \boldsymbol{\lambda}_t = \bar{\boldsymbol{\lambda}} \neq \mathbf{0}'_F$*

where $\mathbf{0}_F$ is a $(F \times 1)$ column vector of zeros, $\mathbf{0}_{F \times K}$ is a $(F \times K)$ matrix of zeros.

Let us detail what the different subparts of **Assumption 1** say. Firstly, the treatment indicator D_{it} is assumed to be correlated with the unobserved component $\boldsymbol{\mu}_i$ (**Assumption 1.i**). This means that essentially D_{it} is endogenous ($\text{Cov}(D_{it}, v_{it}) \neq 0$) and so we cannot estimate τ consistently by simply ignoring the composite structure of the error term and fitting (10). Secondly, we assume that for all individuals i and time periods t the errors ϵ_{it} are iid and zero mean with a finite second moment (**Assumption 1.ii**). So, we do not allow ϵ_{it} to follow more complicated autoregressive processes. Thirdly, we assume that \mathbf{x}_i are uncorrelated with treatment assignment **but** could be correlated with the composite error term v_{it} via the factor loadings (**Assumption 1.iii**). This means that \mathbf{x}_i are endogenous and we cannot estimate their coefficients directly. Lastly, we assume that $\boldsymbol{\mu}_i$ and $\boldsymbol{\lambda}_t$ are respectively stochastic and fixed with means $E[\boldsymbol{\mu}_i] = \boldsymbol{\mu}$ and $\bar{\boldsymbol{\lambda}}$ that are *not* constrained to be 0 necessarily (**Assumptions 1.vi** and **1.v**). The reason for this particular choice stems from a suggestion in Hsiao (2018, p.665) that in cases of short T and long N assuming fixed common factors and stochastic factor loadings is reasonable.

Given this DGP, a natural question is to ponder why bother with SC, given the abundance of estimators for interactive fixed effects models (Hsiao, 2018). For instance, Xu (2017) proposes the Generalised SC which directly fits such a model to the data in

Θ . More generally, we can simply model Θ directly via matrix completion techniques, as in ?.

There are at least three good reasons for not taking this approach. Firstly, in practice, we rarely know the true model generating Θ : it could be an interactive fixed effects model as in Xu (2017) but it could also be a vector autoregressive process (Abadie, 2020b). On the other hand, SC methods can be shown to work well under various DGP (Abadie et al., 2010; Ben-Michael et al., 2021). While with interactive fixed effects models we risk misspecifying the DGP of Θ , SC allows more flexibility with respect to the true DGP. So, we will not risk fitting the wrong model to a data that actually follows a vector autoregression.

Secondly, **Assumption 1.i** allows for a non-zero correlation between our only time-varying covariate D_{it} and the factor loadings μ_i . We can then rewrite model (11) as

$$y_{it} = D_{it}\tau + (\theta_t \quad \lambda_t) \begin{pmatrix} x'_t \\ \mu_i \end{pmatrix} + \epsilon_{it}$$

where we treat θ_t as common factors and x_i as factor loadings. This allows us to apply **Remark 5.9.** from Hsiao (2018) which states that $\hat{\tau}$ will not be estimated consistently by common approaches for estimating interactive fixed effects models. Thus, even if we wanted to use an interactive fixed effects models, it will not estimate the quantity of interest consistently.

Thirdly, SC focuses on imputing the bottom-right panel of Θ rather than predicting all of its entries as many matrix completion techniques would do. Given that often we are not interested in fitting the rest of Θ , SC can be more efficient as it exploits the structure of the matrix and focuses on a smaller prediction task. Of course, if we were faced with a more general matrix where, for example, different individuals are treated at different times, then matrix completion methods will be more appropriate.

With these considerations in mind, we conclude that assuming an interactive fixed effects as DGP and applying SC methods might not be a bad idea.

4.b CSC Bound

Let us consider how we can establish an upper bound on the estimation error of CSC. By estimation error, we mean the absolute value difference between the estimated ATT $\hat{\tau}$ and the true τ , i.e. $|\hat{\tau} - \tau|$, which is also equivalent to the Mean Absolute Error for the case of a single parameter of interest. Unfortunately, given how complicated the optimisation problem in Proposition 1 is there is generally no obvious analytical solution for the weights and so we cannot derive an exact bound for the estimation error. For this reason, we will construct an upper bound for the estimation error, drawing on High-Dimensional Statistics. We will then compare this upper bound to the exact bound, derived for DiD. This comparison can then be interpreted as conservative: even under the worst possible estimation error for CSC, there are still certain conditions, under which it should be chosen over DiD.

Before establishing the bound, we need to make several additional assumptions. Firstly, we will make Exact Fit assumption which states that CSC will find a weighted average of donor units that exactly match treated individual i in terms of both outcomes and covariates. Versions of this assumption for the case of a single treated unit are common in the SC literature, e.g. see [Abadie et al. \(2010\)](#) and [Botosaru and Ferman \(2019\)](#).

Assumption 2 (Exact Fit). *Consider a $(n_0 \times n_1)$ matrix \mathbf{W} such that each weight w_{ji} follows a (correlated) random coefficient model $w_{ji} = \omega_j + \sum_{k=1}^K \alpha_{kj} x_i^{(k)}$ and each column of \mathbf{W} sums up to 1 with every entry being non-negative. Matrix \mathbf{W} satisfies exact fit in **pre-treatment period** if:*

$$(i) \text{ For the outcome variable: } \forall t \in \{1, 2, \dots, T_0\} : \quad y_{it} = \sum_{j=1}^{n_0} w_{ij} y_{jt}$$

$$(ii) \text{ For the covariates: } \forall k \in \{1, 2, \dots, K\} : \quad x_i^k = \sum_{j=1}^{n_0} w_{ij} x_j^{(k)}$$

Note, however, that our statement of the assumption is slightly different from Abadie et al.'s original assumption ([2010](#)). The reason is that when generalised to the many treated unit settings, they assume the existence of *a set* of weights satisfying exact fit for each i . This might be problematic, due to the multiplicity of solutions (See

Section 2). If Abadie's assumption holds, then our Exact Fit assumption will hold, as we can simply make x_i^k be a vector of dummies for each observation. Therefore, one may see our assumption as an extension of Abadie's that also tackles multiplicity of solutions by not assuming a set of weights.

Under the Exact Fit assumption, we can prove an extremely useful lemma that generalises a result from [Abadie et al. \(2010\)](#) for the case of many treated units. This is Lemma 1 Abadie's representation which allow us to rewrite the estimation error without the unobserved component μ . Note that we need to assume invertability of matrix $\lambda'_{pre}\lambda_{pre}$ and a necessary condition for this is $T_0 > F$ which can be interpreted to mean that our DGP is not too complicated relative to how many time periods we observe.

Lemma 1 (Abadie's representation). *Assume that:*

- (i) *The true DGP is the interactive fixed effects model in (10)*
- (ii) *Assumption 2. Exact Fit holds*
- (iii) *The $F \times F$ matrix $\lambda'_{pre}\lambda_{pre}$ is invertible*

Then, we can write for $s > T_0$:

$$\begin{aligned} \hat{\tau}_s - \tau_s = & \frac{1}{n_1} \left[\lambda_s (\lambda'_{pre}\lambda_{pre})^{-1} \lambda'_{pre} \sum_{i=n_0+1}^N \left(\sum_{j=1}^{n_1} \hat{w}_{ij} \epsilon_{j,pre} - \epsilon_{i,pre} \right) \right] \\ & + \frac{1}{n_1} \left[\sum_{i=n_0+1}^N \left(\epsilon_{is} - \sum_{j=1}^{n_0} \hat{w}_{ij} \epsilon_{js} \right) \right] \end{aligned}$$

where \hat{w}_{ij} is the random coefficient weight of donor j on individual i , λ_s are the common factors at post-treatment time s , $\lambda'_{pre}\lambda_{pre}$ is the $F \times F$ matrix of interacted pre-treatment common factors, $\epsilon_{j,pre}$ is a $(T_0 \times 1)$ vector of error terms for observation i in the pretreatment period and ϵ_{js} is the error term for individual j at time s .

Proof. See Appendix B.b for the proof. □

The proof of the lemma involves rewriting the estimation error in terms of the DGP for the outcome variables y_{it} and substituting out the μ_i , using the Exact Fit

assumption. While this may seem complicated, it only consists of several tedious algebraic steps. One other thing to be remarked about the previous expression is that we do not need any assumptions on the DGP beyond the fact that it follows an interacted factor model.

Next, we also assume that the errors ϵ_{it} are iid `subGaussian`(σ^2).²¹ While this assumption may seem restrictive, it only constraints the distribution not to have fat tails relative to a Normal distribution with variance σ^2 . Moreover, it allows us to draw on techniques from High-Dimensional Statistics which are extremely useful for bounding different quantities such as the estimation error in our case (Rigollet, 2015).

Assumption 3 (SubG Errors). ϵ_{it} are iid `subG`(σ^2)

Lastly, we make two further assumptions: the matrix $\lambda'_{pre}\lambda_{pre}$ in the pre-treatment period is invertible (so that we can use Lemma 1 Abadie's Representation) and we work on an Euclidean space (which is useful, as in the proof we need to calculate operator norms of matrices). Then, we can find an upper bound on the estimation error for $\hat{\tau}$ with a high probability:

Proposition 2 (CSC Bound Estimation Error). *Suppose that:*

- (i) *DGP is given by interactive fixed effects model in (10) with Assumption 1.*
- (ii) *Assumption 2 Exact Fit holds.*
- (iii) *Assumption 3 SubG Errors holds.*
- (iv) *$F \times F$ matrix $\lambda'_{pre}\lambda_{pre}$ is invertible*
- (v) *We work on an Euclidean space*

Then, with probability at least $1 - \frac{3}{\exp(0.25h^2)}$ in post-treatment period s the estimation error for the ATT estimated by Correlated Synthetic Control satisfies for $h > 0$:

$$|\hat{\tau}_s^{csc} - \tau_s| < \left(\frac{F\lambda_{max}^2}{\phi_{min}} \sqrt{\frac{2\sigma^2 n_0}{T_0} + \frac{h\sigma}{T_0^{1.5}}} \right) + \left(\frac{h\sigma F\lambda_{min}^2}{\phi_{max}^2} \right) + h\sigma\sqrt{2} \quad (12)$$

²¹See Rigollet (2015) for a more detailed discussion of `subGaussian` variables

where λ_{min} and λ_{max} denote respectively the minimum and maximum common factor λ_{fs} in absolute value for either pre-treatment or post-treatment period. Similarly, ϕ_{min} and ϕ_{max} are respectively the minimum and maximum eigenvalue of matrix $\frac{1}{T_0} \mathbf{\lambda}' \mathbf{\lambda}$

Proof. See Appendix B.d. □

The strategy for the proof is that we use Abadie's representation from Lemma 1 for the estimation error. Then, we use various properties of **subG** variables to bound each quantity in the representation. Note that when proving the theoretical result we focus on CSC *without* an intercept in the objective function for simplicity.

Let us now interpret the expression for the upper bound of CSC (12). Firstly, the parameter h controls the probability with which the bound obtains. The bigger h , the less precise our bound, but the higher the probability it holds. On the other hand, when we increase T_0 or decrease F , our estimate of τ becomes better. This makes intuitive sense, as adding more observations or decreasing the complexity of the DGP should allow us to estimate the weights more precisely. In particular, the first term in the bound will disappear as $T_0 \rightarrow \infty$. Furthermore, the upper bound increases with σ^2 which is approximately the variance of the strictly exogenous error terms in the DGP.²² One reason why this might be happening is that with fixed N and T_0 our SCs can match well the part $\theta_t \mathbf{x}_i + \lambda_t \mu_i$ of the outcomes y_{it} but with bigger σ^2 it gets difficult to match the idiosyncratic shocks ϵ_{it} . So, this result is not so surprising.

The role of n_0 in the bound is ambiguous: one may expect that the more donors we have, the better our SCs will be. However, given how we have constructed the proof, the *upper* bound increases in n_0 , so that the more donors we have, the bigger the estimation error. We shall inspect further the importance of n_0 in the simulation in the next section. Next, let us consider the importance of the common factors λ_{tf} . It is clear that we do not want the biggest common factor $\lambda_{max} = \max_{t \in (1, \dots, T_0, \dots, T), f \in (1, \dots, F)} |\lambda_{tf}|$ to be too large. This requires that no particular time period t experienced a very large shock as reflected in the common factor. In applications, empirical researchers can leverage their domain knowledge to check if this is the case.

²²Approximately, because we assume the ϵ_{it} are **SubG**(σ^2) and so their second moment can very well be smaller.

Lastly, before moving forward, note that the second term in the bound $\left(\frac{h\sigma F\lambda^2}{\phi_{max}^2}\right)$ will usually be quite small, as we are dividing the smallest common factor by the largest eigenvalue. It is, thus, less of a concern than controlling the other two terms.

4.c Estimation Error of DiD

Although the previous section found an *upper* bound for the estimation error of $\hat{\tau}$ for CSC, in the case of DiD we can find an exact analytical solution for the estimation error because the analogical optimisation problem for DiD to the one in Proposition 1 for CSC is unconstrained. So, we can use a more conventional approach such as Frisch-Waugh-Lovell's theorem to derive the asymptotic bias in the ATT without appealing to High-Dimensional Statistics.

We assume that τ^{did} is estimated by the model $y_{it} = \rho + \gamma_i + \delta_t + D_{it}\tau + u_{it}$. In addition, we make the simplifying assumption that we are observing only one post-treatment period, i.e. $T = T_0 + 1$, which considerably simplifies one step in the proof. Thus, Proposition 3 gives an exact asymptotic expression for the estimation error of τ^{did} , as the number of treated units n_1 goes to infinity:

Proposition 3 (DiD Estimation Error). *Suppose:*

1. *DGP is given by the interacted fixed effects model in (14) with Assumption 1.*
2. *$T = T_0 + 1$, i.e. one post-treatment period*
3. *We estimate τ^{did} via $y_{it} = \rho + \gamma_i + \delta_t + D_{it}\tau + u_{it}$*

Then, as $n_1 \rightarrow \infty$ the estimation error for DiD is:

$$|\hat{\tau}^{did} - \tau| \xrightarrow{p} \left| \frac{(\bar{\lambda}_{pre} - \lambda_T)(\bar{\mu}_{don} - E[\mu_i|D_{it} = 1])}{n_0} \right| \quad (13)$$

where $\bar{\lambda}_{pre}$ is $(1 \times F)$ vector of the average value of the common factors in the pre-treatment period, λ_T is $(1 \times F)$ vector of the common factors in the single post-treatment period, the $(F \times 1)$ vector $\bar{\mu}_{don}$ is the average of the factor loadings for the donors and $E[\mu_i|D_{it} = 1]$ is the $(F \times 1)$ vector with expected value of the factor loadings in the treatment group, .

Proof. The proof is in Appendix B.e. □

The main strategy for the proof is to apply the Frisch-Waugh-Lovell Theorem to τ in the two-way fixed effects model and obtain an exact expression for the estimation error. Essentially, (13) gives the expression for the asymptotic bias of $\hat{\tau}^{did}$ when the potential outcomes are generated by an interactive fixed effects. In a sense, the interactive fixed effects assumption generalises the parallel trends assumption which is the main assumption requires for DiD to work. We shall return to the parallel trends assumption in the next section. In any case, the main thing that (13) is telling us is that if there are difference in unobservables μ_i between the treatment group and the control group, then $\hat{\tau}^{did}$ will be biased asymptotically. This would be even more problematic, if we suspect that the common factors experience a structural break between the pre-treatment and post-treatment period. Note that the reason why we have $\bar{\mu}_{don}$ as an average and $E[\mu|D_{it} = 1]$ is an expectation is because we let $n_1 \rightarrow \infty$ and hold n_0 constant.

Expression (13) begs one interesting question: would the estimation error disappear if we also let $n_0 \rightarrow \infty$. Assuming we let both $n_1, n_0 \rightarrow \infty$ but $\frac{n_1}{n_0} \rightarrow c$, then it can be shown that the estimation error will disappear.²³ In applications, however, justifying such asymptotics with respect to both n_1 and n_0 might often be unreasonable: in the German cities example from Section 2 we have $n_1 > n_0$ but both are relatively small ($n_1 = 20$ and $n_0 = 99$) whereas in the Mariel Boatlift example we have $n_1 = 41$ and $n_0 = 1039$. Moreover, if we only let $n_0 \rightarrow \infty$ and hold n_1 fixed, we will get a similar result to (13), meaning that there will be some estimation error. Thus, in practice, it is likely that the bias will not disappear.

Let us now return to the main question of this section: when should CSC be used over DiD. Firstly, CSC would have a smaller estimation error for $\hat{\tau}$ when the treatment assignment D_{it} is strongly correlated with the unobserved factor loadings μ_i . The reason is that even the conservative bound for CSC is independent of this correlation. This points to one of the biggest advantages of SC methods more generally. We construct synthetic individuals with similar characteristics to the treated units and we "throw

²³See the proof to Proposition 3 and in particular equations (40) and (42). The quantity in the first equation (40) would go to infinity whereas (42) would go to some constant. Then, the estimation error would go to 0.

away” bad control units that are very different for our treated individuals. However, unless we apply some additional correction to DiD, these bad control units will bias our estimate of $\hat{\tau}^{DiD}$, as they will form a part of our control group. For instance, in the empirical application to the Mariel Boatlift (Section 6), we are interested in constructing SCs for low-skilled Miami workers. Since Florida is one of the bottom 20% of US states in terms of median wages, we are probably looking at low-wage workers at a low-wage state. It is likely that without any further restrictions on the donor pool DiD will be biased due to a correlation between treatment-assignment and unobserved characteristics of Miami workers: we are including high-wage workers from high-wage states which are not relevant for the estimation of τ but remain a part of the donor pool. So, in that particular case, with individual-level data CSC might be a better idea.

On the other hand, we can see certain similarities between the two expressions for the estimation errors. For example, if the common factors in $\boldsymbol{\lambda}$ are very volatile, both $\hat{\tau}^{did}$ and $\hat{\tau}^{csc}$ can have a big estimation error. So, perhaps, if we suspect that the common factors have experienced structural break, it would be more appropriate to time-series models or other techniques. So, there are circumstances, in which both DiD and CSC will do badly.

Lastly, recall that for CSC to work well we do need to the Exact Fit Assumption. If it fails, then the bound in Proposition 2 will not hold in general. While there is some theoretical work exploring what happens to the single-treated unit case when the assumption fails (Botosaru and Ferman, 2019; Ferman and Pinto, 2021), the case of many treated units when it fails is a fruitful area of further research.

5 Simulations

In this section, we perform a simulation exercise to continue our comparison between Correlated Synthetic Controls (CSC) and Difference-in-Difference (DiD). Additionally, we also contrast their properties against the Penalised Synthetic Control (PSC) of Abadie and L’Hour (2020) which is defined below. In addition, we propose the *infea-*

sible DiD estimator which is a version of DiD that estimates the parameters of interest consistently but is not feasible, as it requires knowledge of unobserved components. Using the infeasible DiD as a benchmark, we shall see under what conditions CSC approximates its performance. In addition, we provide evidence that the framework in Section 3.b which turns the causal inference problem into a prediction problem for the missing values in Θ is indeed a sensible approach for constructing estimators.

To that aim, we begin by specifying the DGP. As in the previous section, we assume that the outcomes in Θ follows the interactive fixed effects models:

$$y_{it} = \beta \mathbf{x}_i' + D_{it}\tau + \lambda_t \mu_i + \epsilon_{it} \quad (14)$$

that was described in the previous section. The only substantial difference that we make relative to Section 4 is that now the coefficients on \mathbf{x}_i are *not* time-varying but are time-invariant, i.e. $\forall t : \beta = \beta_t$. This is helpful, as it allows us to easily introduce the infeasible DiD estimator. In addition, we make the same **Assumption 1**, as in Section 4.

5.a Bias of DiD

Before proceeding with the simulation, we can show why DiD will provide a biased estimate of τ under the DGP, which supplements Proposition 3 that proves the inconsistency of $\hat{\tau}^{did}$. The reason for DiD's biasedness is that it fails the parallel trends assumption which is necessary for the identification of the ATT. As discussed in the previous section, clarifying DiD's assumptions is of vital importance for applied researchers because DiD is often preferred in empirical applications with many treated units (Arkhangelsky et al., 2020, p.2).²⁴

Let us now show the parallel trends assumption more formally (Wooldridge, 2021):

$$E[y_{it}(0) - y_{i1}(0)|D_{is} = 1] = E[y_{it}(0) - y_{i1}(0)|D_{is} = 0]$$

²⁴I would like to thank Barbara Petrongolo and Frank DiTraglia for convincing me of the importance to discuss this questions.

where $y_{it}(0)$ is a potential outcome of individual i at time t without treatment, t indicates either a pre-treatment or post-treatment period ($t \in \{1, 2, \dots, T_0, \dots, T\}$) and D_{is} indicates treatment assignment in some post-treatment period s ($s \in \{T_0 + 1, T_0 + 1, \dots, T\}$). Intuitively, we need the trend in the treatment group to be the same as in the control group. Next, we can plug-in the DiD model $y_{it}(0) = \rho + \gamma_i + \delta_t + u_{it}$ into the assumption:

$$\overline{E[\delta_t - \delta_0 | D_{is} = 1]} + E[u_{it} - u_{i0} | D_{is} = 1] = \overline{E[\delta_t - \delta_0 | D_{is} = 0]} + E[u_{it} - u_{i0} | D_{is} = 0] \quad (15)$$

where in the first row the expressions with δ cancel as they are made up of fixed expressions independent of treatment assignment. Thus, we need the last expression to hold true, in order to identify the $\hat{\tau}^{did}$.

We can now check if (15) holds with data generated from an interactive fixed effect model (as given by (11)) and Assumption 1. We begin by adding and subtracting the mean of factor loadings $\boldsymbol{\mu}_i$ and common factors $\boldsymbol{\lambda}_t$, i.e. demeaning them:

$$\begin{aligned} y_{it} &= \beta \mathbf{x}_i + D_{it}\tau + (\boldsymbol{\lambda}_t - \bar{\boldsymbol{\lambda}} + \bar{\boldsymbol{\lambda}})(\boldsymbol{\mu}_i - E[\boldsymbol{\mu}_i] + E[\boldsymbol{\mu}_i]) + \epsilon_{it} && \iff \\ y_{it} &= \underbrace{(\bar{\boldsymbol{\lambda}} E[\boldsymbol{\mu}_i])}_{\rho} + \underbrace{(\beta \mathbf{x}_i + \bar{\boldsymbol{\lambda}} \boldsymbol{\mu}_i)}_{\gamma_i} + \underbrace{(\boldsymbol{\lambda}_t E[\boldsymbol{\mu}_i])}_{\delta_t} + \tau D_{it} + \underbrace{\tilde{\boldsymbol{\mu}}_i \tilde{\boldsymbol{\lambda}}_t + \epsilon_{it}}_{u_{it}} \end{aligned} \quad (16)$$

where we define the $(F \times 1)$ vector $\tilde{\boldsymbol{\mu}}_i = (\boldsymbol{\mu}_i - E[\boldsymbol{\mu}_i])$ as the demeaned factor loadings and $\tilde{\boldsymbol{\lambda}}_t = (\boldsymbol{\lambda}_t - \bar{\boldsymbol{\lambda}})$ are the demeaned common factors with $\bar{\boldsymbol{\lambda}}$ being the $(1 \times F)$ vector of means and $\tilde{\boldsymbol{\mu}}_i - E[\boldsymbol{\mu}_i]$ are the demeaned factor loadings with $E[\boldsymbol{\mu}_i]$ being the $(F \times 1)$ vector of means. Note that the second expression illustrates how the potential outcome y_{it} can be represented as the additive fixed effects structure assumed by the parallel trends assumption of DiD.

Nevertheless, the condition required by Parallel Trends in (15) reduces to:

$$E[\boldsymbol{\mu}_i | D_{is} = 0] \neq E[\boldsymbol{\mu}_i | D_{is} = 1]$$

after cancelling the common factors and using the structural model's DGP in (16) and

Assumption 1.²⁵ However, since $u_{it} = \tilde{\boldsymbol{\mu}}_i \tilde{\boldsymbol{\lambda}}_t + \epsilon_{it}$ and the demeaned factor loadings $\tilde{\boldsymbol{\mu}}_i$ are not independent from treatment assignment (Assumption 1), the last equality will not hold in general. Thus, the parallel trends assumption does not hold and τ^{did} is biased.

5.b Infeasible DiD

So, DiD is both inconsistent²⁶ and biased because $\boldsymbol{\mu}_i$ enters the error term u_{it} in (16) which generates a non-zero correlations between the error term u_{it} and D_{it} . However, what if we can control for this non-zero correlation? Then τ^{DiD} would be consistent and unbiased. In particular, if we subtract the demeaned interacted fixed effects $\tilde{\boldsymbol{\mu}}_i \tilde{\boldsymbol{\lambda}}_t$ from both sides of (16), then DiD should be consistent. This is clearly infeasible, as we do not observe $\tilde{\boldsymbol{\mu}}_i \tilde{\boldsymbol{\lambda}}_t$. Nevertheless, since we are running a simulation, we can define the Infeasible Difference-in-Difference estimator (iDiD).

Definition 1 (Infeasible DiD). *The infeasible $\hat{\tau}^{idid}$ is given by subtracting the demeaned interactive fixed effects from the outcome variables:*

$$y_{it} - \tilde{\boldsymbol{\mu}}_i \tilde{\boldsymbol{\lambda}}_t = \alpha + \delta_t + \gamma_i + D_{it}\tau + u_{it}$$

where $\tilde{\boldsymbol{\mu}}_i \tilde{\boldsymbol{\lambda}}_t$ are the demeaned interactive fixed effects.

Given this definition, it is possible to show that $\hat{\tau}^{idid}$ will estimate consistently the true quantity of interest τ .

Proposition 4 (Consistency of iDiD). *Suppose that:*

- (i) *Data is generated by $y_{it} = \boldsymbol{\beta} \mathbf{x}'_i + \boldsymbol{\lambda}_t \boldsymbol{\mu}_i + D_{it}\tau + \epsilon_{it}$ under Assumption 1.*
- (ii) *$\hat{\tau}^{idid}$ is given by Definition 1. Infeasible DiD*

Then, Difference-in-Difference will estimate ATT consistently: $\hat{\tau}^{idid} \xrightarrow{P} \tau$

²⁵For more theoretical results on DiD with an interacted fixed effects model, please refer to Proposition 3 for an explicit expression for the asymptotic bias of $\hat{\tau}^{did}$ relative to the true τ .

²⁶There follows from Proposition 3. However, we can show it in another way. Estimating $y_{it} = \rho + \gamma_i + \delta_t + \tau D_{it} + u_{it}$ will yield an inconsistent estimate of the true τ , since $Cov(D_{it}, u_{it}) \neq 0$, meaning that D_{it} is endogenous. Section A.e in the Appendix discusses this point further.

Proof. The proof is in Appendix B.f. □

The main strategy for the proof is to use the Frish-Waugh-Lovell theorem for $\hat{\tau}^{idid}$ and it is analogical to the more involved proof for the exact bias of *feasible* DiD in Proposition 3.

Our key motivation for including iDiD in the simulation is that it provides a strong benchmark, against which we can compare the other models. So, if the feasible DiD performs very similarly to iDiD, this would be evidence that the bias induced by non-zero $Cov(\boldsymbol{\mu}_i, D_{it})$ is not too serious and so there is not much benefit to using SC techniques.

5.c Distributional Assumptions on DGP

In order to generate data from the interactive fixed effects model under Assumption 1, we need to make some further distributional assumptions on how the stochastic parameter are determined. The full DGP for the simulation is specified as:

$$\begin{aligned}
(\textit{Outcome}) \quad & y_{it} = \beta \mathbf{x}'_i + \boldsymbol{\lambda}_t \boldsymbol{\mu}_i + D_{it} \tau + \epsilon_{it} \\
(\textit{Common Factor}) \quad & \boldsymbol{\mu}_i = \boldsymbol{\gamma} \mathbf{x}'_i + \mathbf{v}_i \\
(\textit{Shock to } \mu_i) \quad & \mathbf{v}_i \stackrel{iid}{\sim} N_F(\boldsymbol{\mu}, \mathbf{I}_F) \\
(\textit{Idiosyncratic shock}) \quad & \epsilon_{it} \stackrel{iid}{\sim} N(0, 1) \\
(\textit{Categorical Predictor}) \quad & x_i^{(1)} \stackrel{iid}{\sim} U[1, 2, \dots, K] \\
(\textit{Continuous Predictor}) \quad & x_i^{(2)} \stackrel{iid}{\sim} N(0, 1)
\end{aligned} \tag{17}$$

Let us now detail what this DGP suggests. For $\boldsymbol{\Theta}$, we assume that every outcome follows an interactive fixed effects model, as in (*Outcome*). In addition, the common factors $\boldsymbol{\mu}_i$ are allowed to be correlated with the time-invariant covariates, as required by Assumption 1.(iii). Moreover, $\boldsymbol{\gamma}$ is a fixed $1 \times K$ vector controlling the correlation between the factor loadings and the time-invariant covariates, e.g. if we set $\boldsymbol{\gamma} = 0$ then there is no correlation. We allow $\boldsymbol{\mu}_i$ to have an expectation different from 0 since the mean of \mathbf{v} is not 0, as \mathbf{v} 's mean is actually $\boldsymbol{\mu}$. The idiosyncratic shocks ϵ_{it} are drawn from a standard normal distribution. On the other hand, we assume that there

are two covariates in the DGP: a categorical covariate $x_i^{(1)}$ and a continuous predictor $x_i^{(2)}$.²⁷ The motivation for including both continuous and discrete covariates stemmed from the restriction on CSC that only allows for discrete predictors.²⁸ So, including $x_i^{(2)}$ allows us to explore what happens to CSC when we simply recode a continuous predictor as a binary one, given that CSC cannot handle categorical predictors (see Section 3.f). On the other hand, we assume that $x_i^{(1)}$ is drawn from a discrete uniform distribution with K categories, meaning that every observation is equally likely to belong to any category in $\{1, 2, \dots, K\}$. Lastly, we assume v_i , u_{it} , x_i^1 and x_i^2 are all uncorrelated with each other, so that we have no correlations between them.

The parameters β , λ and γ are treated as fixed in the assumptions of the model. So, they are constant across all simulations. In particular, we set the vector β with coefficients for $x_i^{(1)}$ and the $H = 5$ categories of $x_i^{(2)}$ as:

$$\beta = (1, 0.4, 0.6, 0.8, 1, 1.2)$$

For the values of the $(F \times T_0)$ matrix λ and the $(F \times H)$ matrix γ , please refer to Appendix A.f. In general, after generating some initial values for these parameters, we hold them to be the same across all of our simulations and for all estimators that we compare in order to avoid making the performance of the different estimators dependent on the particular values being generated in each case.²⁹

Lastly, treatment assignment D_{it} is generated to be correlated with unobservables, following a similar approach to Arkhangelsky et al. (2020, p.14). Given the set-up in Section 3.b. the treatment can only be 1 in the last T_0 periods. The treatment indicator is drawn from a Bernoulli distribution with mean π_i . We introduce a dependence of treatment assignment D_{it} on μ_i via π_i . While π_i is independent from the observables

²⁷Both of them are time-invariant in order to avoid the need to make transformations such as taking the average of

$$x_i^{(2)}$$

which is required by SC-like methods. Integrating time-varying predictors in SC methods and more generally the place of covariates in SC is still work in progress (Botosaru and Ferman, 2019).

²⁸See Section 3.f

²⁹For instance, suppose we decide to generate different values these parameters when simulating the data for different estimators. However, we may get very unlucky with our draws of values for iDiD but very lucky with our values of CSC, so that in the end it may seem that CSC performs better but in reality this is simply due to the randomness.

in \mathbf{x}_i , it is related to the factor loadings $\boldsymbol{\mu}_i$ according to a hierarchical model:

$$D_{iT}|\boldsymbol{\mu}_i, v_i \sim \text{Ber}(\pi_i) \quad \pi_i = \frac{\exp(\boldsymbol{\mu}_i\boldsymbol{\phi} + \varepsilon_i)}{1 + \exp(\boldsymbol{\mu}_i\boldsymbol{\phi} + \varepsilon_i)} \quad (18)$$

where ε_i are some iid $N(0, 1)$ shocks. If $\boldsymbol{\phi} \neq 0$, then D_{it} does not occur at random, as it is correlated with the unobserved $\boldsymbol{\mu}_i$. We can control this correlation via our choice of $\boldsymbol{\phi}$. If set to 0, then treatment occurs independently of factor loadings. This matters because studies assuming random treatment assignment tend to overestimate how well causal inference techniques such as DiD tend to estimate the true treatment effect (Arkhangelsky et al., 2020), e.g. in terms of the deviation of the estimate from the true treatment effect. Moreover, since in practice with observational studies treatment is unlikely to be assigned completely at random, it seems reasonable to assume that D_{it} is not independent of unobserved characteristics, as reflected in $\boldsymbol{\mu}_i$.

After specifying the details of the DGP, the simulation will contrast and compare four estimators: the feasible difference-in-difference (fDiD), the infeasible difference-in-difference (iDiD), the Correlated Synthetic Controls (CSC) and the Penalised Synthetic Control (PSC) of Abadie and L'Hour (2020) which is CSC's closest sibling from the family of SC estimators. The motivation for including PSC is that it is the only other estimator known to us that addresses the same set-up of many treated units by developing a SC estimator. **Algorithm 1** summarises the flow of the simulation.

Algorithm 1: Pseudo-code for the flow of our simulation	
<pre> for j <i>in</i> $1 : 1000$ do Generate each $y_{it}(j)$ in the $(N \times T)$ matrix $\boldsymbol{\Theta}(j)$ and K covariates $\mathbf{x}_i(j)$ if <i>treat at random</i>(j) == <i>TRUE</i> then Set $\boldsymbol{\phi} = 0$ else Set $\boldsymbol{\phi} \neq 0$ end Divide $\boldsymbol{\Theta}(j)$ into pre-treatment data and post-treatment data Fit <i>CSC</i>, <i>feasible DID</i>, <i>infeasible DID</i> and <i>PSC</i> in pre-treatment dataset Estimate performance metrics in post-treatment period end </pre>	

Each of the 1000 replications begins by simulating the matrix $\boldsymbol{\Theta}$. Next, the treat-

ment is assigned either at random or not at random via the choice of ϕ . After the data is divided into pre-treatment and post-treatment datasets, we implement the four estimators and estimate the ATT $\hat{\tau}$ for each of them. Afterwards, we calculate two performance metrics: 1) the estimation error of the estimated ATT $\hat{\tau}$, i.e. the difference $\hat{\tau} - \tau$; and 2) RMSE between the synthetic $\hat{y}_{it}(0)$ and the true unobserved potential outcome $y_{it}(0)$ for the treatment group, i.e. $\forall i \in \{n_0 + 1, n_0 + 2, \dots, n_0 + n_1\}$ in the post-treatment period $\forall t \in \{T_0 + 1, \dots, T\}$. Note that this will not be possible with actual data, as we do not observe the two potential outcomes simultaneously due to the Fundamental Problem of Causal Inference, nor do we know the true ATT. Nevertheless, if indeed the matrix completion framework for matrix Θ discussed in Section 3.b is appropriate, then we would expect the results for the point estimate of ATT and the RMSE of the synthetic controls to be similar.

Before moving forward, the reader may wonder how PSC (Abadie and L'Hour, 2020) works. As mentioned in Section 3.d, it creates separate synthetic controls for every treated observation in a way that tackles multiplicity of solutions. Out of all possible combinations of donors that exactly fit the time-series for treated observation i , we chose the one combination, in which the donors are most similar to i in terms of observables in addition to matching the weighted average of unit i . We can show this more formally via the optimisation problem solved by PSC. For every treated unit i in $\{n_0 + 1, n_0 + 2, \dots, N\}$, we have in scalar form³⁰:

$$\min_{\mathbf{w}_i \in \Delta} \underbrace{\sum_{t=1}^{T_0} \left(y_{it} - \sum_{j=1}^{n_0} y_{jt} w_{ij} \right)^2 + \sum_{k=1}^K \left(x_i^{(k)} - \sum_{j=1}^{n_0} x_j^{(k)} w_{ij} \right)^2}_{\text{Synthetic Control}} + \underbrace{\lambda \sum_{j=1}^{n_0} w_{ij} \left(\sum_{t=1}^{T_0} (y_{it} - y_{jt})^2 + \sum_{k=1}^K (x_i^{(k)} - x_j^{(k)})^2 \right)}_{\text{Penalty}}$$

where λ is a tuning parameter that can be chosen via cross-validation.³¹ The Synthetic

³⁰For the purpose of clarity, we diverge from the widely used notation in Section 3

³¹When we have data on outcomes y_{it} for $T_0 \geq 2$, we can calculate n_1 (penalised) synthetic controls using data on the time-invariant covariates x_i^k and the first $T_0 - 1$ outcomes y_1 . We can then evaluate the performance of the n_1 different SC at time T_0 .

Control part of the objective function ensures that the weighted average matches the treated unit’s outcomes and covariates. However, on its own, it is not sufficient to select a unique solution. So, the penalty part ensures that we select the set of donors that are also similar to the treated unit in terms of observable characteristics (and not just with a similar weighted average), i.e. PSC prefers similar donors rather than different donors. A numerical example of how PSC works is presented in Appendix A.c.³²

5.d Results

The results for the average estimation error between the true τ and the estimated $\hat{\tau}$ are presented in Table 2. Since we report results for the same quantity of interest as in Section 4, that is *avr. est. error* = $\frac{\sum_{r=1}^{1000}(\hat{\tau}_r - \tau)}{1000}$, we can interpret Table 2 in light of the theoretical properties we derived.³³

Table 2 begins with a baseline specification where we set $N = 100$ and $T = 5$ for the dimensions of the matrix Θ , a three-factor model $F = 3$, treatment-not-at-random, i.e. $Cov(D_{it}, \mu_i) \neq 0$, and the average probability of treatment is set at 0.15, or in other words we set $E[\pi_i] = 0.15$ in (18), so that we expect with $N = 100$ on average $n_1 = 15$ treated observations and $n_0 = 85$ donors. In Table 2, we progressively change the values of the five parameters (random treatment allocation, N , π_i , T , and F) relative to the baseline specification.

Several intriguing things can be seen from the results in Table 2. Firstly, across all specifications the iDiD dominates the other estimators, which is unsurprising given that it is unbiased and consistent in contrast to other estimators. However, across the feasible estimators, CSC tends to outperform fDiD and PSC with PSC usually performing a bit worse than CSC. In particular, an (average) estimation error of -0.11 in the baseline case³⁴ for CSC means that with the true τ being 1, CSC estimates

³²As in the case of CSC, we coded the optimisation problem in R using CVXR package (Fu et al., 2019).

³³Except that here we calculate the *average* estimation error over 1000 simulated datasets rather than the *single* estimation error for one dataset.

³⁴For reasons of time, I was unable to rerun the same analysis using the absolute value of the estimation error which would have made the results from the Simulation and the Theoretical Section exactly comparable.

Table 2: *Average estimation error for ATT ($\frac{\sum_{r=1}^{1000}(\hat{\tau}_r - \tau)}{1000}$) under different DGPs*

Change	CSC	fDiD	PSC	iDiD
Baseline*	-0.11	-0.94	-0.25	-0.01
Less Correlation	-0.05	-1.05	-0.22	-0.02
Random Assignment ($\phi = 0$)	0.05	-0.01	-0.03	-0.01
Increase $N = 150$	-0.04	-0.95	-0.24	-0.01
Increase $N = 200$	-0.03	-0.96	-0.18	-0.01
Increase $E[\pi_i] = 0.25$	-0.07	-0.86	-0.23	0.00
Increase $E[\pi_i] = 0.40$	-0.12	-0.79	-0.36	-0.01
Decrease $T = 4$	2.43	3.53	4.71	0.00
Increase $T = 8$	0.09	0.50	0.31	-0.01
Decrease $F = 2$	-0.31	-2.18	-0.78	-0.01
Increase $F = 4$	-0.05	-0.94	-0.30	-0.01

*Baseline refers to specifying {Non – Random Assign., $N = 100, T = 6, F = 3, E[\pi_i] = 0.15$ }. All tables were created with R package **stargazer** (Hlavac, 2018)

$\hat{\tau}^{CSC} = 0.89$ on average across 1000 simulations. So, CSC is doing quite well and even with a relatively small sample of just 100 observations it performs similarly to iDiD. Secondly, let us consider what happens once we make treatment less correlated with the unobserved factor loadings. As suggested by Proposition 3, fDiD indeed does better under random treatment assignment than under non-random treatment assignment and outperforms CSC. However, even under a weak correlation between D_{it} and μ_i , fDiD still does not do too well relative to CSC. So, even if we suspect the treatment is allocated slightly not at random, CSC should be used, confirming a key prediction made by our theoretical results.

Thirdly, increasing N has little effect on the estimation errors for either one of the four estimators. So, perhaps we need to experiment with bigger values of N to see any effects or consider changing the share of treated observations, as done in the next panel of the table. However, increasing the probability of treatment π_i also has no substantial effect on our results. This is seemingly surprising in light of Proposition 2 which implies that the bias of CSC is positively related to the number of donors n_0 . We may expect the estimation error of CSC to fall, as we decrease the proportion of donors, holding total N fixed. Nevertheless, we can explain this by the fact that what

we found in Proposition 2 is an upper bound. So, perhaps, changing the number of donors n_0 is not such a bad thing for CSC.

Fourthly, the results for changing T in the case of CSC are as expected: it improves its performance, even though we mostly gain from increasing $T = 4$ to $T = 6$. The next set of results from the simulation in Table 3 paints a clearer picture of what happens if we increase T . Lastly, the results for changing the number of factors F are puzzling. The theoretical properties of fDiD and CSC in Section 4 suggest that the estimation error of both is increasing in F but here we do not find any significant differences when we change F .

Before moving forward, note that while the point estimate for the estimation error of $\hat{\tau}$ look very close to its true value for CSC, its variance might be very big. The reason why the variance is not accounted for in the estimation error is a consequence of our decision to work with the unadjusted difference between the true and estimated ATT. For this reason, Appendix A.g reports the results for the *RMSE* between the true τ and the estimated $\hat{\tau}$ rather than the pure estimation error from the same simulation exercise, as discussed above. Recall that the RMSE is $\frac{\sum_{r=1}^{1000}(\hat{\tau}-\tau)^2}{1000}$ and the estimation error is just $\frac{\sum_{r=1}^{1000}(\hat{\tau}-\tau)}{1000}$. Overall, the results appear to be largely unchanged.

Besides studying our estimates of the ATT, we may also be interested in how well the synthetic controls resemble the unobserved potential outcomes for the treated units without treatment in the post-treatment period (Table 3). As in the previous table, we begin with the baseline specification {Non – Random Assign., $N = 100, T = 6, F = 3, E[\pi_i] = 0.15$ } and progressively change some of the key parameters. The results are very similar to those for the estimation error. For example, while iDiD creates the best possible synthetic control, the best performing estimator among those that are feasible is CSC in all cases when treatment does not occur at random. This is good news, as it suggests that the framework in 3.b is meaningful: we can indeed turn the causal inference problem into a prediction problem. One other important result from Table 3 is the fact that we can observe more clearly what happens as we increase T : all estimators improve significantly and start to approximate the RMSE for iDiD. With $T = 8$, CSC and fDiD perform very similarly. While in practice we may often observe

shorter time periods, this result confirms the fact that SC methods work very well with long panels (Abadie et al., 2010).

Table 3: *RMSE between \hat{y}_{it}^0 and true y_{it}^0 under different DGPs across 1000 simulations*

Change	CSC	FDID	PSC	iDiD
Baseline*	1.80	2.28	2.51	0.92
Weak Correlation	1.75	2.24	2.41	0.92
Random Assignment	1.75	2.03	2.35	0.92
Increase $N = 150$	1.80	2.28	2.36	0.92
Increase $N = 200$	1.80	2.29	2.30	0.92
Increase $E[\pi_i] = 0.25$	1.84	2.26	2.63	0.92
Increase $E[\pi_i] = 0.40$	1.90	2.24	2.75	0.92
Decrease $T = 4$	5.40	5.93	9.86	0.88
Increase $T = 8$	1.33	1.34	2.12	0.94
Decrease $F = 2$	1.86	2.97	2.71	0.92
Increase $F = 4$	1.95	2.43	2.98	0.92

*Baseline refers to specifying $\{\text{Non} - \text{Random Assign.}, N = 100, T = 6, F = 3, E[\pi_i] = 0.15\}$

6 Empirical Application

This section illustrates how the Correlated Synthetic Controls (CSC) estimator can be used in practice: we study the effect of immigration on natives' labour market outcomes. We hope to gain some new insight on this question by revisiting the Mariel Boatlift (Card, 1990) and analysing it with a new dataset: the Panel Study of Income and Dynamics (PSID). The motivation for this exercise stems from the fact that recent work by Clemens and Hunt (2019) has identified key compositional changes to the conventional data sources used for evaluating the effects of the Mariel Boatlift. However, given the small sample size of PSID, we should be careful, when interpreting the results, as their external validity might be questioned. In any case, we illustrate how researchers can obtain heterogeneous treatment for the effects of immigration on different type of workers. In addition, we perform a Cross-Validation exercise, showing that CSC does slightly better at predicting treated group's outcomes than its main alternative from the SC family - the Penalised Synthetic Control (PSC) of Abadie and

L'Hour (2020).

6.a Mariel Boatlift

The textbook model of labour markets with perfect competition (Carlin and Soskice, 2014) postulates that immigration will increase labour supply and so will bring down the wages of natives. This can be seen in Figure 1 where wages fell from w down to w' following the shift of the Labour Supply (LS) curve. However, testing the hypothesis for a negative effect on wages raises many econometric challenges. Suppose we specify the model:

$$y_{it} = \eta + \tau m_{it} + \beta x_{it} + u_{it} \quad (19)$$

where i indexes a local labour market³⁵, y_{it} are log wages (although it could be another labour market outcome), η is an intercept, m_{it} is the number of immigrants, x_{it} are time varying characteristics and u_{it} are error terms. The parameter of interest is τ which is the effect of immigration on wages and theory predicts that it should be negative. Arguably the biggest econometric challenge with (19) is the fact that, immigrants self-select into areas with greater economic opportunities, suggesting that $\text{Corr}(m_{it}, u_{it}) > 0$ which will make m_{it} endogenous and so the OLS estimate $\hat{\tau}$ inconsistent.³⁶

Card's Mariel Boatlift study (1990) overcomes this identification challenge by exploiting the effect of a large exogenous shock to labour supply. On 20th April 1980, Fidel Castro announced that Cubans would be allowed to leave the country and migrate to the US from the Mariel port. As a result, almost 125,000 Cubans fled the country and arrived in the US before the policy was reversed in late September. Many of them settled in Florida (and Miami in particular), because previous Cuban immigrants have established themselves there and not because of the economic opportunities there. Thus, the Mariel Boatlift provides an exogenous increase of 8% to labour supply

³⁵We could base i on geography or national education-experience cells, for example.

³⁶Other challenges stem from how we define local labour markets. If we define i to be geographical local labour markets such as counties, we can be concerned about displacement effects (Borjas, 2017). On the other hand, if we use national education-experience cells, we should still think about what instrument to use to tackle the selection effects (Dustmann et al., 2016).

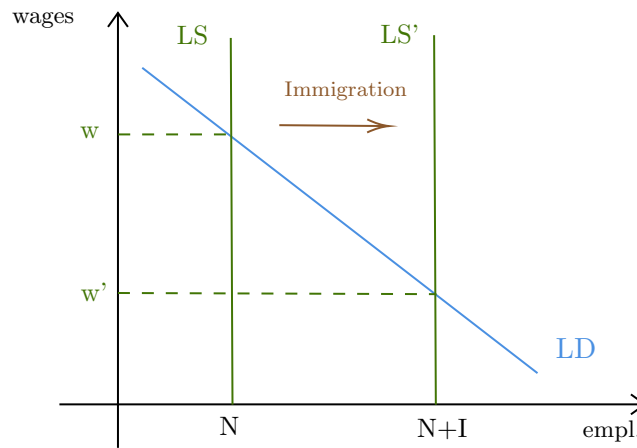


Figure 1: The effect of immigration in the classical labour market model with inelastic labour supply and flexible wages. *Note:* We have inelastic LS, due to perfect competition between firms: if a firm decreases its wages, then all of its workers can leave and immediately find a job, meaning that wage elasticity of labour supply is infinite Manning (2006).

in Miami and so an excellent opportunity to test the model in Figure 1. The reason is that the choice of settling in Miami was unrelated to the economic opportunities there.

Card evaluates the effect of immigration via a difference-in-difference (DiD) approach, studying how wages and unemployment of natives evolved in Miami compared to a control group of cities in the pre-treatment period and post-treatment periods. He finds that while wages in Miami fell for white, black and Hispanic native workers between 1979 and 1983, they also fell in the control group by a similar amount. Thus, he concludes that there was no negative effect of the increase in labour supply to wages, contrary to the prediction of the model in Figure 1.

Over the past five years, there has been a revival³⁷ in interest in evaluating the Mariel Boatlift (Borjas, 2017; Peri and Yasenov, 2018). These recent papers have found markedly different results: Borjas (2017) claims that the increase in immigration had a negative effect on natives' wages whereas Peri and Yasenov (2018) find no effects. However, all of these papers point to some important limitations of Card's original study:

1. Firstly, the choice of control group consisting of Atlanta, Los Angeles, Houston and Tampa-St. Peterbourg seems a bit ad hoc and is not well justified (Borjas, 2017). Since the publication of the paper, better techniques for constructing

³⁷For an excellent review (much clearer than mine), click [here](#).

counterfactuals such as Synthetic Control have been developed and should be preferred.

2. Secondly, and perhaps most fundamentally, the Miami CPS datasets, versions of which are used in all of [Card \(1990\)](#), [Borjas \(2017\)](#), and [Peri and Yasenov \(2018\)](#), have undergone significant compositional changes that were not related to the Mariel Boatlift ([Clemens and Hunt, 2019](#)). In particular, the number of low-skilled black workers in Miami with very low wages surveyed have been increased substantially in the 1980s, thereby giving a greater survey weight to this group in early 1980s relative to late 1970s. This renders CPS inappropriate for comparison across different years, especially given that the treatment occurs at 1980, given the important changes in survey weights.
3. Thirdly, since the CPS data is a repeated cross-section, it has to be aggregate on a city-level in order to find the trend in wages in Miami. This can mask important individual level differences in labour supply that might have resulted from the Mariel Boatlift, e.g. at the intensive margin rather than the extensive margin. Note that the classical labour markets model in [Figure 1](#) does not allow any labour supply effects, as the LS is assumed to be inelastic. However, given the strong evidence for sticky wages, some labour supply effects can be expected ([Dustmann et al., 2016](#)).

6.b PSID

Based on the second critique, it seems that using CPS data should be avoided and this is the reason why we revisit the Mariel Boatlift using the Panel Study of Income and Dynamics (PSID). When we combine PSID with CSC, we can also automate the selection of a control group (thus, tackling the first problem) and exploit PSID's panel structure to explore changes in labour supply at the intensive margin (thus, tackling the third problem).

The PSID is the longest running panel survey of households, stemming from late 1960s to the present day. It contains longitudinal data on roughly 5000 families, living

in the US. PSID's composition did not change considerably which allows us to tackle problem 2. above. Moreover, the fact that it is panel study opens the door to using CSC which allows us to construct good counterfactuals for every treated worker and which tackles problem 1. The PSID sample that we use consists of 11 waves (1974-1984) and is restricted to heads of households who are male. The justification for these restrictions and the variable selected can be found in Appendix A.h.

Unfortunately, using PSID comes at a cost: our final sample contains only 42 workers. However, while our results may not be generalisable for the overall effect of the Mariel Boatlift, we can still interpret them as yielding useful insights for our particular sample and worry about their external validity separately. Moreover, the sample sizes used in Borjas (2017) and Peri and Yasenov (2018) also contain only respectively around 20 and 65 people per year, suggesting that small sample sizes is a persistent issue when estimating the effects of immigration.³⁸

Perhaps, a seemingly more fundamental issue with PSID is that it does not include metropolitan area indicators but only state indicators in its public release. So, we use as treatment group not residents of Miami but those of the state of Florida. This is problematic, as the impact of the Mariel Boatlift was most strongly felt in Miami. Nevertheless, there is some evidence that some Mariel Cubans also settled in other areas in Florida such as Tempa and West Palm Beach³⁹, as shown by Skop (2001). Previous Cuban immigrants (e.g. Golden Exiles in 1960-64 and Freedom Flight refugees in 1965-1974) were mostly white and had higher socio-economic status than immigrants in 1980 (McCoy and Gonzalez, 1985, p.457). In contrast, Mariel Cubans were more likely to be non-white and more generally had different locational preferences, regarding where to settle (Skop, 2001). Quantitatively, Miami did host the majority of Cuban immigrants, but there were spillovers of migrants to other cities in Florida. Specifically, 1990 US Census data shows that more than 10% of Mariel Cubans based in Florida were not located in Miami but in the early 1980s the number was higher due to within-Florida

³⁸However, note that they define low-skilled workers as workers without a high school degree whereas we include workers with a high school degree into low-skilled workers, implying that their sample sizes could have been greater if they included the latter group.

³⁹At the time, Palm Beach was not considered a part of Miami Metropolitan Area (US Department of Commerce, 1982, p.30). Since this is the definition that Card's data uses (1990, footnote), he did not include Mariel immigrants located there.

migration to Miami in the 1980s (Skop, 2001, p.464). Lastly, if we are able to find a negative effect for low-skilled workers in Florida, then it is likely that this effect would be even stronger for the same workers in Miami. Overall, then, the lack of metropolitan area identifiers do not invalidate our analysis. Although the small sample size remains a concern, the analysis can still yield valuable information.

6.c Penalised vs Correlated Synthetic Controls

The first set of results with PSID that we present compares CSC and PSC using a cross-validation approach (Abadie et al., 2015). We begin by dividing the data into post-treatment data (1980-1984) and pre-treatment data (1975-1979) which is further divided into (pre-treatment) training data and (pre-treatment) testing data. In all specifications, the testing period consists of the last two pre-treatment years, namely 1978 and 1979, but we vary the length of the training period. Then, PSC and CSC are fit on the training data and we evaluate their performance by calculating the Root Mean Squared Error (RMSE) between the predicted labour market outcome (either wages or hours worked) and the true outcomes. Lastly, we calculate the individual treatment effects and the overall ATT which we store.⁴⁰

The results from this cross-validation approach are reported in Table 4. Overall, CSC performs slightly better than PSC for $T_{train} = 3$ and $T_{train} = 2$, as it predicts labour supply more accurately in the training period (1979-1980). In light of the simulation results in Section 5, this could reflect the fact that sample size is big enough so that PSC approaches CSC or that treatment assignment is not too strongly correlated with unobserved components.⁴¹

Interestingly, as we increase T_{train} , CSC becomes better at predicting total hours worked but worse at predicting (log) wages. This could perhaps be explained by the fact that wages are more volatile than hours worked which exhibits stronger autocorrelation.

Note that our best prediction for log wages involves using $T_{train} = 1$ and for hours

⁴⁰When fitting the models, we use the covariates for occupation, industry, education, being white, being married and having being ill a lot that were coded as discrete variables, so that CSC can accommodate them, as discussed by the restriction in Section 3.f.

⁴¹However, note that *some* correlation is expected, given that Florida is in the bottom quantile in terms of median wage and that we are interested primarily in the effect on wages of the low-skilled workers. See the discussion in the end of Section 4.c for more details

Table 4: RMSE in testing period for different lengths of training period

	<i>Year</i>	<i>Lab. Supply</i>		<i>Log wages*</i>	
		<i>CSC</i>	<i>PSC</i>	<i>CSC</i>	<i>PSC</i>
$T_{train} = 4$	1978	569.92	666.43	0.63	0.57
	1979	592.62	670.27	0.86	0.76
	<i>Total</i>	581.38	668.35	0.75	0.67
$T_{train} = 3$	1978	616.17	663.19	0.56	0.57
	1979	638.05	687.91	0.77	0.73
	<i>Total</i>	627.21	675.67	0.67	0.65
$T_{train} = 2$	1978	637.04	655.14	0.49	0.60
	1979	677.68	725.55	0.69	0.68
	<i>Total</i>	657.67	691.24	0.60	0.64
$T_{train} = 1$	1978	739.45	671.06	0.40	0.60
	1979	851.64	687.11	0.59	0.68
	<i>Total</i>	797.52	679.13	0.50	0.64

* We applied transformation $\log(1 + w_{it})$ to wages because some individuals had $w_{it} < 1$ so $\log(w_{it})$ yields a non-positive value. It would have been better to use the inverse hyperbolic sine transformation but time constraints prevented us for this.

worked using $T_{train} = 1$ and that both use CSC rather than PSC. Moreover, this result points to the fact that using more lags of the outcome variable is helpful only if it is informative for future values. Note, however, that this is not the case for PSC where the results seem largely invariant to the number of lags we include. Overall, we can conclude that CSC definitely does not do worse than PSC, when we compare the two estimators with this data, and our best performers in terms of predicting wages and hours worked utilise CSC.

6.d Heterogeneous Treatment Effects

In Figure 2, we present the treatment effects of the Mariel Boatlift on labour supply and log wages for low-skilled and high-skilled native workers. To get the individual treatment effects, we use the synthetic controls created by CSC in the previous section with $T_{train} = 4$ for labour supply and with $T_{train} = 1$ for wages. These choices of T_{train} seem the best option based on the results in Table 4. Moreover, results for other choices of T_{train} are provided in the Appendix A.i and are quite similar.

Figure 2 plots the evolution of the 95% for the estimated heterogeneous treatment effects on low-skilled workers and high-skilled workers over time. Following Acemoglu and Autor (2011), we define low-skilled workers as those individuals who do not have a college degree and so do not enjoy the college premium. We can see that there are no effects on any of high-skilled workers' outcomes or low-skilled workers' total hours worked. This is not surprising, given that many Mariels were classified as being low-skilled and so were competing with low-skilled natives for jobs rather than high-skilled natives (Skop, 2001). The lack of labour supply effects also implies that the LS curve in Figure 1 is probably quite inelastic, albeit not perfectly inelastic. However, there is a significant decrease in wages for low-skilled workers, which is surprising relative to previous studies using CPS (Card, 1990; Peri and Yasenov, 2018) that find no effect on workers without a college degree.

While the interpretation above makes sense relative to what we would expect to happen after a positive shock to labour supply of low-skilled workers in a competitive labour market model, note that the sample sizes are very small. We observe just 25 low-skilled and 16 high-skilled workers. At most, what we can say is that for the selected sample we observe a negative effect on wages of low-skilled workers. If PSID is to be used for future evaluations of the Mariel Boatlift, the researchers should consider how to increase sample size, e.g. include women.

Lastly, the confidence intervals in Figure 2 could in reality be wider. The reason is that the treatment effects do not incorporate the uncertainty connected with the weights of the synthetic control. In general, when doing causal inference with any SC method with many treated units, there are (at least) two sources of uncertainty: from the weights of the synthetic controls and from the estimates of the individual treatment effects τ_i . The confidence intervals above only account for the second source of uncertainty and not for the uncertainty connected to the weights themselves. So, overall inference is likely to be unreliable in this case. At present, accounting for both sources of uncertainty is an open question.



Figure 2: 95% CI for Mariel Boatlift's effect on low-skilled and high-skilled workers.

7 Discussion, Limitations and Conclusions

The main purpose of this paper was to develop an estimator from the family of Synthetic Control for the case of many treated units. The estimator we propose Correlated Synthetic Control (CSC) creates synthetic doppelgangers that are correlated across treated individuals similar in terms of observables. We show that CSC has good theoretical properties when treatment assignment is correlated with unobservable characteristics. In such cases, it even dominates the difference-in-difference estimator which is the most widely used technique when researchers are faced with a panel containing many treated units. We confirmed this key result in our simulation exercise: making treatment correlated with unobservables has little effect for estimates of treatment effects from CSC but estimates from difference-in-difference become much less precise. Moreover, we explored two applications where CSC can be preferred over other causal inference techniques: the effect of the German division on population growth in German cities and the effect of immigration on Miami natives' wages and labour supply.

With that said, CSC currently has some important limitations. At present, it only allows for time-invariant discrete characteristics. This key limitation stems partly from the fact that we have constrained the (correlated) random coefficients model of the weights to allow only for time-invariant predictors. What happens if we relax this restriction? Theoretically, the optimisation problem in Proposition 1 that we solve to find the weights can accommodate time-varying discrete covariates in the weights (e.g. changing occupation of treated workers).⁴² This is a fascinating possibility, as it will allow the weights on donors to be evolving over time, according to the time-varying covariates of a given treated individual. For instance, if an individual loses her job, we can construct different synthetic controls for her before and after this event.

On the other hand, the conventional (*SC Constraints*) that the weights sum up to 1 and are non-negative is definitely useful in macroeconomic interpretations where we can interpret the weights as probabilities. However, our results from the simulation showed good support for the framework introduced in Section 3 that turns the causal inference problem into a prediction problem. So, perhaps, if prediction is the end game, this restriction can be relaxed. Maybe we can impose other constraints that also find sparse solutions but do not restrict the set of weights as much such as Lasso.

Further limitations include the fact that we do not use *at all* donors' covariates. In particular, when specifying the correlated random coefficients structure of the weights, we only used treated individuals' covariates. This is not entirely satisfactory, as we are losing an important source of information. In any case, there are many open questions related to CSC and other SC estimators that this paper could not address. However, as noted by Guido Imbens in the Sargan lecture last month, SC-like methods are growing into an important part of the toolkit of the empirical researchers. Therefore, pursuing such open questions further is a fruitful area of research.

⁴²See restriction (9). With time-varying covariates we still need for it to hold. However, it does not preclude the possibility for x_{it}^k to be time-varying.

References

- A. Abadie. Using synthetic controls: Feasibility, data requirements, and methodological aspects. *Forthcoming in Journal of Economic Literature*, 2020a. URL <https://economics.mit.edu/files/17847>.
- A. Abadie. Discussion session. In *Chamberlain Seminar*. Symposium on Synthetic Control, 2020b. URL https://www.youtube.com/watch?v=Wta8N6FaM4Y&t=5205s&ab_channel=ChamberlainSeminar.
- A. Abadie and J. L’Hour. A penalized synthetic control estimator for disaggregated data. *MIT Working Paper*, 2020. URL <http://economics.mit.edu/files/18642>.
- A. Abadie, A. Diamond, and J. Hainmueller. Synthetic control methods for comparative case studies: Estimating the effect of california’s tobacco control program. *Journal of the American Statistical Association*, 105(490):493–505, 2010.
- A. Abadie, A. Diamond, and J. Hainmueller. Comparative politics and the synthetic control method. *American Journal of Political Science*, 59(2):495–510, 2015. doi:<https://doi.org/10.1111/ajps.12116>. URL <https://onlinelibrary.wiley.com/doi/abs/10.1111/ajps.12116>.
- D. Acemoglu and D. Autor. Skills, tasks and technologies: Implications for employment and earnings. In D. Card and O. Ashenfelter, editors, *Handbook of Labor Economics*, volume 4b, page 1043–1171. Elsevier, 2011.
- J. D. Angrist and J.-S. Pischke. *Mostly harmless econometrics: An empiricist’s companion*. Princeton university press, 2008.
- D. Arkhangelsky, S. Athey, D. A. Hirshberg, G. W. Imbens, and S. Wager. Synthetic difference in differences. *ArXiv preprint No. 1812.09970*, 2020. URL <https://arxiv.org/abs/1812.09970>.
- S. Athey, M. Bayati, G. Imbens, and Z. Qu. Ensemble methods for causal effects in panel data settings. *AEA Papers and Proceedings*, 109:65–70, 2019. URL <https://arxiv.org/abs/1903.10079>.

- J. Bai. Panel data models with interactive fixed effects. *Econometrica*, 77 (4):1229–1279, 2009. doi:<https://doi.org/10.3982/ECTA6135>. URL <https://onlinelibrary.wiley.com/doi/abs/10.3982/ECTA6135>.
- P. Balestra and J. Krishnakumar. Fixed effects models and fixed coefficients models. In L. Mátyás and P. Sevestre, editors, *The Econometrics of Panel Data*, chapter 2. Springer, 3rd edition, 2008.
- E. Ben-Michael, A. Feller, and J. Rothstein. The augmented synthetic control method. *ArXiv preprint No. 1811.04170*, 2020. URL <https://arxiv.org/abs/1811.04170>.
- E. Ben-Michael, A. Feller, and J. Rothstein. Synthetic controls with staggered adoption. *ArXiv preprint No. 1912.03290*, 2021. URL <https://arxiv.org/abs/1912.03290>.
- D. S. Bernstein. *Matrix Mathematics*. Princeton University Press, 2009. doi:[doi:10.1515/9781400833344](https://doi.org/10.1515/9781400833344). URL <https://doi.org/10.1515/9781400833344>.
- G. J. Borjas. The wage impact of the marielitos: A reappraisal. *ILR Review*, 70 (5):1077–1110, 2017. doi:[10.1177/0019793917692945](https://doi.org/10.1177/0019793917692945). URL <https://doi.org/10.1177/0019793917692945>.
- B. Born, G. J. Müller, M. Schularick, and P. Sedláček. The Costs of Economic Nationalism: Evidence from the Brexit Experiment. Technical report, Centre for Macroeconomics Working paper, 2017. URL <http://www.centreformacroeconomics.ac.uk/Discussion-Papers/2017/CFMDP2017-38-Paper.pdf>.
- B. Born, G. J. Müller, M. Schularick, and P. Sedláček. The Costs of Economic Nationalism: Evidence from the Brexit Experiment. *The Economic Journal*, 129(623):2722–2744, 05 2019. doi:[10.1093/ej/uez020](https://doi.org/10.1093/ej/uez020). URL <https://doi.org/10.1093/ej/uez020>.
- I. Botosaru and B. Ferman. On the role of covariates in the synthetic control method. *The Econometrics Journal*, 22(2):117–130, 01 2019. doi:[10.1093/ectj/utz001](https://doi.org/10.1093/ectj/utz001). URL <https://doi.org/10.1093/ectj/utz001>.

- D. Card. The impact of the mariel boatlift on the miami labor market. *ILR Review*, 43(2):245–257, 1990. doi:[10.1177/001979399004300205](https://doi.org/10.1177/001979399004300205). URL <https://doi.org/10.1177/001979399004300205>.
- D. Card and D. R. Hyslop. Female earnings inequality: The changing role of family characteristics on the extensive and intensive margins. Technical report, NBER Working Paper No. 25387, 2018. URL https://www.nber.org/system/files/working_papers/w25387/w25387.pdf.
- W. Carlin and D. W. Soskice. *Macroeconomics: Institutions, instability, and the financial system*. Oxford University Press, USA, 2014.
- G. Chamberlain. Multivariate regression models for panel data. *Journal of econometrics*, 18(1):5–46, 1982.
- V. Chernozhukov, K. Wuthrich, and Y. Zhu. An exact and robust conformal inference method for counterfactual and synthetic controls. *ArXiv preprint No. 1712.09089*, 2020. URL <https://arxiv.org/abs/1712.09089>.
- M. A. Clemens and J. Hunt. The labor market effects of refugee waves: Reconciling conflicting results. *ILR Review*, 72(4):818–857, 2019. doi:[10.1177/0019793918824597](https://doi.org/10.1177/0019793918824597). URL <https://doi.org/10.1177/0019793918824597>.
- B. Crépon and J. Mairesse. The chamberlain approach to panel data: An overview and some simulations. In L. Mátyás and P. Sevestre, editors, *The Econometrics of Panel Data*, chapter 5. Springer, 3rd edition, 2008.
- D. R. Davis and D. E. Weinstein. Bones, bombs, and break points: the geography of economic activity. *American economic review*, 92(5):1269–1289, 2002.
- C. De Chaisemartin and X. d’Haultfoeuille. Two-way fixed effects estimators with heterogeneous treatment effects. *American Economic Review*, 110(9):2964–96, 2020.
- N. Doudchenko and G. Imbens. Balancing, regression, difference-in-differences and synthetic control methods: A synthesis. *ArXiv preprint No. 1610.07748*, 2018. URL <https://arxiv.org/abs/1610.07748>.

- C. Dustmann, U. Schönberg, and J. Stuhler. The impact of immigration: Why do studies reach such different results? *Journal of Economic Perspectives*, 30(4): 31–56, November 2016. doi:[10.1257/jep.30.4.31](https://doi.org/10.1257/jep.30.4.31). URL <https://www.aeaweb.org/articles?id=10.1257/jep.30.4.31>.
- B. Ferman. On the properties of the synthetic control estimator with many periods and many controls. *ArXiv preprint No. 1906.06665*, 2020. URL <https://arxiv.org/pdf/1906.06665.pdf>.
- B. Ferman and C. Pinto. Synthetic controls with imperfect pre-treatment fit. *ArXiv preprint No. 1911.08521*, 2021. URL <https://arxiv.org/pdf/1911.08521.pdf>.
- B. Ferman, C. Pinto, and V. Possebom. Cherry picking with synthetic controls. *Journal of Policy Analysis and Management*, 39(2):510–532, 2020. doi:<https://doi.org/10.1002/pam.22206>. URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/pam.22206>.
- A. Fu, B. Narasimhan, and S. Boyd. Cvxr: An r package for disciplined convex optimization. *Journal of Statistical Software*, 2019. URL https://web.stanford.edu/~boyd/papers/cvxr_paper.html.
- D. L. Hanson and F. T. Wright. A Bound on Tail Probabilities for Quadratic Forms in Independent Random Variables. *The Annals of Mathematical Statistics*, 42(3):1079 – 1083, 1971. doi:[10.1214/aoms/1177693335](https://doi.org/10.1214/aoms/1177693335). URL <https://doi.org/10.1214/aoms/1177693335>.
- M. Hlavac. **Stargazer**: beautiful latex, html and ascii tables from r statistical output. Technical report, CRAN Repository, 2018. URL <https://cran.r-project.org/web/packages/stargazer/vignettes/stargazer.pdf>.
- C. Hsiao. Panel models with interactive effects. *Journal of Econometrics*, 206(2):645–673, 2018. doi:<https://doi.org/10.1016/j.jeconom.2018.06.017>. URL <https://www.sciencedirect.com/science/article/pii/S0304407618301131>. Special issue on Advances in Econometric Theory: Essays in honor of Takeshi Amemiya.

- C. Hsiao and M. H. Peseran. Random coefficients models. In L. Mátyás and P. Sevestre, editors, *The Econometrics of Panel Data*, chapter 6. Springer, 3rd edition, 2008.
- D. Hsu, S. M. Kakade, and T. Zhang. A tail inequality for quadratic forms of sub-gaussian random vectors. *ArXiv preprint No. 1110.2842v1*, 2011. URL <https://arxiv.org/abs/1110.2842>.
- S. M. Iacus, G. King, and G. Porro. Causal inference without balance checking: Coarsened exact matching. *Political analysis*, pages 1–24, 2012.
- A. Kaul, S. Klöbner, G. Pfeifer, and M. Schieler. Synthetic control methods: never use all pre-intervention outcomes together with covariates. *Saarland Working Paper*, 2018. URL https://www.gregor-pfeifer.net/files/SCM_Predictors_MC.pdf.
- P. Krugman. Increasing returns and economic geography. *Journal of political economy*, 99(3):483–499, 1991.
- A. Manning. A generalised model of monopsony*. *The Economic Journal*, 116(508):84–100, 2006. doi:<https://doi.org/10.1111/j.1468-0297.2006.01048.x>. URL <https://onlinelibrary.wiley.com/doi/abs/10.1111/j.1468-0297.2006.01048.x>.
- C. McCoy and D. Gonzalez. Cuban immigration and immigrants in florida and the united states: Implications for immigration policy. Technical report, University of Florida, Gainesville: Bureau of Economic and Business, 1985. referenced in Skop (2001).
- H. R. Moon and M. Weidner. Linear regression for panel with unknown number of factors as interactive fixed effects. *Econometrica*, 83(4):1543–1579, 2015. doi:<https://doi.org/10.3982/ECTA9382>.
- F. Oswald. **psidR**: Build panel data sets from psid raw data. Technical report, CRAN Repository, 2020. URL <https://cran.r-project.org/web/packages/psidR/psidR.pdf>.
- G. Peri and V. Yassenov. The labor market effects of a refugee wave. *Journal of*

- Human Resources*, 54(2):267–309, Jan. 2018. doi:[10.3368/jhr.54.2.0217.8561r1](https://doi.org/10.3368/jhr.54.2.0217.8561r1). URL <https://doi.org/10.3368/jhr.54.2.0217.8561r1>.
- S. J. Redding and D. M. Sturm. The costs of remoteness: Evidence from german division and reunification. *American Economic Review*, 98(5):1766–97, 2008.
- P. Rigollet. MIT OpenCourseWare 18.S997: High Dimensional Statistics Chapter 2, 2015. URL https://ocw.mit.edu/courses/mathematics/18-s997-high-dimensional-statistics-spring-2015/lecture-notes/MIT18_S997S15_CourseNotes.pdf.
- E. H. Skop. Race and place in the adaptation of marinel exiles. *The International Migration Review*, 35(2):449–471, 2001. URL <http://www.jstor.org/stable/2675876>.
- T. Suri. Selection and comparative advantage in technology adoption. *Econometrica*, 79(1):159–209, 2011. doi:<https://doi.org/10.3982/ECTA7749>. URL <https://onlinelibrary.wiley.com/doi/abs/10.3982/ECTA7749>.
- U. US Department of Commerce. 1980 census of population: Florida. Technical report, US Department of Commerce: Census Bureau, 1982. URL https://www2.census.gov/prod2/decennial/documents/1980a_flABCs1-01.pdf.
- J. Wooldridge. Two-way fixed effects, the two-way mundlak regression, and event study estimators. Technical report, Working paper, 2021. URL https://www.dropbox.com/sh/zj91darudf2fica/AADj_jaf5ZuS1muobgsnxS6Za?dl=0&preview=two_way_mundlak_20210405.pdf.
- J. M. Wooldridge. Further results on instrumental variables estimation of average treatment effects in the correlated random coefficient model. *Economics Letters*, 79(2):185–191, 2003. doi:[https://doi.org/10.1016/S0165-1765\(02\)00318-X](https://doi.org/10.1016/S0165-1765(02)00318-X). URL <https://www.sciencedirect.com/science/article/pii/S016517650200318X>.
- Y. Xu. Generalized synthetic control method: Causal inference with interactive fixed effects models. *Political Analysis*, 25(1):57–76, 2017. doi:[10.1017/pan.2016.2](https://doi.org/10.1017/pan.2016.2).

Appendix

A Supplementary results and examples

A.a Matching covariates versus matching outcomes

As discussed in the main text, the model specified by SC methods for the weights:

$$y_{it} = \sum_{i=1}^{n_0} \hat{w}_i y_{it} + e_{it}$$

looks similar to a linear regression model, except for the constraint that the weights sum up to 1 and are greater than 0. Based on this intuition, there have emerged two contrasting approaches in the literature emphasising the matching of synthetic controls on different things: 1) matching primarily on *covariates* advocated by Alberto Abadie and coauthors (Abadie et al., 2010; Abadie and L’Hour, 2020; Abadie, 2020a) versus 2) matching primarily on *outcomes* advocated by Athey et al. (2019); Doudchenko and Imbens (2018); Ben-Michael et al. (2020). The main reason why this distinction matters is that it leaves open the question of how to use covariates in synthetic control, as illustrated by (Botosaru and Ferman, 2019).

To understand the differences, suppose that $n_1 = 1$. It is useful to consider a linear transformation of the stacked row vector of outcomes and covariates $(\mathbf{y}_{n_0+1}^{pre}, \mathbf{x}_{n_0+1})'$ via the $(H \times (T_0 + k))$ matrix \mathbf{M} , given by the multiplication $\mathbf{M}(\mathbf{y}_{n_0+1}^{pre}, \mathbf{x}_{n_0+1})'$. The two approaches differ in the choice of \mathbf{M} and its dimension H . As common in the literature (e.g. Ferman and Pinto, 2021), suppose that Δ denotes the n_0 -dimensional standard simplex $\Delta \equiv \{(w_1, \dots, w_{n_0})' \in \mathbb{R}^{n_0} \mid \sum_{j=1}^{n_0} w_j = 1 \text{ and } w_j \geq 0\}$. Simplifying slightly,⁴³ both approaches can be reduced to solving a version of the constrained

⁴³We are simplifying because in practice Alberto Abadie (2020a) advocates that the weights \mathbf{V} and \mathbf{W} should be estimated in two optimisation problems.

quadratic optimisation problem in matrix form:

$$\min_{\mathbf{w} \in \Delta} \left(\mathbf{M} \begin{pmatrix} \mathbf{y}_{n_0+1}^{pre} \\ \mathbf{x}_{n_0+1} \end{pmatrix} - \mathbf{M} \begin{pmatrix} (\mathbf{Y}_{n_0}^{pre})' \\ \mathbf{X}_{n_0} \end{pmatrix} \mathbf{w} \right)' \mathbf{V} \left(\mathbf{M} \begin{pmatrix} \mathbf{y}_{n_0+1}^{pre} \\ \mathbf{x}_{n_0+1} \end{pmatrix} - \mathbf{M} \begin{pmatrix} (\mathbf{Y}_{n_0}^{pre})' \\ \mathbf{X}_{n_0} \end{pmatrix} \mathbf{w} \right) \quad (20)$$

where $\mathbf{w} = (w_1, \dots, w_{n_0})'$ is the $(n_0 \times 1)$ vector of weights and \mathbf{X}_{n_0} is the $n_0 \times K$ matrix of time-invariant covariates for the donors.⁴⁴ The $(H \times H)$ matrix \mathbf{V} is a diagonal matrix with different *predictors'* weights (that are different from *donors'* weights \mathbf{w}) on the linear combination of outcomes and covariates that we use.

Matching on outcomes approaches consider predicting covariates a secondary concern and instead aims at imputing the missing elements of Θ . As a result, the matrix \mathbf{M} is chosen so that the covariates are excluded and we solve the problem:

$$\min_{\mathbf{w} \in \Delta} (\mathbf{y}_{n_0+1}^{pre} - \mathbf{Y}_{n_0}^{pre} \mathbf{w})' (\mathbf{y}_{n_0+1}^{pre} - \mathbf{Y}_{n_0}^{pre} \mathbf{w}) \quad (21)$$

For this approach, imputing the matrix Θ is the most important factor when constructing synthetic controls. If our synthetic controls match well the time-series of the outcome in the pre-treatment period, then we should not be too worried about balance on covariates. In this case, \mathbf{V} is often chosen to be the identity matrix. We discuss further the exact choice of \mathbf{M} delivering the optimisation problem (21) and other details on the literature review in Appendix ??.

Instead of matching on the full set of outcomes separately $\mathbf{y}_{n_0+1}^{pre}$, *matching on covariates* takes an average of the pretreatment outcomes $\bar{y}_{n_0+1,pre} = \sum_{t=1}^{T_0} y_{n_0+1,t} / T_0$ and stack this average on top of the covariates. By choosing an appropriate \mathbf{M} ,⁴⁵ the general optimisation problem in (20) reduces to:

$$\min_{\mathbf{w} \in \Delta} \left(\begin{pmatrix} \bar{y}_{n_0+1}^{pre} \\ \mathbf{x}_{n_0+1} \end{pmatrix} - \begin{pmatrix} \bar{\mathbf{y}}_{n_0}^{pre} \\ \mathbf{X}_{n_0} \end{pmatrix} \mathbf{w} \right)' \mathbf{V} \left(\begin{pmatrix} \bar{y}_{n_0+1}^{pre} \\ \mathbf{x}_{n_0+1} \end{pmatrix} - \begin{pmatrix} \bar{\mathbf{y}}_{n_0}^{pre} \\ \mathbf{X}_{n_0} \end{pmatrix} \mathbf{w} \right) \quad (22)$$

⁴⁴We need to take the transpose of $\mathbf{Y}_{n_0}^{pre}$ to achieve consistency with the matrix Θ above

⁴⁵See Appendix ??

where $\bar{\mathbf{y}}_{n_0}^{pre}$ is a row vector of averages for the donors. The diagonal $(K + 1 \times K + 1)$ matrix \mathbf{V} can be chosen to ensure we have the same scale on all covariates and the average outcome or can optimise some other criteria, e.g. predictors are proportional to their importance in explaining the outcome variable (Abadie, 2020a).

So, should we matching outcomes or covariates? Kaul et al. (2018) warns that if we decide to ignore covariates, we may get misleading results in certain cases even though our synthetic control follows the true time series better in the pre-treatment period. Nevertheless, Botosaru and Ferman (2019) show that SC methods have good theoretical properties in cases when we fail to match on covariates. Furthermore, unless T_0 is very long, there seem to be ample possibilities for specification-searching of counterfactuals under both matching on covariates and outcomes (Ferman et al., 2020). Therefore, the place of covariates in synthetic control methods is ambiguous at present. As described below, we will take another approach to using the information from covariates for situations with many tread units. In particular, we shall use covariates to form synthetic controls that are correlated across individuals that are similar in terms of observable characteristics.

A.b Details on Estimation

Proposition 1 in the main body of the text is the first half of Proposition 5. The last proposition reformulates into a quadratic optimisation problem the Correlated Synthetic Controls (CSC) model:

$$\begin{aligned}
 y_{it}(0) &= \eta_i + \sum_{j=1}^{n_0} w_{ij} y_{jt}(0) + e_{it} \quad s.t. \\
 w_{ij} &= \underbrace{\omega_j}_{ind.-invariant} + \underbrace{\mathbf{x}_i \boldsymbol{\alpha}^j}_{ind.-specific} \quad (Random\ Coef.) \\
 \sum_{j=1}^{n_0} w_{ij} &= 1 \quad \forall j : w_{ij} \geq 0 \quad (SC\ Constraints)
 \end{aligned} \tag{23}$$

where we also allow for an intercept η_i , following suggestions in [Ferman and Pinto \(2021\)](#) and [Doudchenko and Imbens \(2018\)](#). Proposition 5 gives us two equivalent formulations of the same constrained quadratic optimisation problem from (23): in scalar form (24) and in matrix form (25). First, we can consider the scalar formulation. One intuition for estimating the model given by CSC in (23) is that we are essentially regressing the vector of outcomes for the treatment group on the outcomes for the donors and an interactions between donors' outcomes and treatment group's covariates, given some constraints on the coefficients. In (24) below, this can be seen as a regression of y_{it} for the treated group on donors' y_{jt} and the interaction between y_{jt} and treated group's covariates \mathbf{x}_i .

To understand Proposition 5, let us define \otimes to be the Kronecker product of two matrices and $\mathbf{1}_{n_1}$ to be a $(n_1 \times 1)$ column vector of 1-s. We also introduce the operator $Vec(\cdot)$ for vectorising some matrix \mathbf{A} . In particular, $Vec(\mathbf{A})$ takes the columns of \mathbf{A} and stack them on top of each other.⁴⁶

Proposition 5. *The parameters $\boldsymbol{\alpha}$ and $\boldsymbol{\omega}$ can be estimated by solving the optimisation*

⁴⁶For example, $Vec\left(\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix}\right) = (1, 3, 2, 4)'$

problem:

$$\begin{aligned}
& \max_{\alpha_j^k, \omega_j} \sum_{i=1}^{n_1} \sum_{t=1}^{T_0} \left(y_{it} - \eta_i - \sum_{j=1}^{n_0} y_{jt} \omega_j - \sum_{j=1}^{n_0} \sum_{k=1}^K \alpha_j^k y_{jt} x_i^k \right)^2 \quad s.t. \\
& \forall i : \sum_{j=1}^{n_0} \left(\omega_j + \sum_{k=1}^K x_i^k \alpha_j^k \right) = 1 \quad \forall (i, j) : \omega_j + \sum_{k=1}^K x_i^k \alpha_j^k \geq 0
\end{aligned} \tag{24}$$

where \mathbf{x}_i and $\boldsymbol{\alpha}_j$ are $(1 \times K)$ row vectors. Equivalently, in matrix notation, we have:

$$\begin{aligned}
& \max_{\boldsymbol{\omega}, \boldsymbol{\alpha}} \text{Vec} \left((\mathbf{Y}_{n_1}^{pre})' - \boldsymbol{\iota}_{T_0} \otimes \boldsymbol{\eta}' - (\mathbf{Y}_{n_0}^{pre})' \boldsymbol{\omega} \otimes \boldsymbol{\iota}_{n_1}' - (\mathbf{Y}_{n_0}^{pre})' \boldsymbol{\alpha} \otimes \mathbf{X}_{n_1}' \right)' \\
& \text{Vec} \left((\mathbf{Y}_{n_1}^{pre})' - \boldsymbol{\iota}_{T_0} \otimes \boldsymbol{\eta}' - (\mathbf{Y}_{n_0}^{pre})' \boldsymbol{\omega} \otimes \boldsymbol{\iota}_{n_1}' - (\mathbf{Y}_{n_0}^{pre})' \boldsymbol{\alpha} \otimes \mathbf{X}_{n_1}' \right) \quad s.t. \\
& (\boldsymbol{\iota}_{n_1} \otimes \boldsymbol{\omega}' + \mathbf{X}_{n_1} \boldsymbol{\alpha}') \boldsymbol{\iota}_{n_0} = \boldsymbol{\iota}_{n_1} \quad \boldsymbol{\iota}_{n_1} \otimes \boldsymbol{\omega}' + \mathbf{X}_{n_1} \boldsymbol{\alpha}' \geq \mathbf{O}_{(n_1 \times n_0)}
\end{aligned} \tag{25}$$

where as discussed $\mathbf{Y}_{n_1}^{pre}$ is the observed $(n_1 \times T_0)$ matrix of pre-treatment outcomes for the treated group, $\boldsymbol{\eta}$ is a $(n_1 \times 1)$ column vector of intercepts, $\mathbf{Y}_{n_0}^{pre}$ is the observed $(n_0 \times T_0)$ matrix of pre-treatment outcomes for the donors, $\boldsymbol{\omega}$ is the $(n_0 \times 1)$ vector of individual invariant weights, $\boldsymbol{\alpha}$ is $(n_0 \times K)$ matrix of coefficients and $\mathbf{O}_{(n_1 \times n_0)}$ is a $(n_1 \times n_0)$ matrix of zeros.

Proof. All proofs are relegated to Appendix B. See Appendix B.a for the proof of this Proposition. \square

The main intuition for the proof is that we are stacking (23) horizontally over i and vertically over t . This allows us to obtain the matrix inside the $\text{Vec}(\cdot)$ operator. Note also that even though the constraints in the two formulations look different at first, they represent the same expressions.

Let us now explore further the matrix representation in (25). While (25) may look daunting at first given these definitions, it is extremely useful when we would like to implement the estimator in practice via some optimisation software. The reason is that it is written in matrices and vectors that are directly observed in the data. For example, $\mathbf{Y}_{n_1}^{pre}$ and $\mathbf{Y}_{n_0}^{pre}$ can be directly taken from $\boldsymbol{\Theta}$ - the matrix of outcomes. As a result, we can apply some relatively simply transformations on the data such as finding transposes and vectorising matrices, in order to achieve the desired formulation.

Lastly, regarding the implementation of CSC, we code the optimisation problem (25) in R using the package CVXR. The key advantage of CVXR relative to other quadratic programming solvers is that it allows us to formulate the problem in a natural mathematical language rather than the restrictive formulations that are required by other solvers (Fu et al., 2019). Moreover, since CVXR is a wrapper around other solvers, we can directly compare the performance of different solvers in terms of how accurately and how quickly they find the weights.

A.c Example of multiplicity of solutions

Example 2 (Multiplicity of solutions). *Suppose that we have the dataset in Table 5 and we would like to construct a synthetic control for Tobi. Normal synthetic control that matches on outcomes only will be infeasible, as we have two weighted averages that match perfectly Tobi’s pretreatment wage in 1978 and 1979 and his education:*

$$\begin{aligned} Tobi &= \frac{1}{2}Max + \frac{1}{2}Micol \\ Tobi &= \frac{2}{3}Dirk + \frac{1}{3}Yi \end{aligned}$$

Moreover, for any $\lambda \in [0, 1]$ we have that $\lambda(\frac{1}{2}Max + \frac{1}{2}Micol) + (1-\lambda)(\frac{2}{3}Dirk + \frac{1}{3}Yi)$ will also work. Thus, we have an infinite number of solutions. The question then become how to choose one synthetic control among the infinite number of potential synthetic control. [Abadie and L’Hour \(2020\)](#) suggest selecting the synthetic control with the most similar outcomes and covariates. In this case, this will be Max and Micol, as the difference between their outcomes and covariates and Tobi’s outcomes and covariates is smaller in absolute terms than the same difference but for Dirk and Yi. Go back to [Section 3.d](#) by clicking on the number.

Table 5: Pseudo Individual-level sample for the Mariel Boatlift

Person	City	wage 1978	wage 1979	education
Tobi	Miami	12	15	11
Mobarak	Miami	21	23	12
...	
Max	Nashville	11	13	12
Dirk	Nashville	9	10	9
...	
Micol	Chicago	13	17	10
Yi	Chicago	18	30	15

A.d City-level vs Pooled SC

This section is aimed at giving some intuition towards why the pooled synthetic control nests the city-level synthetic control as a special case. We can prove the following proposition:

Proposition 6. *Suppose that:*

- (i) *We have a panel which is balanced across cities, e.g. for every treated and untreated city we observe the same number of individuals n .*
- (ii) *We have no covariates, so that we only observe Θ and city membership*
- (iii) *We have a total of C cities, of which the first c_0 cities are untreated and the rest c_1 cities are treated*

Then, if we impose the additional requirement $w_{dj} = \frac{w_d}{n}$, the pooled SC estimator and the city-level SC estimator solve the same optimisation problem.

Proof. The pooled SC solves the problem:

$$\max_{w_{jc}} \sum_{c=c_1+1}^C \sum_{i=1}^n \sum_{t=1}^T \left(y_{cit} - \sum_{d=1}^{c_0} \sum_{j=1}^n w_{dj} y_{djt} \right)^2 \quad s.t. \quad \sum_{d=1}^{c_0} \sum_{j=1}^n w_{dj} = 1 \quad w_{dj} \geq 0$$

Let us ignore the constraints and impose the additional requirement that every individual within a city gets the same weight $w_{dj} = \frac{w_d}{n}$. Then, we can rewrite the objective function as:

$$\min_{\mathbf{w}} \sum_{c=c_1+1}^C \sum_{i=1}^n \sum_{t=1}^T \underbrace{\left(y_{cit} - \sum_{d=1}^{c_0} w_d \bar{y}_{dt} \right)^2}_{h_t(w)}$$

where we use $\bar{y}_{dt} = \frac{\sum_{j=1}^n y_{djt}}{n}$ and $h_t(w) = \sum_{d=1}^{c_0} w_d \bar{y}_{dt}$. If we open the squares and drop the terms that do not involve \mathbf{w} , we obtain:

$$\min_{\mathbf{w}} \sum_{t=1}^T h_t(w) \left(Ch_t(w) - 2 \sum_{c=c_1+1}^C \bar{y}_{ct} \right) \quad (26)$$

where we use the same definition of \bar{y}_{ct}

In contrast, the city-level SC solves:

$$\min_{\mathbf{w}} \sum_{c=c_1+1}^C \sum_{t=1}^T (\bar{y}_{ct} - \sum_{d=1}^{c_0} w_d \bar{y}_{dt})^2$$

Once we upon the square and drop the terms that are irrelevant for optimisation, we can rewrite this as:

$$\min_{\mathbf{w}} \sum_{t=1}^T h_t(w) \left(Ch_t(w) - 2 \sum_{c=c_1+1}^C \bar{y}_{ct} \right) \quad (27)$$

which is the same as the (26), so that the objective function is the same in the two cases.⁴⁷ □

While the balance of the panel across cities may seem like a restrictive condition, we should be able to relax it by imposing in the pooled synthetic control that each city has an equal weight overall, even though different cities contain a different number of observations. For instance, if we have many individuals within two cities Miami and Atlanta with Miami having 100 individuals and Atlanta just 50, then we can impose that each individual in Miami gets a weight $\frac{50}{150}$ and each individual in Atlanta gets a weight of $\frac{100}{150}$. For reasons of time, however, I was unable to show this more general case more formally. Nevertheless, the main result still stands: the pooled synthetic control nests the city-level synthetic control.

⁴⁷This is related to the regression equivariance property of the Least Squares Objective function.

A.e DiD Inconsistency

Here we sketch another argument in addition to Proposition 13 for showing that $\hat{\tau}^{did}$ will be inconsistent in light of the discussion in Section 5.a. Note that estimating $y_{it} = \rho + \gamma_i + \delta_t + \tau D_{it} + u_{it}$ will yield an inconsistent estimate of the true τ , since $Cov(D_{it}, u_{it}) \neq 0$, meaning that D_{it} is endogenous. Naturally, in the spirit of the Arellano-Bond estimator, the question arises if we can apply some form of transformation to the data (e.g. first differences) and instrument D_{it} with some of its lagged values D_{is} in order to deal with the endogeneity issue. Unfortunately, this will not work, due to the interactive fixed effects structure. Note that we also have a further problem with consistency: $Cov(\gamma_i, u_{it}) \neq 0$ as factor loadings μ_i enter into both (reduced-form) terms, complicating further consistency of $\hat{\tau}^{did}$ in (16).

A.f Fixed parameters for simulation

We need to fix the parameters for three matrices: the common factors $\boldsymbol{\lambda}$, the partial correlation between unobservables and covariates $\boldsymbol{\gamma}$ and the partial correlation between unobservables and treatment assignment $\boldsymbol{\phi}$.

Firstly, every value in the $(T \times F)$ matrix of common factor $\boldsymbol{\lambda}$ was independently drawn once from $N(3, 2)$. The resulting values are in Table 28. The rows are the time periods and the columns are the factors.

$$\begin{pmatrix} 1.79 & 2.44 & 2.49 & 2.31 \\ 3.27 & 2.11 & 2.09 & 1.55 \\ 4.08 & 2.52 & 2.16 & 3.58 \\ 0.65 & 2.00 & 5.41 & 1.98 \\ 3.43 & 2.22 & 3.13 & 2.99 \\ 3.51 & 3.06 & 2.51 & 2.06 \\ 2.43 & 3.96 & 2.56 & 4.10 \\ 2.45 & 2.89 & 3.46 & 2.52 \end{pmatrix} \quad (28)$$

Secondly, we give the fixed values of the $K \times F$ matrix $\boldsymbol{\gamma}$ in (29). These are the values, controlling the (partial) correlation between the time-invariant covariates \boldsymbol{x}_i and the unobserved factor loadings $\boldsymbol{\mu}$ from (17):

$$\boldsymbol{\mu}_i = \boldsymbol{\gamma}' \boldsymbol{x}_i + \boldsymbol{v}_i$$

We draw each member of $\boldsymbol{\gamma}$ independently from a standard normal $N(1, 0)$ distribution. So, we have the matrix in Table 29.

$$\begin{pmatrix} -1.48 & -0.32 & -0.78 & 0.51 \\ 1.58 & -0.63 & 0.01 & -0.29 \\ -0.96 & -0.11 & -0.15 & 0.22 \\ -0.92 & 0.43 & -0.70 & 2.01 \\ -2.00 & -0.78 & 1.19 & 1.01 \\ -0.27 & -1.29 & 0.34 & -0.30 \end{pmatrix} \quad (29)$$

Lastly, we have for the parameter ϕ that determined the degree of correlation between the factor loadings and treatment assignment (See [18](#)). In particular, we set the parameter as

$$\phi = (-1.12, -0.46, 3.12, 0.14)'$$

Note that we hold it constant across all of our simulations, even though generated it initially from a $N(0, 1)$ distribution.

A.g Further Simulation Results

Table 6 gives the results from the same simulation in **Algorithm 1** that we used in the main body of the test. The general features of the results are the same, as discussed in Section 5. One interesting feature of the results, however, is that once we take account of the variance we can clearly see that iDiD does a better job at predicting the ATT relative to the case where we study the pure estimation error.

Table 6: *RMSE between $\hat{\tau}$ and true τ under different DGPs across 1000 simulations*

Change	CSC	fDiD	PSC	iDiD
Baseline	1.91	2.40	2.68	0.98
Less Correlation	1.87	2.35	2.56	0.99
Random Assignment	1.88	2.15	2.51	1.00
$N = 150$	1.88	2.37	2.45	0.98
$N = 200$	1.88	2.37	2.37	0.97
$T = 4$	5.70	6.35	10.29	0.94
$T = 8$	1.42	1.40	2.28	1.01
$F = 2$	1.98	3.08	3.05	0.99
$F = 4$	2.06	2.55	3.15	0.99
$E[\pi_i] = 0.25$	1.92	2.35	2.76	0.98
$E[\pi_i] = 0.40$	1.97	2.30	2.86	0.97

A.h Variables from PSID

A.h.1 Sample restrictions

The PSID sample that we use consists of 11 waves (1974-1984) and is restricted to male heads of households. Until recently, the heads of households in PSID were considered to be the husband in a traditional heterosexual marriage. This formulation has been controversial and was justifiably described as "anachronistic" (Card and Hyslop, 2018) and so it has been changed in recent years but not for the time period studied here. As a result, our sample is restricted to males of working age (19-64), for whom we have data on all 11 years.

Excluding women can be justified in two ways: i) the data for them in PSID is very sparse for the period we studied and ii) controlling for the general increase in women's participation in the labour force at the time is challenging given that we work with time-invariant covariates (Borjas, 2017). However, further work on the Mariel Boatlift with PSID should strive to include women in the analysis.

A.h.2 Variable Selection

Turning to the variables selected in the sample, the main outcomes of interest are total hours worked per year as an indicator of labour supply and hourly (log) wages. While previous studies have found no effect on aggregate unemployment, Dustmann et al. (2016) has recently stated the importance of relaxing the assumption of perfectly inelastic labour supply and studying labour supply effect in response to immigration. This would amount to making LS in Figure 1 downward sloping rather than a vertical line. In addition, we also include various covariates on occupation, industry, health, marriage status, age and education.

Table 7 below presents the variables that we include in the analysis of the empirical application the MPhil thesis. It also discusses any recoding that we have used to ensure that the variables can be meaningfully included in the analysis. For example, days being ill has been recoded to a binary variable, indicating being one of the top 25% people who spent the longest period of time being ill. More importantly, we

turn the covariates below into time-invariant variables by taking their average over the pre-treatment period. In addition, we give the name of the variable only in 1975. The reason is that variable names differ across years but from PSID's Variable Search engine one can identify which are the comparable variables across years.

Next, it should be remarked that the original sample we used was from 1973 to 1985. However, we need to impose restrictions, due to the fact that one of the outcomes (hours worked) was retrospective for the past year and the other outcome (wages) was recorded for the week of interview. Lastly, since PSID's website is a bit hard to navigate, we used R package `psidR` to download the data ([Oswlad, 2020](#)).

Table 7: Variables used in the analysis for the empirical application

1975 Var.	Interpretation	Use	Recoding and Notes
V3823	Tot. Hours Worked	Outcome	Hours worked in <i>previous</i> year
V4093	Hourly Wage	Outcome	Hourly wage in <i>current</i> year
V3803	State Residence	Treatment	Used to find who's in Florida
V4204	Race	Covariate	Coded as white vs non-white
V4093	Education	Covariate	Coded as college vs no college
V4194	Married	Covariate	Married or not
V3969A	Head Industry	Covariate	Coded as dummy variables
V3968A	Head Occupation	Covariate	Coded as dummy variables
V3921	Age	Covariate	Used to restrict sample
V3825	Hours Ill	Covariate	Coded as binary variable

A.i HTT for other Specifications

As discussed in Section 6.d, we also calculated heterogeneous treatment effects for other lengths of the training period. The overall impression remains, however, that the only variable for which we consistently find a significant negative effects is wages of low-skilled workers. For other combinations (e.g. labour supply for high-skilled workers), we cannot reject the null hypothesis of no effect. The results can be found in Figure 3, Figure 4 and Figure 5.

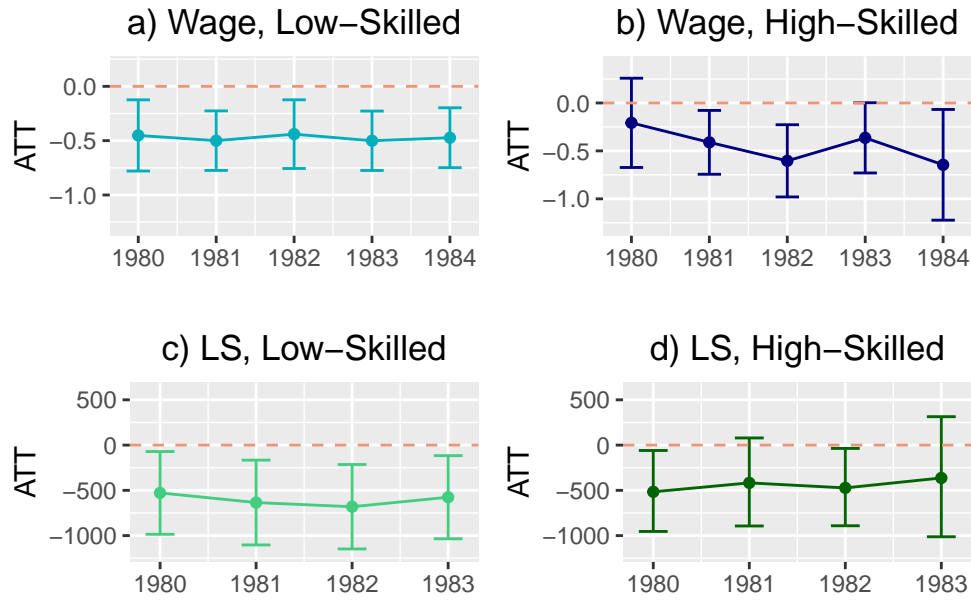


Figure 3: Treatment Effects when $T_{train} = 2$



Figure 4: Treatment Effects when $T_{train} = 3$

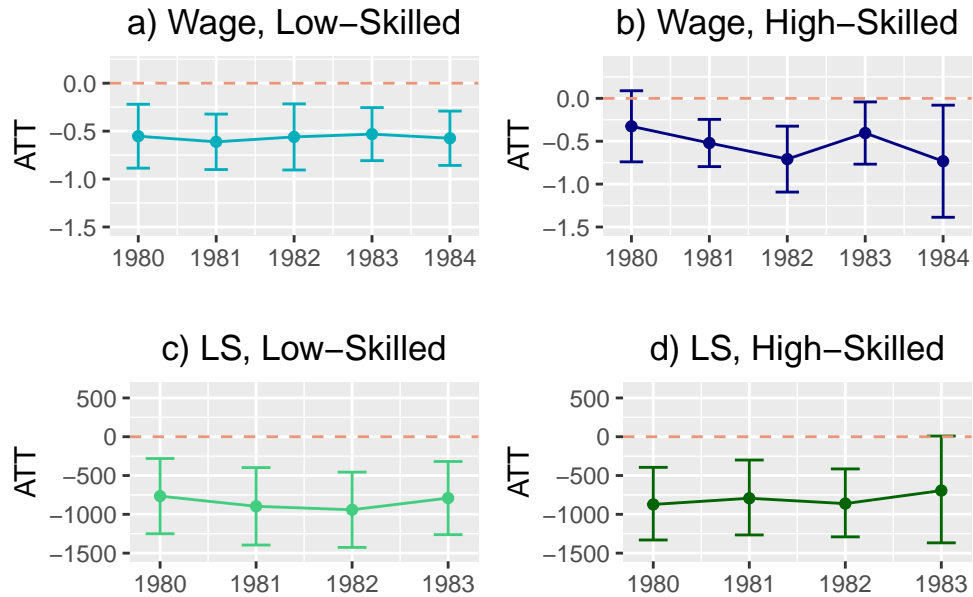


Figure 5: Treatment Effects when $T_{train} = 4$ for wages and $T_{train} = 1$ for LS

B Proofs of main results

B.a Proof of Proposition 5

Define \otimes to be the Kronecker product of two matrices and $\boldsymbol{\iota}_{n_1}$ to be a $(n_1 \times 1)$ column vector of 1-s. We also introduce the operator $Vec(\cdot)$ for vectorising some matrix \mathbf{A} . In particular, $Vec(\mathbf{A})$ takes the columns of \mathbf{A} and stack them on top of each other.⁴⁸

Proposition 5. *The parameters $\boldsymbol{\alpha}$ and $\boldsymbol{\omega}$ can be estimated by solving the optimisation problem:*

$$\begin{aligned} \max_{\boldsymbol{\alpha}_j^k, \boldsymbol{\omega}_j} \sum_{i=1}^{n_1} \sum_{t=1}^{T_0} \left(y_{it} - \eta_i - \sum_{j=1}^{n_0} y_{jt} \omega_j - \sum_{j=1}^{n_0} \sum_{k=1}^K \alpha_j^k y_{jt} x_i^k \right)^2 \quad s.t. \\ \forall i : \sum_{j=1}^{n_0} \left(\omega_j + \sum_{k=1}^K x_i^k \alpha_j^k \right) = 1 \quad \forall (i, j) : \omega_j + \sum_{k=1}^K x_i^k \alpha_j^k \geq 0 \end{aligned} \quad (24)$$

where \mathbf{x}_i and $\boldsymbol{\alpha}_j$ are $(1 \times K)$ row vectors. Equivalently, in matrix notation, we have:

$$\begin{aligned} \max_{\boldsymbol{\omega}, \boldsymbol{\alpha}} \quad & Vec \left((\mathbf{Y}_{n_1}^{pre})' - \boldsymbol{\iota}_{T_0} \otimes \boldsymbol{\eta}' - (\mathbf{Y}_{n_0}^{pre})' \boldsymbol{\omega} \otimes \boldsymbol{\iota}_{n_1}' - (\mathbf{Y}_{n_0}^{pre})' \boldsymbol{\alpha} \otimes \mathbf{X}_{n_1}' \right)' \\ & Vec \left((\mathbf{Y}_{n_1}^{pre})' - \boldsymbol{\iota}_{T_0} \otimes \boldsymbol{\eta}' - (\mathbf{Y}_{n_0}^{pre})' \boldsymbol{\omega} \otimes \boldsymbol{\iota}_{n_1}' - (\mathbf{Y}_{n_0}^{pre})' \boldsymbol{\alpha} \otimes \mathbf{X}_{n_1}' \right) \quad s.t. \\ & (\boldsymbol{\iota}_{n_1} \otimes \boldsymbol{\omega}' + \mathbf{X}_{n_1} \boldsymbol{\alpha}') \boldsymbol{\iota}_{n_0} = \boldsymbol{\iota}_{n_1} \quad \boldsymbol{\iota}_{n_1} \otimes \boldsymbol{\omega}' + \mathbf{X}_{n_1} \boldsymbol{\alpha}' \geq \mathbf{O}_{(n_1 \times n_0)} \end{aligned} \quad (25)$$

where as discussed $\mathbf{Y}_{n_1}^{pre}$ is the observed $(n_1 \times T_0)$ matrix of pre-treatment outcomes for the treated group, $\boldsymbol{\eta}$ is a $(n_1 \times 1)$ column vector of intercepts, $\mathbf{Y}_{n_0}^{pre}$ is the observed $(n_0 \times T_0)$ matrix of pre-treatment outcomes for the donors, $\boldsymbol{\omega}$ is the $(n_0 \times 1)$ vector of individual invariant weights, $\boldsymbol{\alpha}$ is $(n_0 \times K)$ matrix of coefficients and $\mathbf{O}_{(n_1 \times n_0)}$ is a $(n_1 \times n_0)$ matrix of zeros.

Proof. The strategy for the proof involves three steps, in which we are turning model (7) into a constrained quadratic optimisation problem by stacking equations over i and t .

⁴⁸For example, $Vec \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} = (1, 3, 2, 4)'$

Firstly, we substitute the constraint (*Random. Coef*) for some given i and t :

$$y_{it} = \eta_i + \sum_{j=1}^{n_0} \underbrace{\omega_j}_{ind.-invariant} y_{jt} + \sum_{j=1}^{n_0} \underbrace{\mathbf{x}_i \boldsymbol{\alpha}^j}_{ind.-specific} y_{jt} + \epsilon_{it} \quad s.t.$$

$$\sum_{j=1}^{n_0} \omega_j + \mathbf{x}_i \boldsymbol{\alpha}^j = 1 \quad \forall j : \omega_j + \mathbf{x}_i \boldsymbol{\alpha}^j \geq 0$$

We can rewrite this expression as the first constrained quadratic optimisation problem in the proposition:

$$\max_{\alpha_j^k, \omega_j} \sum_{i=1}^{n_1} \sum_{t=1}^{T_0} \left(y_{it} - \eta_i - \sum_{j=1}^{n_0} y_{jt} \omega_j - \sum_{j=1}^{n_0} \sum_{k=1}^K y_{jt} \alpha_j^k x_i^k \right)^2 \quad s.t.$$

$$\sum_{j=1}^{n_0} \left(\omega_j + \sum_{k=1}^K x_i^k \alpha_j^k \right) = 1 \quad \forall j : \omega_j + \sum_{k=1}^K x_i^k \alpha_j^k \geq 0$$

where every symbol represents a scalar.

Secondly, we will stack y_{it} over time periods t and treated units index i . Let us leave the constraints in their present form and focus on the expression for y_{it}^0 rewritten in vector form:

$$y_{it} = \eta_i + \mathbf{y}'_{t, n_0} \boldsymbol{\omega} + \mathbf{y}'_{t, n_0} \boldsymbol{\alpha} \mathbf{x}'_i + e_{it}$$

where \mathbf{y}_{t, n_0} is a $(n_0 \times 1)$ vector of outcomes for the donor pool at time t , $\boldsymbol{\omega}$ is $(n_0 \times 1)$ vector of ind.-invariant weights, \mathbf{x}_i a $(1 \times K)$ row vector, $\boldsymbol{\alpha}$ is a $(n_0 \times K)$ vector of coefficients and \mathbf{y}_{t, n_0} is a $(n_0 \times 1)$ vector of outcomes for donor pool. The next step is to stack over pre-treatment time period $t \leq T_0$ for some given individual i :

$$\underbrace{(\mathbf{y}_i^{pre})'}_{T_0 \times 1} = \underbrace{\boldsymbol{\nu}_{T_0}}_{T_0 \times 1} \eta_i + \underbrace{(\mathbf{Y}_{n_0}^{pre})'}_{T_0 \times n_0} \underbrace{\boldsymbol{\omega}}_{n_0 \times 1} + \underbrace{(\mathbf{Y}_{n_0}^{pre})'}_{T_0 \times n_0} \underbrace{\boldsymbol{\alpha}}_{n_0 \times K} \underbrace{\mathbf{x}_i}_{K \times 1} + \underbrace{\boldsymbol{\epsilon}_i}_{T_0 \times 1}$$

where \mathbf{y}_i^{pre} is a row $1 \times T_0$ vector of outcomes for the given treated unit i in the pre-treatment period, $\mathbf{Y}_{n_0}^{pre}$ is a $n_0 \times T_0$ matrix of pre-treatment outcomes for the donors and $\boldsymbol{\epsilon}_i$ is a $T_0 \times 1$ vector of errors. Here we can remark that $\mathbf{Y}_{n_0}^{pre}$ is the same matrix that appears in the full matrix of observed outcomes $\boldsymbol{\Theta}$ in (2). The matrix $\boldsymbol{\Theta}$ is the

reason we pick the specific dimensions of the stacked \mathbf{y} vectors. Next, we will combine the information over all treated individuals i to form the following mapping:

$$\underbrace{(\mathbf{Y}_{n1}^{pre})'}_{T_0 \times n_1} = \underbrace{\boldsymbol{\iota}_{T_0}}_{T_0 \times 1} \otimes \underbrace{\boldsymbol{\eta}'}_{1 \times n_1} + \underbrace{(\mathbf{Y}_{n0}^{pre})'\boldsymbol{\omega}}_{T_0 \times 1} \otimes \underbrace{\boldsymbol{\iota}'_{n_1}}_{1 \times n_1} + \underbrace{(\mathbf{Y}_{n0}^{pre})'}_{T_0 \times n_0} \underbrace{\boldsymbol{\alpha}}_{n_0 \times K} \otimes \underbrace{(\mathbf{X}_{n1})'}_{K \times n_1} + \underbrace{\boldsymbol{\epsilon}}_{T_0 \times n_1} \quad (30)$$

where \mathbf{Y}_{n1}^{pre} is the $(n_1 \times T_0)$ matrix of pre-treatment outcomes that can also be found in (2), \mathbf{X}_{n1} is a $(n_1 \times K)$ matrix of time-invariant covariates for all treated units and $\boldsymbol{\epsilon}$ is a $(T_0 \times n_0)$ matrix. Both sides of the last equation are actually $(T_0 \times n_1)$ matrices.

Thirdly, to turn (30) into a quadratic optimisation problem, we can introduce the operator $Vec(\cdot)$ for vectorising some matrix \mathbf{A} . In particular, $Vec(\mathbf{A})$ takes the columns of \mathbf{A} and stack them on top of each other, as discussed above. Thus, if we apply this to (30), we will be able to identify the parameters via solving the constrained quadratic optimisation problem:

$$\begin{aligned} \max_{\boldsymbol{\omega}, \boldsymbol{\alpha}} \quad & Vec\left((\mathbf{Y}_{n1}^{pre})' - \boldsymbol{\iota}_{T_0} \otimes \boldsymbol{\eta}' - (\mathbf{Y}_{n0}^{pre})'\boldsymbol{\omega} \otimes \boldsymbol{\iota}'_{n_1} - (\mathbf{Y}_{n0}^{pre})'\boldsymbol{\alpha} \otimes \mathbf{X}_{n1}'\right)' \\ & Vec\left((\mathbf{Y}_{n1}^{pre})' - \boldsymbol{\iota}_{T_0} \otimes \boldsymbol{\eta}' - (\mathbf{Y}_{n0}^{pre})'\boldsymbol{\omega} \otimes \boldsymbol{\iota}'_{n_1} - (\mathbf{Y}_{n0}^{pre})'\boldsymbol{\alpha} \otimes \mathbf{X}_{n1}'\right) \quad s.t. \\ & (\boldsymbol{\iota}_{n_1} \otimes \boldsymbol{\omega}' + \mathbf{X}_{n1} \boldsymbol{\alpha}') \boldsymbol{\iota}_{n_0} = \boldsymbol{\iota}_{n_1} \quad \boldsymbol{\iota}_{n_1} \otimes \boldsymbol{\omega}' + \mathbf{X}_{n1} \boldsymbol{\alpha}' \geq \mathbf{0} \end{aligned}$$

where as discussed \mathbf{Y}_{n1}^{pre} is the observed $(n_1 \times T_0)$ matrix of pre-treatment outcomes for the treated group, $\boldsymbol{\eta}$ is a $(n_1 \times 1)$ column vector of intercepts, \mathbf{Y}_{n1}^{pre} is the observed $(n_0 \times T_0)$ matrix of pre-treatment outcomes for the donors, $\boldsymbol{\omega}$ is the $(n_0 \times 1)$ vector of individual invariant weights, $\boldsymbol{\alpha}$ is $(n_0 \times K)$ matrix of coefficients. \square

B.b Proof of Lemma 1

Lemma 1 (Abadie's representation). *Assume that:*

- (i) *The true DGP is the interactive fixed effects model in (10)*
- (ii) *Assumption 2. Exact Fit holds*
- (iii) *The $F \times F$ matrix $\boldsymbol{\lambda}'_{pre}\boldsymbol{\lambda}_{pre}$ is invertible*

Then, we can write for $s > T_0$:

$$\begin{aligned}\hat{\tau}_s - \tau_s = & \frac{1}{n_1} \left[\boldsymbol{\lambda}_s (\boldsymbol{\lambda}'_{pre}\boldsymbol{\lambda}_{pre})^{-1} \boldsymbol{\lambda}'_{pre} \sum_{i=n_0+1}^N \left(\sum_{j=1}^{n_1} \hat{w}_{ij} \boldsymbol{\epsilon}_{j,pre} - \boldsymbol{\epsilon}_{i,pre} \right) \right] \\ & + \frac{1}{n_1} \left[\sum_{i=n_0+1}^N \left(\epsilon_{is} - \sum_{j=1}^{n_0} \hat{w}_{ij} \epsilon_{js} \right) \right]\end{aligned}$$

where \hat{w}_{ij} is the random coefficient weight of donor j on individual i , $\boldsymbol{\lambda}_s$ are the common factors at post-treatment time s , $\boldsymbol{\lambda}'_{pre}\boldsymbol{\lambda}_{pre}$ is the $F \times F$ matrix of interacted pre-treatment common factors, $\boldsymbol{\epsilon}_{j,pre}$ is a $(T_0 \times 1)$ vector of error terms for observation i in the pretreatment period and ϵ_{js} is the error term for individual j at time s .

Proof. The proof is a generalisation for $n_1 > 1$ of a result in [Abadie et al. \(2010, 503-504\)](#). We begin by writing for for all pre-treatment and post-treatment $t \in \{1, 2, \dots, T_0, \dots, T\}$:

$$\begin{aligned}\hat{\tau}_t - \tau_t = & \frac{1}{n_1} \sum_{i=n_0+1}^N (y_{it}(1) - \hat{y}_{it}(0)) - \tau \mathbb{1}\{t > T_0\} = \\ = & \frac{1}{n_1} \sum_{i=n_0+1}^{n_1} \left(y_{it}(0) + D_{it}\tau - \sum_{j=1}^{n_0} \hat{w}_{ij} y_{jt}(0) \right) - \tau \mathbb{1}\{t > T_0\} = \\ = & \frac{1}{n_1} \sum_{i=n_0+1}^{n_1} \left(y_{it}(0) - \sum_{j=1}^{n_0} \hat{w}_{ij} y_{jt}(0) \right)\end{aligned}$$

where D_{it} are all 0 in the pre-treatment period and all 1 in the post-treatment period. Then, we can drop the notation (0), as all potential outcomes in the last expression

are for the case of no treatment. We have for all $t \in \{1, 2, \dots, T_0, \dots, T\}$ that:

$$\begin{aligned} & \frac{1}{n_1} \sum_{i=n_0+1}^N (y_{it} - \sum_{j=1}^{n_0} \hat{w}_{ij} y_{jt}) = \\ & \frac{1}{n_1} \sum_{i=n_0+1}^N \left(\boldsymbol{\theta}_t \mathbf{x}_i + \boldsymbol{\lambda}_t \boldsymbol{\mu}_i + \epsilon_{it} - \left(\sum_{j=1}^{n_0} \hat{w}_{ij} (\boldsymbol{\theta}_t \mathbf{x}_j + \boldsymbol{\lambda}_t \boldsymbol{\mu}_j + \epsilon_{jt}) \right) \right) = \\ & = \frac{1}{n_1} \sum_{i=n_0+1}^N \left(\boldsymbol{\theta}_t \left(\mathbf{x}_i - \sum_{j=1}^{n_0} \hat{w}_{ij} \mathbf{x}_j \right) + \boldsymbol{\lambda}_t \left(\boldsymbol{\mu}_i - \sum_{j=1}^{n_0} \hat{w}_{ij} \boldsymbol{\mu}_j \right) + \left(\epsilon_{it} - \sum_{j=1}^{n_0} \hat{w}_{ij} \epsilon_{jt} \right) \right) \end{aligned}$$

where the first equality follows from the assumption for the data generating process.

We can then apply **Assumption 2** Exact Fit for the covariates \mathbf{x}_i in the last expression:

$$\frac{1}{n_1} \sum_{i=n_0+1}^N (y_{it} - \sum_{j=1}^{n_0} \hat{w}_{ij}) = \frac{1}{n_1} \sum_{i=n_0+1}^N \left(\boldsymbol{\lambda}_t \left(\boldsymbol{\mu}_i - \sum_{j=1}^{n_0} \hat{w}_{ij} \boldsymbol{\mu}_j \right) + \left(\epsilon_{it} - \sum_{j=1}^{n_0} \hat{w}_{ij} \epsilon_{jt} \right) \right) \quad (*)$$

Next, following [Abadie et al. \(2010\)](#), we will eliminate from $(*)$ the expression involving $\boldsymbol{\mu}$ -s. So, we stack over $t \in \{1, 2, \dots, T_0\}$ for the pretreatment period:

$$\begin{aligned} \sum_{i=n_0+1}^N \left(\mathbf{y}_{i,pre} - \sum_{j=1}^{n_0} \hat{w}_{ij} \mathbf{y}_{j,pre} \right) &= \sum_{i=n_0+1}^N \left[\boldsymbol{\lambda}_{pre} \left(\boldsymbol{\mu}_i - \sum_{j=1}^{n_1} \hat{w}_{ij} \boldsymbol{\mu}_j \right) + \left(\boldsymbol{\epsilon}_{i,pre} - \sum_{j=1}^{n_1} \hat{w}_{ij} \boldsymbol{\epsilon}_{j,pre} \right) \right] \\ \sum_{i=n_0+1}^N \left(\sum_{j=1}^{n_1} \hat{w}_{ij} \boldsymbol{\epsilon}_{j,pre} - \boldsymbol{\epsilon}_{i,pre} \right) &= \boldsymbol{\lambda}_{pre} \sum_{i=n_0+1}^N \left(\boldsymbol{\mu}_i - \sum_{j=1}^{n_1} \hat{w}_{ij} \boldsymbol{\mu}_j \right) \\ (\boldsymbol{\lambda}'_{pre} \boldsymbol{\lambda}_{pre})^{-1} \boldsymbol{\lambda}'_{pre} \sum_{i=n_0+1}^N \left(\sum_{j=1}^{n_1} \hat{w}_{ij} \boldsymbol{\epsilon}_{j,pre} - \boldsymbol{\epsilon}_{i,pre} \right) &= \sum_{i=n_0+1}^N \left(\boldsymbol{\mu}_i - \sum_{j=1}^{n_1} \hat{w}_{ij} \boldsymbol{\mu}_j \right) \end{aligned}$$

where the second row follows from the first due to **Assumption 2** Exact Fit for the outcomes in the pre-treatment period and the third row follows due to assumption of invertible $\boldsymbol{\lambda}'_{pre} \boldsymbol{\lambda}_{pre}$ and after multiplyg both sides by $(\boldsymbol{\lambda}'_{pre} \boldsymbol{\lambda}_{pre})^{-1} \boldsymbol{\lambda}'_{pre}$. Next, we substitute the last expression back into $(*)$ to obtain for $s > T_0$, i.e. in the post-

treatment period:

$$\begin{aligned}
& \frac{1}{n_1} \sum_{i=n_0+1}^N (y_{it} - \sum_{j=1}^{n_0} \hat{w}_{ij}) = \\
&= \frac{1}{n_1} \lambda_t \sum_{i=n_0+1}^N \left((\lambda'_{pre} \lambda_{pre})^{-1} \lambda'_{pre} \left(\sum_{j=1}^{n_1} \hat{w}_{ij} \epsilon_{j,pre} - \epsilon_{i,pre} \right) + \left(\epsilon_{it} - \sum_{j=1}^{n_0} \hat{w}_{ij} \epsilon_{jt} \right) \right) \\
&= \frac{1}{n_1} \left[\lambda_t (\lambda'_{pre} \lambda_{pre})^{-1} \lambda'_{pre} \sum_{i=n_0+1}^N \left(\sum_{j=1}^{n_1} \hat{w}_{ij} \epsilon_{j,pre} - \epsilon_{i,pre} \right) \right] + \frac{1}{n_1} \left[\sum_{i=n_0+1}^N \left(\epsilon_{it} - \sum_{j=1}^{n_0} \hat{w}_{ij} \epsilon_{jt} \right) \right]
\end{aligned}$$

as required. \square

B.c Proof of Lemma 2

Lemma 2. Suppose Z_1, Z_2, \dots, Z_n are all $\text{subG}(\sigma_i^2)$. Then, their average satisfies:

$$\frac{\sum_{i=1}^n Z_i}{n} \sim \text{SubG}\left(\frac{1}{n^2} \left(\sum_{i=1}^n \sigma_i\right)^2\right)$$

Proof. We will show this by induction. The expression holds trivially for $n = 1$, as if $Z_1 \sim \text{SubG}(\sigma^2)$, then $aZ_1 \sim \text{SubG}(a^2\sigma^2)$. Suppose for $k - 1$, it holds that:

$$\sum_{i=1}^{k-1} Z_i \sim \text{SubG}\left(\left(\sum_{i=1}^{k-1} \sigma_i\right)^2\right) \equiv \text{SubG}(\sigma^2)$$

Then, for k we have:

$$\begin{aligned} & E \left[e^{(h \sum_{i=1}^{k-1} Z_i + Z_k)} \right] \leq \\ & \leq \left[E \left(\exp \left(h \frac{\sum_i Z_i (\sigma + \sigma_{k+1})}{\sigma} \right) \right) \right]^{\frac{\sigma}{\sigma + \sigma_k}} \left[E \left(h \exp \left(\frac{Z_k (\sigma + \sigma_{k+1})}{\sigma_k} \right) \right) \right]^{\frac{\sigma_k}{\sigma + \sigma_k}} \leq \\ & \leq \left[\exp \left(\frac{h^2 (\sigma + \sigma_k)^2}{2\sigma^2} \sigma^2 \right) \right]^{\frac{\sigma}{\sigma + \sigma_k}} \left[\exp \left(\frac{h^2 (\sigma + \sigma_k)^2}{2\sigma_k^2} \sigma_k^2 \right) \right]^{\frac{\sigma_k}{\sigma + \sigma_k}} = \\ & = \left[\exp \left(\frac{h^2 \sigma (\sigma + \sigma_k)}{2} \right) \right] \left[\exp \left(\frac{h^2 \sigma_k (\sigma + \sigma_k)}{2} \right) \right] = \exp \left(\frac{(h(\sigma + \sigma_k))^2}{2} \right) \end{aligned}$$

where the second row follows due to Holder's inequality for $p = \frac{\sum_i Z_i (\sigma + \sigma_{k+1})}{\sigma}$ and $q = \frac{\sum_k Z_k (\sigma + \sigma_{k+1})}{\sigma}$. The third row follows due to the induction step for $k - 1$ and the definition of subG . The last expression gives the necessary result for the induction:

$$\sum_{i=1}^k Z_i \sim \text{SubG}\left(\left(\sum_{i=1}^k \sigma_i\right)^2\right)$$

Next, we multiply by $1/k$ to obtain:

$$\frac{\sum_{i=1}^k Z_i}{k} \sim \text{SubG}\left(\left(\frac{1}{k} \sum_{i=1}^k \sigma_i\right)^2\right)$$

□

B.d Proof of Proposition 2

Proposition 2 (CSC Bound Estimation Error). *Suppose that:*

- (i) *DGP is given by interactive fixed effects model in (10) with Assumption 1.*
- (ii) *Assumption 2 Exact Fit holds.*
- (iii) *Assumption 3 SubG Errors holds.*
- (iv) *$F \times F$ matrix $\lambda'_{pre} \lambda_{pre}$ is invertible*
- (v) *We work on an Euclidean space*

Then, with probability at least $1 - \frac{3}{\exp(0.25h^2)}$ in post-treatment period s the estimation error for the ATT estimated by Correlated Synthetic Control satisfies for $h > 0$:

$$|\hat{\tau}_s^{csc} - \tau_s| < \left(\frac{F\lambda_{max}^2}{\phi_{min}} \sqrt{\frac{2\sigma^2 n_0}{T_0} + \frac{h\sigma}{T_0^{1.5}}} \right) + \left(\frac{h\sigma F\lambda_{min}^2}{\phi_{max}^2} \right) + h\sigma\sqrt{2} \quad (12)$$

where λ_{min} and λ_{max} denote respectively the minimum and maximum common factor λ_{fs} in absolute value for either pre-treatment or post-treatment period. Similarly, ϕ_{min} and ϕ_{max} are respectively the minimum and maximum eigenvalue of matrix $\frac{1}{T_0} \lambda' \lambda$

Proof. The strategy for proof is related to the result in Abadie et al. (2010) and Theorem 2 in Ben-Michael et al. (2021). However, the final bound is different. We proceed in four steps where we prove the claims:

- We begin by proving that

Claim 1. *The estimation error for $\hat{\tau}_s$ for post-treatment period $s > T_0$ is:*

$$\hat{\tau}_s - \tau_s = \quad (31)$$

$$\frac{1}{n_1} \left[\lambda_t (\lambda'_{pre} \lambda_{pre})^{-1} \lambda'_{pre} \sum_{i=n_0+1}^N \left(\sum_{j=1}^{n_1} \hat{w}_{ij} \epsilon_{j,pre} \right) \right] - \quad (32)$$

$$\frac{1}{n_1} \left[\lambda_t (\lambda'_{pre} \lambda_{pre})^{-1} \lambda'_{pre} \sum_{i=n_0+1}^N \epsilon_{i,pre} \right] + \quad (33)$$

$$\frac{1}{n_1} \sum_{i=n_0+1}^N \left[\epsilon_{is} - \sum_{j=1}^{n_0} \hat{w}_{ij} \epsilon_{js} \right] \quad (34)$$

- We will bound each of (32), (33) and (34) in the last expression, similarly to Ben-Michael et al. (2021). Firstly, for 32 we find an upper bound:

Claim 2. *With probability $1 - \exp(-h^2/4)$, it holds that*

$$\frac{1}{n_1} \sum_{i=n_0+1}^{n_1} \lambda_s(\lambda'_{pre} \lambda_{pre})^{-1} \lambda'_{pre} \left[\sum_{j=1}^{n_0} \hat{w}_{ij} e_{j,pre} \right] \leq \frac{F \tilde{\lambda}^2}{\phi_{min} \left(\frac{1}{T_0} \lambda'_{pre} \lambda_{pre} \right)} \sqrt{\frac{2\sigma^2 n_0}{T_0} + \frac{h\sigma}{T_0^{1.5}}}$$

- Secondly, for (33) we can found a lower bound, as it is subtracted:

Claim 3. *With probability $1 - \exp(-h^2/2)$ we have:*

$$\frac{1}{n_1} \left[\lambda_t(\lambda'_{pre} \lambda_{pre})^{-1} \lambda'_{pre} \sum_{i=n_0+1}^N \epsilon_{i,pre} \right] > \frac{-hF^2 \lambda^2 \sigma^2}{\phi_{max} \left(\frac{1}{T_0} \lambda'_{pre} \lambda_{pre} \right)}$$

- Thirdly, (34) an upper bound can be found

Claim 4. *With prob. $1 - \exp(-h^2/2)$ we have $\frac{1}{n_1} \sum_{i=n_0+1}^N \left[\epsilon_{is} - \sum_{j=1}^{n_0} \hat{w}^j \epsilon_{js} \right] \leq h\sigma\sqrt{2}$*

- We finish the proof by combining the last three claims via Frechet's inequality and obtaining the statement in the initial proposition.

Claim 1 follows after a slight algebraic manipulation of Abadie's representation in Lemma 1.

Now, we will show Claim 2. Note that \hat{w}_{ij} are not independent from ϵ_{jt} for some

pretreatment $t < T_0$. Let us now proceed with the proof:

$$\begin{aligned}
& \frac{1}{n_1} \sum_{i=n_0+1}^{n_1} \lambda_s (\lambda'_{pre} \lambda_{pre})^{-1} \lambda'_{pre} \left[\sum_{j=1}^{n_0} \hat{w}_{ij} e_{j,pre} \right] \stackrel{\text{min eigenvalue}}{\leq} \\
& \frac{1}{n_1 T_0} \sum_{i=n_0+1}^N \left(\frac{\sum_{j=1}^{n_0} \hat{w}_{ij} (\lambda_s \lambda'_{pre} \epsilon_{j,pre})}{\phi_{min} \left(\frac{1}{T_0} \lambda'_{pre} \lambda_{pre} \right)} \right) \stackrel{\text{Cauchy-Schwarz Ineq.}}{\leq} \\
& \frac{1}{n_1 T_0 \phi_{min} \left(\frac{1}{T_0} \lambda'_{pre} \lambda_{pre} \right)} \sum_{i=n_0+1}^N \|\hat{w}_i\|_2 \|\lambda_s \lambda'_{pre} \epsilon_{n_0,pre}\|_2 \stackrel{\text{def. } \hat{w}}{\leq} \\
& \frac{1}{n_1 T_0 \phi_{min} \left(\frac{1}{T_0} \lambda'_{pre} \lambda_{pre} \right)} \sum_{i=n_0+1}^N \left\| \sum_{f=1}^F \sum_{t=1}^{T_0} \lambda_{sf} \lambda_{tf} \epsilon_{n_0,t} \right\|_2 \stackrel{\max_{t,f} |\lambda_{tf}|}{\leq} \\
& \frac{n_1}{n_1 T_0 \phi_{min} \left(\frac{1}{T_0} \lambda'_{pre} \lambda_{pre} \right)} \left\| F \tilde{\lambda}^2 \sum_{t=1}^{T_0} \epsilon_{n_0,t} \right\|_2 = \\
& \frac{F \tilde{\lambda}^2}{T_0 \phi_{min} \left(\frac{1}{T_0} \lambda'_{pre} \lambda_{pre} \right)} \left\| \sum_{t=1}^{T_0} \epsilon_{n_0,t} \right\|_2 \quad (*)
\end{aligned}$$

where the second inequality follows as we define $\phi_{min} \left(\frac{1}{T_0} \lambda'_{pre} \lambda_{pre} \right)$ as the smallest eigenvalues of the matrix $\frac{1}{T_0} \lambda'_{pre} \lambda_{pre}$. Now, note that each row $\sum_{t=1}^{T_0} \epsilon_{n_0,t}$ is a sum of T_0 iid variables, so that $\sum_{t=1}^{T_0} \epsilon_{j,t} \sim \text{SubG}(T_0^2 \sigma^2)$. However, then $\left\| \sum_{t=1}^{T_0} \epsilon_{n_0,t} \right\|_2$ is the Euclidean norm of n_0 SubG variables. Thus, we can use a corollary of Theorem 1 in [Hsu et al. \(2011\)](#) which is a version of the more general Hanson-Wright inequality ([Hanson and Wright, 1971](#)):

Corollary (to [Hsu et al. \(2011\)](#)'s Theorem 1). *Suppose \mathbf{A} is a $n_0 \times n_0$ matrix. If $z \in \mathbb{R}^{n_0}$ is $\text{subG}_{n_0}(\xi^2)$ vector, then:*

$$P(z' \mathbf{A} z < 2\xi^2 \text{trace}(\mathbf{A}) + h) \geq 1 - \exp \left(-\frac{h^2}{4\xi^2 \|\mathbf{A}\|_{op}} \right)$$

where $\|\mathbf{A}\|_{op}$ is the operator norm of matrix \mathbf{A} .

Let us apply this result with

$$z = \sum_{t=1}^{T_0} \epsilon_{n_0,t} \quad \xi^2 = T_0 \sigma^2 \quad \mathbf{A} = \mathbf{I}_{n_0}$$

where we have $\text{trace}(\mathbf{I}_{n_0}) = n_0$ and, since we are on Euclidian space, $\|\mathbf{I}_{n_0}\|_{op} = \phi_{\max}(\mathbf{I}'_{n_0} \mathbf{I}_{n_0}) = 1$, i.e. the operator norm is the largest eigenvalue. Thus, choosing $h = h\sqrt{T_0}\sigma$, we can find:

$$P\left(\left\|\sum_{t=1}^{T_0} \epsilon_{n_0,t}\right\|_2^2 < 2T_0\sigma^2 n_0 + h\sqrt{T_0}\sigma\right) \leq 1 - \exp\left(-\frac{h^2}{4}\right)$$

Although the last result is with respect to the *squared* Euclidean norm, we have $\left\|\sum_{t=1}^{T_0} \epsilon_{n_0,t}\right\|_2 \geq 0$ and so we can write:

$$P\left(\left\|\sum_{t=1}^{T_0} \epsilon_{n_0,t}\right\|_2 < \sqrt{2T_0\sigma^2 n_0 + h\sqrt{T_0}\sigma}\right) \leq 1 - \exp\left(-\frac{h^2}{4}\right)$$

Plugging in the last result to (*), we obtain with high probability that:

$$\begin{aligned} & \frac{F\tilde{\lambda}^2}{T_0\phi_{\min}\left(\frac{1}{T_0}\boldsymbol{\lambda}'_{pre}\boldsymbol{\lambda}_{pre}\right)} \left\|\sum_{t=1}^{T_0} \epsilon_{n_0,t}\right\|_2 \leq \\ & \frac{F\tilde{\lambda}^2}{T_0\phi_{\min}\left(\frac{1}{T_0}\boldsymbol{\lambda}'_{pre}\boldsymbol{\lambda}_{pre}\right)} \sqrt{2T_0\sigma^2 n_0 + h\sqrt{T_0}\sigma} = \\ & \frac{F\tilde{\lambda}^2}{\phi_{\min}\left(\frac{1}{T_0}\boldsymbol{\lambda}'_{pre}\boldsymbol{\lambda}_{pre}\right)} \sqrt{\frac{2\sigma^2 n_0}{T_0} + \frac{h\sigma}{T_0^{1.5}}} \end{aligned}$$

as required by Claim 2.

We will proceed to show Claim 3. Note that we are seeking a lower bound for expression (33). Analogically to the last derivation for an *upper* bound, here we can take the *largest* eigenvalue of $\left(\frac{1}{T_0}\boldsymbol{\lambda}'_{pre}\boldsymbol{\lambda}_{pre}\right)$ and the *smallest* $\tilde{\lambda}$:

$$\begin{aligned} & \frac{1}{n_1} \sum_{i=n_0+1}^N \lambda_s \left(\frac{\boldsymbol{\lambda}'_{pre}\boldsymbol{\lambda}_{pre}}{T_0}\right)^{-1} \boldsymbol{\lambda}'_{pre} \mathbf{e}_{i,pre} \geq \\ & \frac{1}{T_0 n_1} \sum_{i=n_0+1}^N \sum_{t=1}^{T_0} \frac{F\tilde{\lambda}^2 e_{it}}{\phi_{\max}\left(\frac{1}{T_0}\boldsymbol{\lambda}'_{pre}\boldsymbol{\lambda}_{pre}\right)} = \frac{F\tilde{\lambda}^2}{T_0 n_1 \phi_{\max}\left(\frac{1}{T_0}\boldsymbol{\lambda}'_{pre}\boldsymbol{\lambda}_{pre}\right)} \underbrace{\sum_{i=n_0+1}^N \sum_{t=1}^{T_0} e_{it}}_{n_1 T_0 \text{sub}G(\sigma^2)} \equiv Q \end{aligned}$$

where we define the maximum eigenvalue of $\left(\frac{1}{T_0}\boldsymbol{\lambda}'_{pre}\boldsymbol{\lambda}_{pre}\right)$ as $\phi_{\max}\left(\frac{1}{T_0}\boldsymbol{\lambda}'_{pre}\boldsymbol{\lambda}_{pre}\right)$. The

subG-ness of the sum of error follows from their independence. Note that we define Q as the last expression we found. Since Q is a sum of $n_1 T$ independent RV with $\sigma_Q^2 = \left(\frac{\sigma_F \lambda^2}{\phi_{max}^2}\right)^2$, we find a Chernoff bound for any h :

$$P(Q < -h) = P(-Q > h) = P(e^{-Qv} > e^{hv}) \leq \frac{E[e^{-Qv}]}{e^{hv}} \leq \frac{\exp\left(\frac{\sigma_Q^2 v^2}{2}\right)}{\exp(hv)} = \exp\left(\frac{\sigma_Q^2 v^2}{2} - hv\right)$$

where the first inequality follows from Markov's inequality and the second inequality is due to the definition of *SubG* random variables. Then, we find the tightest bound by maximising $\frac{\sigma_Q^2 v^2}{2} - hv$. FOC is $\sigma_Q^2 v = h$. So, the Chernoff bound is:

$$P(Q < -h) \leq \exp\left(-\frac{h^2}{2\sigma^2}\right)$$

Since the above holds for any h , take $h\sigma_Q$:

$$P\left(Q > -\frac{h\sigma_F \lambda^2}{\phi_{max}^2}\right) = 1 - \exp(-h^2/2)$$

as suggested by Claim 3, with probability at least $1 - \exp\left(-\frac{h^2}{2}\right)$ we have:

$$\frac{1}{n_1} \sum_{i=n_0+1}^N \lambda_s \left(\frac{\lambda'_{pre} \lambda_{pre}}{T_0}\right)^{-1} \lambda'_{pre} e_{i,pre} \geq Q > -\frac{h\sigma_F \lambda^2}{\phi_{max}^2}$$

Now, we will show the last Claim 4 for (34). For some given i in the treated group $\{n_0 + 1, \dots, n_0 + n_1\}$, we have that ϵ_{is} is independent from ϵ_{js} for the control group, $j \in \{1, 2, \dots, n_0\}$. So, we can write for some h :

$$\begin{aligned} E\left[e^{h(\epsilon_{is} - \sum_{j=1}^{n_0} \hat{w}_i^j \epsilon_{js})}\right] &\stackrel{ind.}{=} E[e^{h\epsilon_{is}}] \prod_{j=1}^{n_0} E\left[e^{-s\hat{w}_{ij}\epsilon_{js}}\right] \\ &\stackrel{subG}{\leq} e^{\frac{\sigma^2 h^2}{2}} \prod_{j=1}^{n_0} e^{\frac{\sigma^2 (-h\hat{w}_{ij})^2}{2}} = \exp\left(\frac{h^2 \sigma^2 (1 + \sum_j \hat{w}_{ij})}{2}\right) \end{aligned}$$

which means by definition of **subG** that:

$$\epsilon_{is} - \sum_{j=1}^{n_0} \hat{w}_i^j \epsilon_{js} \sim \text{SubG} \left(\sigma^2 \left(1 + \sum_{j=1}^{n_0} \hat{w}_{ij}^2 \right) \right)$$

However, if we further aggregate over i in the treated group, we have a problem. Define $z_i \equiv \epsilon_{is} - \sum_{j=1}^{n_0} \hat{w}_i^j \epsilon_{js}$. For some k and i , we do not have independence of z_i and z_k as for $s > T_0$:

$$\begin{aligned} E[z_i z_k] &= E \left[\left(\epsilon_{is} - \sum_{j=1}^{n_0} \hat{w}_{ij} \epsilon_{js} \right) \left(\epsilon_{ks} - \sum_{j=1}^{n_0} \hat{w}_{kj} \epsilon_{js} \right) \right] = \\ &= E[\epsilon_{is} \epsilon_{ks}] - E[\epsilon_{is} \sum_{j=1}^{n_0} \hat{w}_{kj} \epsilon_{js}] - E[\epsilon_{ks} \sum_{j=1}^{n_0} \hat{w}_{ij} \epsilon_{js}] + E \left[\left(\sum_{j=1}^{n_0} \hat{w}_{ij} \epsilon_{js} \right) \left(\sum_{j=1}^{n_0} \hat{w}_{kj} \epsilon_{js} \right) \right] = \\ &= 0 - 0 - 0 + \sum_{i=1}^{n_0} \hat{w}_{ij}^2 E[\epsilon_{is}^2] > 0 \end{aligned}$$

due to independence of ϵ_{is} , $E[\epsilon_{is}^2] \neq 0$ and the fact that every \hat{w}_{kj} is independent of all post-treatment errors ϵ_{is} .

We can still apply Lemma 2 for the sum of not necessarily independent **subG** variables, as $z_i = \epsilon_{is} - \sum_{j=1}^{n_0} \hat{w}_i^j \epsilon_{js}$ are **subG**. Thus, we obtain:

$$\frac{1}{n_1} \sum_{i=1}^{n_1} \epsilon_{is} - \sum_{j=1}^{n_0} \hat{w}_i^j \epsilon_{js} \sim \text{SubG} \left(\frac{\sigma^2}{n_1^2} \left(\sum_{i=n_0+1}^N \sqrt{1 + \sum_{j=1}^{n_0} \hat{w}_{ij}^2} \right)^2 \right)$$

As in the case of Claim 3, we can take a Chernoff bound and use the same argument.

So, with probability $1 - \exp(-h/2)$:

$$\frac{1}{n_1} \sum_{i=1}^{n_1} \epsilon_{is} - \sum_{j=1}^{n_0} \hat{w}_{ij} \epsilon_{js} \leq \frac{h\sigma}{n_1} \left(\sum_{i=n_0+1}^N \sqrt{1 + \sum_{j=1}^{n_0} \hat{w}_{ij}^2} \right) \leq h \frac{\sigma}{n_1} n_1 \sqrt{2}$$

and the last inequality follows due to the properties of the weights.

Now, let us call the events as follows:

$$\begin{aligned}
E_1 &= \left\{ \left(\frac{1}{n_1} \sum_{i=n_0+1}^{n_1} \lambda_s (\lambda'_{pre} \lambda_{pre})^{-1} \lambda'_{pre} \left(\sum_{j=1}^{n_0} \hat{w}_{ij} e_{j,pre} \right) \right) \leq \frac{F \tilde{\lambda}^2}{\phi_{min}} \sqrt{\frac{2\sigma^2 n_0}{T_0} + \frac{h\sigma}{T_0^{1.5}}} \right\} \\
E_2 &= \left\{ \frac{1}{n_1} \sum_{i=n_0+1}^N \lambda_s \left(\frac{\lambda'_{pre} \lambda_{pre}}{T_0} \right)^{-1} \lambda'_{pre} e_{i,pre} \geq -\frac{h\sigma F \tilde{\lambda}^2}{\phi_{max}^2} \right\} \\
E_3 &= \left\{ \frac{1}{n_1} \sum_{i=n_0+1}^N \left[\epsilon_{is} - \sum_{j=1}^{n_0} \hat{w}^j \epsilon_{js} \right] \leq h\sigma\sqrt{2} \right\}
\end{aligned}$$

where we use the shorthand to denote the minimum eigenvalues $\phi_{min} = \phi_{min} \left(\frac{1}{T_0} \lambda'_{pre} \lambda_{pre} \right)$ and maximum eigenvalue $\phi_{max} = \phi_{max} \left(\frac{1}{T_0} \lambda'_{pre} \lambda_{pre} \right)$

Then, we can use our results from the four claims and apply Frechet's inequality to show that:

$$\begin{aligned}
Pr \left[\hat{\tau}_s - \tau_s < \frac{F \tilde{\lambda}^2}{\phi_{min}} \sqrt{\frac{2\sigma^2 n_0}{T_0} + \frac{h\sigma}{T_0^{1.5}}} - \left(-\frac{h\sigma F \tilde{\lambda}^2}{\phi_{max}^2} \right) + h\sigma\sqrt{2} \right] \\
\geq Pr(E_1 \cap E_2 \cap E_3) \geq \\
\max \left(0, 1 - \frac{1}{\exp(0.25h^2)} - \frac{2}{\exp(0.5h^2)} \right) \geq \max \left(0, 1 - \frac{3}{\exp(0.25h^2)} \right)
\end{aligned}$$

where h will be set such that the last expression is bigger than 0. Lastly, we apply a union bound to the event that that E_1 , E_2 and E_3 hold simultaneously for the absolute value of the estimation error (and not just for $\hat{\tau}_s - \tau_s$ as done so far), then it follows:

$$|\hat{\tau}_s - \tau_s| < \frac{F \tilde{\lambda}^2}{\phi_{min}} \sqrt{\frac{2\sigma^2 n_0}{T_0} + \frac{h\sigma}{T_0^{1.5}}} + \left(\frac{h\sigma F \tilde{\lambda}^2}{\phi_{max}^2} \right) + h\sigma\sqrt{2}$$

with probability $1 - \frac{3}{\exp(0.25h^2)}$, due to Frechet inequality for logical disjunction, i.e. $P(A \cup B) \geq \max(P(A), P(B))$. Note that in the statement of the theorem we amend the notation, so that for clarity λ_{min} and λ_{max} denote respectively the minimum and maximum common factor λ_{fs} in absolute value. However, we have $\lambda_{min} \equiv \underline{\lambda}$ and $\lambda_{max} \equiv \tilde{\lambda}$ \square

B.e Proof of Proposition 3

Proposition 3 (DiD Estimation Error). *Suppose:*

1. *DGP is given by the interacted fixed effects model in (14) with Assumption 1.*
2. *$T = T_0 + 1$, i.e. one post-treatment period*
3. *We estimate τ^{did} via $y_{it} = \rho + \gamma_i + \delta_t + D_{it}\tau + u_{it}$*

Then, as $n_1 \rightarrow \infty$ the estimation error for DiD is:

$$|\hat{\tau}^{did} - \tau| \xrightarrow{p} \left| \frac{(\bar{\lambda}_{pre} - \lambda_T)(\bar{\mu}_{don} - E[\mu_i|D_{it} = 1])}{n_0} \right| \quad (13)$$

where $\bar{\lambda}_{pre}$ is $(1 \times F)$ vector of the average value of the common factors in the pre-treatment period, λ_T is $(1 \times F)$ vector of the common factors in the single post-treatment period, the $(F \times 1)$ vector $\bar{\mu}_{don}$ is the average of the factor loadings for the donors and $E[\mu_i|D_{it} = 1]$ is the $(F \times 1)$ vector with expected value of the factor loadings in the treatment group, .

Proof. We follow an analogical approach to Proposition 4. So, we begin by estimating $\hat{\tau}$ via difference-in-difference and so fit the model:

$$y_{it} = \rho + \gamma_i + \delta_t + D_{it}\tau + e_{it}$$

where ρ is an intercept, γ_i is an individual fixed effect and δ_t is a time fixed effect. In matrix form, let us denote by \mathbf{Z} the $NT \times (1 + N - 1 + T - 1)$ matrix of dummies, used for estimating $\rho + \gamma_i + \delta_t$, so that

$$\mathbf{y} = \mathbf{Z} \begin{pmatrix} \rho \\ \gamma \\ \delta \end{pmatrix} + \mathbf{D}\tau + \mathbf{e}$$

where \mathbf{D} is a $(NT \times 1)$ vector of treatment indicators. By Frisch-Waugh-Lovell's

Theorem, we have that:

$$\begin{aligned}
\hat{\tau} &= (\hat{\mathbf{u}}'_{D \cdot \rho, \gamma, \delta} \hat{\mathbf{u}}_{D \cdot \rho, \gamma, \delta})^{-1} \hat{\mathbf{u}}'_{D \cdot \rho, \gamma, \delta} \mathbf{u}_{y \cdot \rho, \gamma, \delta} = \\
&= (\mathbf{D}'(\mathbf{I}_{NT} - \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}')\mathbf{D})^{-1} \mathbf{D}'(\mathbf{I}_{NT} - \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}')\mathbf{y} = \\
&= (\hat{\mathbf{u}}'_{D \cdot \rho, \gamma, \delta} \mathbf{D})^{-1} \hat{\mathbf{u}}'_{D \cdot \rho, \gamma, \delta} \mathbf{y}
\end{aligned}$$

where $\hat{\mathbf{u}}'_{D \cdot \rho, \gamma, \delta}$ are the residuals⁴⁹ from regressing D_{it} on ρ , γ_i and δ_t and analogically for $\hat{\mathbf{u}}'_{y \cdot \rho, \gamma, \delta}$ being the residuals from regressing y_{it} on ρ , γ_i and δ_t . The second row follows due to idempotence of matrix $\mathbf{M} \equiv (\mathbf{I}_{NT} - \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}')$, i.e. $\mathbf{M}'\mathbf{M} = \mathbf{M}$. Next, we substitute the factor model generating \mathbf{y} to obtain:

$$\begin{aligned}
\hat{\tau} &= (\mathbf{D}'\mathbf{M}\mathbf{D})^{-1} \mathbf{D}'\mathbf{M}\mathbf{y} = \\
&= (\mathbf{D}'\mathbf{M}\mathbf{D})^{-1} \mathbf{D}'\mathbf{M}(\text{Vec}(\boldsymbol{\theta}\mathbf{X}') + \text{Vec}(\boldsymbol{\lambda}\boldsymbol{\mu}') + \mathbf{D}\tau + \mathbf{e}) = \\
&= \tau + \left(\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT}\right)^{-1} \frac{\mathbf{D}'\mathbf{M}\text{Vec}(\boldsymbol{\theta}\mathbf{X}')}{NT} + \left(\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT}\right)^{-1} \frac{\mathbf{D}'\mathbf{M}\text{Vec}(\boldsymbol{\lambda}\boldsymbol{\mu}')}{NT} + \left(\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT}\right)^{-1} \frac{\mathbf{D}'\mathbf{M}\mathbf{e}}{NT}
\end{aligned}$$

where we need to use the $\text{Vec}(\mathbf{A})$ operator which stacks all columns of matrix \mathbf{A} . Given the assumption on the DGP, let us hold T and n_0 fixed and let $n_1 \rightarrow \infty$, so that $N \rightarrow \infty$. Then, we have for all $(i, j) \in \{1, 2, \dots, N\}$ and $(j, s) \in \{1, 2, \dots, T\}$ by the Law of Large Numbers:

$$\frac{\mathbf{D}'\mathbf{M}\mathbf{e}}{NT} \xrightarrow{p} E[g(D_{it})e_{js}] = 0 \quad \frac{\mathbf{D}'\mathbf{M}\text{Vec}(\boldsymbol{\theta}\mathbf{X}')}{NT} \xrightarrow{p} E[f(D_{it})x_j] = 0$$

where $g(\cdot)$ and $f(\cdot)$ are some measurable functions. Crucially, we note that while the elements of \mathbf{M} could be correlated with the elements of \mathbf{D} , what matters here is the fact that both \mathbf{x} and $\boldsymbol{\epsilon}$ are assumed to be independent from \mathbf{M} and \mathbf{D} . Lastly, to get the probability limits we need either $n_1 \rightarrow \infty$ or $n_0 \rightarrow \infty$ or both. Note that in the expression above $\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT}$ is a scalar, as \mathbf{D} is $(NT \times 1)$ vector. Assuming (as we

⁴⁹We do not need to worry about predicted values outside of $[0, 1]$ range, as we are regressing only on fixed effects and not on any continuous variable. I would like to thank Frank DiTraglia for pointing this out.

shall show below) that $\left(\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT}\right)^{-1} \xrightarrow{p} c \neq 0$, then we get as $n_1 \rightarrow \infty$:

$$\hat{\tau} - \tau \xrightarrow{p} \left(\text{plim}_{N \rightarrow \infty} \left(\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT} \right) \right)^{-1} \text{plim}_{N \rightarrow \infty} \left(\frac{\mathbf{D}'\mathbf{M}Vec(\boldsymbol{\lambda}\boldsymbol{\mu}')}{NT} \right) \quad (35)$$

with the result following from the Law of Large Numbers and Slutsky's Theorem. Next, we will obtain a more precise estimate of the probability limit. We note that \mathbf{M} is a function of the matrix of dummies \mathbf{Z} and so only of n_1 , n_0 , T_0 and T . To see this, consider the case for $N = 3$ and $T = 2$ where \mathbf{Z} is 6×4 matrix:

$$\begin{pmatrix} 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 1 & 0 & 0 \\ 1 & 1 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 1 \end{pmatrix}$$

More generally, the first column contains 1-s for the intercept, the next $N - 1$ columns contain dummies for all individuals in the data and the last $T - 1$ are with dummies for the treated periods.

With this insight in mind, we shall find an *exact* expression for $\mathbf{M} = (\mathbf{I}_{NT} - \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}')$.

We begin by showing:

$$\mathbf{Z}'\mathbf{Z} = \begin{pmatrix} NT & \boldsymbol{\iota}'_{N-1}T & \boldsymbol{\iota}'_{T-1}N \\ \boldsymbol{\iota}_{N-1}T & T\mathbf{I}_{N-1} & \mathbf{1}_{(N-1) \times (T-1)} \\ \boldsymbol{\iota}_{T-1}N & \mathbf{1}_{(T-1) \times (N-1)} & N\mathbf{I}_{T-1} \end{pmatrix} \quad (36)$$

where $\boldsymbol{\iota}_K$ is $1 \times K$ column vector of ones and $\mathbf{1}_{K \times Q}$ is a $(K \times Q)$ matrix of ones. This matrix be found by guess and verify. Next, we define by \mathbf{H} the matrix from (36):

$$\mathbf{H} = \begin{pmatrix} T\mathbf{I}_{N-1} & \mathbf{1}_{(N-1) \times (T-1)} \\ \mathbf{1}_{(T-1) \times (N-1)} & N\mathbf{I}_{T-1} \end{pmatrix}$$

We can then apply the matrix inversion lemma ([Bernstein, 2009](#), p.108):

$$\mathbf{B}^{-1} = \begin{pmatrix} \mathbf{P} & \mathbf{Q} \\ \mathbf{R} & \mathbf{S} \end{pmatrix}^{-1} = \begin{pmatrix} \mathbf{P}^{-1} + \mathbf{P}^{-1}\mathbf{Q}(\mathbf{S} - \mathbf{R}\mathbf{P}^{-1}\mathbf{Q})^{-1}\mathbf{R}\mathbf{P}^{-1} & -\mathbf{P}^{-1}\mathbf{Q}(\mathbf{S} - \mathbf{R}\mathbf{P}^{-1}\mathbf{Q})^{-1} \\ -(\mathbf{S} - \mathbf{R}\mathbf{P}^{-1}\mathbf{Q})^{-1}\mathbf{R}\mathbf{P}^{-1} & (\mathbf{S} - \mathbf{R}\mathbf{P}^{-1}\mathbf{Q})^{-1} \end{pmatrix} \quad (37)$$

where all the inverted matrices are assumed to be non-singular and also use the Sherman-Morrison-Woodbury lemma ([Bernstein, 2009](#), p.141):

$$(\mathbf{C} + \mathbf{a}\mathbf{b}')^{-1} = \mathbf{C}^{-1} - \frac{\mathbf{C}^{-1}\mathbf{a}\mathbf{b}'\mathbf{C}^{-1}}{1 + \mathbf{b}'\mathbf{C}^{-1}\mathbf{a}} \quad (38)$$

where \mathbf{a} and \mathbf{b} are column vectors. If we use these two results, we can find:

$$\mathbf{H}^{-1} = \begin{pmatrix} \frac{1}{T} \left(\mathbf{I}_{n-1} + \frac{T-1}{T+N-1} \mathbf{1}_{(N-1) \times (N-1)} \right) & -\frac{1}{T+N-1} \mathbf{1}_{(N-1) \times (T-1)} \\ -\frac{1}{T+N-1} \mathbf{1}_{(T-1) \times (N-1)} & \frac{1}{N} \left(\mathbf{I}_{T-1} + \frac{N-1}{T+N-1} \mathbf{1}_{(T-1) \times (T-1)} \right) \end{pmatrix}$$

The calculations are tedious, but not so challenging mentally. Next, we can use \mathbf{H}^{-1} in finding the inverse of $\mathbf{Z}'\mathbf{Z}$. We achieve this in three steps. Firstly, use the matrix inversion lemma (37) to the expression for $\mathbf{Z}'\mathbf{Z}$ in (36) where we use \mathbf{H} as our bottom-right block \mathbf{S} in (37). Secondly, we use \mathbf{H}^{-1} and the Sherman-Morrison-Woodbury lemma to find the bottom right panel of the inverted lemma, that is $(\mathbf{S} - \mathbf{R}\mathbf{P}^{-1}\mathbf{Q})^{-1}$. Thirdly, we plug-in this expression in the other three panels, noting that it appears in all of them. The derivations are significantly simplified by the fact that in $\mathbf{Z}'\mathbf{Z}$ the top-left panel is a scalar and so the top-right and bottom-left blocks are actually vectors.

Using this procedure, we can find:

$$(\mathbf{Z}'\mathbf{Z}) = \begin{pmatrix} \frac{N+T-1}{NT} & -\frac{1}{T}\boldsymbol{\iota}'_{N-1} & -\frac{1}{N}\boldsymbol{\iota}'_{T-1} \\ -\frac{1}{T}\boldsymbol{\iota}_{N-1} & \frac{1}{T} (\mathbf{I}_{N-1} + \mathbf{1}_{(N-1) \times (N-1)}) & \mathbf{O}_{(N-1) \times (T-1)} \\ -\frac{1}{N}\boldsymbol{\iota}_{T-1} & \mathbf{O}_{(T-1) \times (N-1)} & \frac{1}{N} (\mathbf{I}_{T-1} + \mathbf{1}_{(T-1) \times (T-1)}) \end{pmatrix}$$

The next step in the derivation is to guess and verify the $(NT \times NT)$ matrix $\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'$

has a block symmetric structure:

$$\begin{pmatrix} \mathbf{A} & \mathbf{B} & \mathbf{B} & \dots & \mathbf{B} & \mathbf{B} \\ \mathbf{B} & \mathbf{A} & \mathbf{B} & \dots & \mathbf{B} & \mathbf{B} \\ \mathbf{B} & \mathbf{B} & \mathbf{A} & \dots & \mathbf{B} & \mathbf{B} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ \mathbf{B} & \mathbf{B} & \mathbf{B} & \dots & \mathbf{A} & \mathbf{B} \\ \mathbf{B} & \mathbf{B} & \mathbf{B} & \dots & \mathbf{B} & \mathbf{A} \end{pmatrix}$$

where $\mathbf{A} = \frac{1}{NT} (T\mathbf{I}_T + (N-1)\mathbf{1}_T)$ and $\mathbf{B} = \frac{1}{NT} (T\mathbf{I}_T - \mathbf{1}_{T \times T})$. On each block row, we have $N-1$ matrices \mathbf{B} and one matrix \mathbf{A} on the block diagonal for some treated observation. In specific, it is given by:

$$\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}' = \frac{1}{NT} \times \overbrace{\left(\begin{array}{ccccc} T\mathbf{I}_T + (N-1)\mathbf{1}_T & \overbrace{T\mathbf{I}_T - \mathbf{1}_{T \times T}}^{T(N-1)} & T\mathbf{I}_T - \mathbf{1}_{T \times T} & \dots & T\mathbf{I}_T - \mathbf{1}_{T \times T} \\ T\mathbf{I}_T - \mathbf{1}_{T \times T} & T\mathbf{I}_T + (N-1)\mathbf{1}_T & \overbrace{T\mathbf{I}_T - \mathbf{1}_{T \times T}}^{T(N-2)} & \dots & T\mathbf{I}_T - \mathbf{1}_{T \times T} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ T\mathbf{I}_T - \mathbf{1}_{T \times T} & \dots & \dots & \dots & T\mathbf{I}_T + (N-1)\mathbf{1}_T \end{array} \right)}^{NT}$$

It is then easy to calculate $\mathbf{M} = (\mathbf{I}_{NT} - \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}')$:

$$\frac{1}{NT} \begin{pmatrix} T(N-1)\mathbf{I}_T - (N-1)\mathbf{1}_{T \times T} & -T\mathbf{I}_T + \mathbf{1}_{T \times T} & \dots & -T\mathbf{I}_T + \mathbf{1}_{T \times T} \\ -T\mathbf{I}_T + \mathbf{1}_{T \times T} & T(N-1)\mathbf{I}_T - (N-1)\mathbf{1}_{T \times T} & \dots & -T\mathbf{I}_T + \mathbf{1}_{T \times T} \\ \vdots & \vdots & \ddots & \vdots \\ -T\mathbf{I}_T + \mathbf{1}_{T \times T} & -T\mathbf{I}_T + \mathbf{1}_{T \times T} & \dots & T(N-1)\mathbf{I}_T - (N-1)\mathbf{1}_{T \times T} \end{pmatrix} \quad (39)$$

Given this expression, we will now find the probability limit of:

$$\left(\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT} \right)^{-1} \left(\frac{\mathbf{D}'\mathbf{M}\mathbf{Vec}(\lambda\boldsymbol{\mu}')}{NT} \right)$$

Let us first find an expression for the scalar $\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT}$. At this point, we realise that the vector \mathbf{D} contains only zeros and ones. Since $T_0 = T - 1$, we have:

$$\mathbf{D} = (\underbrace{0, \dots, 0}_{T}, \underbrace{0, \dots, 0}_{T}, \dots, \underbrace{0, \dots, 0}_{T}, \underbrace{0, \dots, 0}_{T_0}, \underbrace{1, 0, \dots, 0}_{T_0}, \underbrace{1, 0, \dots, 0}_{T_0}, \dots, \underbrace{0, \dots, 0}_{T_0}, \underbrace{1, \dots, 1}_{T_0})'$$

So, using the structure of \mathbf{D} , we can find:

$$\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT} = \sum_{p=1}^{NT} \sum_{r=1}^{NT} \mathbb{1}\{(D_p = 1) \cap (D_r = 1)\} m_{pr} = n_1 m_{kk} + n_1(n_1 - 1) m_{qk}$$

where m_{kk} is any diagonal entry of \mathbf{M} and m_{qk} is any diagonal entry of the off-diagonal matrix $(-T\mathbf{I}_T + \mathbf{1}_{T \times T})$ in the expression for \mathbf{M} in (39). Next, the last expression can be written as $n_1 \rightarrow \infty$ in the following way:

$$\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT} = n_1 \frac{[T(N-1) - (N-1)]}{NT} + n_1(n_1 - 1) \frac{(-T+1)}{NT} \quad (40)$$

$$\begin{aligned} &= \frac{(T-1)n_0 n_1}{(n_1 + n_0)T} = \frac{(Tn_0 - n_0)}{(1 + \frac{n_0}{n_1})T} \\ &\xrightarrow{p} \frac{(T-1)n_0}{T} \end{aligned} \quad (41)$$

Note that in addition if we let $n_0 \rightarrow \infty$ we have $\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT} \xrightarrow{p} \frac{T-1}{T} n_1$ and if we have both $n_0 \rightarrow \infty$ and $n_1 \rightarrow \infty$ then $\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT} \xrightarrow{p} \frac{T-1}{T} \frac{1}{0} \approx \infty$.⁵⁰ By Slutsky's Theorem, we have $\left(\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT}\right) \xrightarrow{p} \frac{T}{T_0 n_0}$.

Next, assuming that $n_1 > n_0$, we can write:

$$\begin{aligned} \frac{1}{NT} \mathbf{D}'\mathbf{M}\mathbf{V}ec(\boldsymbol{\lambda}\boldsymbol{\mu}') &= \frac{1}{NT} \sum_{t=1}^{T_0} \left[\left(\frac{n_1}{n_0} \sum_{j=1}^{n_0} \boldsymbol{\mu}_j \right) - \left(\sum_{i=n_0+1}^N \boldsymbol{\mu}_i \right) \right] \boldsymbol{\lambda}_t \\ &\quad - \frac{T_0}{NT} \left[\left(\frac{n_1}{n_0} \sum_{j=1}^{n_0} \boldsymbol{\mu}_j \right) - \left(\sum_{i=n_0+1}^N \boldsymbol{\mu}_i \right) \right] \boldsymbol{\lambda}_T \end{aligned}$$

given our expression for \mathbf{M} and recalling that \mathbf{D} contains of ones and zeros. Then, by

⁵⁰Interestingly, see what happens when $n_0 \rightarrow \infty$, $n_1 \rightarrow \infty$ and $\frac{n_0}{n_1} \rightarrow c$

the Law of Large Numbers and our assumptions, we can find as $n_1 \rightarrow \infty$:

$$\begin{aligned}
& \frac{1}{NT} \sum_{t=1}^{T_0} \left[\left(\frac{n_1}{n_0} \sum_{j=1}^{n_0} \mu_j \right) \lambda_t \right] \xrightarrow{p} \bar{\mu}_{don} \bar{\lambda}_{pre} \frac{T_0}{T} \\
& \frac{1}{NT} \sum_{t=1}^{T_0} \left[\left(\sum_{i=n_0+1}^N \mu_i \right) \lambda_t \right] \xrightarrow{p} E[\mu_{tr} | D_{it} = 1] \bar{\lambda}_{pre} \frac{T_0}{T} \\
& \frac{T_0}{NT} \left[\left(\frac{n_1}{n_0} \sum_{j=1}^{n_0} \mu_j \right) \lambda_T \right] \xrightarrow{p} \bar{\mu}_{don} \lambda_T \frac{T_0}{T} \\
& \frac{T_0}{NT} \left[\left(\sum_{i=n_0+1}^N \mu_i \right) \lambda_T \right] \xrightarrow{p} E[\mu_{tr} | D_{it} = 1] \lambda_T \frac{T_0}{T}
\end{aligned}$$

We can then combine by Slutsky's Theorem to obtain:

$$\mathbf{D}' MVec(\lambda \mu') \xrightarrow{p} \frac{T_0}{T} (\hat{\mu}_{don} - E[\hat{\mu} | D_{it} = 1]) (\bar{\lambda}_{pre} - \lambda_T) \quad (42)$$

Lastly, via Slutsky's Theorem we can combine with (41) to obtain the final result in the proposition, using the representation (35):

$$\hat{\tau} - \tau \xrightarrow{p} \frac{(\bar{\mu}_{don} - E[\mu | D_{it} = 1]) (\bar{\lambda}_{pre} - \lambda_T)}{n_0}$$

□

B.f Proof of Proposition 4

Proposition 4 (Consistency of iDiD). *Suppose that:*

(i) *Data is generated by $y_{it} = \beta \mathbf{x}'_i + \lambda_t \boldsymbol{\mu}_i + D_{it}\tau + \epsilon_{it}$ under Assumption 1.*

(ii) *$\hat{\tau}^{idid}$ is given by Definition 1. Infeasible DiD*

Then, Difference-in-Difference will estimate ATT consistently: $\hat{\tau}^{idid} \xrightarrow{p} \tau$

Proof. We follow an analogical approach to Proposition 3. So, we begin by estimating $\hat{\tau}$ via infeasible difference-in-difference and so fit the model:

$$\tilde{y}_{it} \equiv y_{it} - \tilde{\boldsymbol{\mu}}_i \tilde{\boldsymbol{\lambda}}_t = \rho + \gamma_i + \delta_t + D_{it}\tau + e_{it}$$

where ρ is an intercept, γ_i is an individual fixed effect and δ_t is a time fixed effect and where we define as \tilde{y}_{it} the values of the outcomes after subtracting the demeaned interactive fixed effects. In matrix form, let us denote by \mathbf{Z} the $NT \times (1 + N - 1 + T - 1)$ matrix of dummies, used for estimating $\rho + \gamma_i + \delta_t$, so that

$$\tilde{\mathbf{y}} = \mathbf{Z} \begin{pmatrix} \rho \\ \boldsymbol{\gamma} \\ \boldsymbol{\delta} \end{pmatrix} + \mathbf{D}\tau + \mathbf{e}$$

By Frisch-Waugh-Lovell's Theorem, we have that:

$$\begin{aligned} \hat{\tau} &= (\hat{\mathbf{u}}'_{D \cdot \rho, \gamma, \delta} \hat{\mathbf{u}}_{D \cdot \rho, \gamma, \delta})^{-1} \hat{\mathbf{u}}'_{D \cdot \rho, \gamma, \delta} \mathbf{u}_{\tilde{y} \cdot \rho, \gamma, \delta} = \\ &= (\mathbf{D}'(\mathbf{I}_{NT} - \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}')\mathbf{D})^{-1} \mathbf{D}'(\mathbf{I}_{NT} - \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}')\tilde{\mathbf{y}} = \\ &= (\hat{\mathbf{u}}'_{D \cdot \rho, \gamma, \delta} \mathbf{D})^{-1} \hat{\mathbf{u}}'_{D \cdot \rho, \gamma, \delta} \tilde{\mathbf{y}} \end{aligned}$$

where $\hat{\mathbf{u}}'_{D \cdot \rho, \gamma, \delta}$ are the residuals⁵¹ from regressing D_{it} on ρ , γ_i and δ_t and analogically for $\hat{\mathbf{u}}'_{\tilde{y} \cdot \rho, \gamma, \delta}$ being the residuals from regressing \tilde{y}_{it} on ρ , γ_i and δ_t . The second row follows due to idempotence of matrix $\mathbf{M} \equiv (\mathbf{I}_{NT} - \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}')$, i.e. $\mathbf{M}'\mathbf{M} = \mathbf{M}$.

⁵¹We do not need to worry about predicted values outside of $[0, 1]$ range, as we are regressing only on fixed effects and not on any continuous variable. I would like to thank Frank DiTraglia for pointing this out.

Next, we substitute the interactive fixed effects model generating $\tilde{\mathbf{y}}$ to obtain:

$$\begin{aligned}\hat{\tau} &= (\mathbf{D}'\mathbf{M}\mathbf{D})^{-1} \mathbf{D}'\mathbf{M}\tilde{\mathbf{y}} = \\ &= (\mathbf{D}'\mathbf{M}\mathbf{D})^{-1} \mathbf{D}'\mathbf{M}(\text{Vec}(\beta\mathbf{X}') + \mathbf{D}\tau + \boldsymbol{\epsilon}) = \\ &= \tau + \left(\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT}\right)^{-1} \frac{\mathbf{D}'\mathbf{M}\text{Vec}(\beta\mathbf{X}')}{NT} + \left(\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT}\right)^{-1} \frac{\mathbf{D}'\mathbf{M}\boldsymbol{\epsilon}}{NT}\end{aligned}$$

where we need to use the $\text{Vec}(\mathbf{A})$ operator which stacks all columns of matrix \mathbf{A} (Crépon and Mairesse, 2008). Given the assumption on the DGP, let us hold T and let $n_0 \rightarrow \infty$ or $n_1 \rightarrow \infty$, so that $N \rightarrow \infty$. Then, we have for all $(i, j) \in \{1, 2, \dots, N\}$ and $(j, s) \in \{1, 2, \dots, T\}$ by the Law of Large Numbers:

$$\frac{\mathbf{D}'\mathbf{M}\boldsymbol{\epsilon}}{NT} \xrightarrow{p} E[g(D_{it})e_{js}] = 0 \quad \frac{\mathbf{D}'\mathbf{M}\text{Vec}(\boldsymbol{\theta}\mathbf{X}')}{NT} \xrightarrow{p} E[f(D_{it})x_j] = 0$$

where $g(\cdot)$ and $f(\cdot)$ are some measurable functions. Note that in the expression above $\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT}$ is a scalar, as \mathbf{D} is $(NT \times 1)$ vector. Assuming (as we shall show below in Proposition 3) that $\left(\frac{\mathbf{D}'\mathbf{M}\mathbf{D}}{NT}\right)^{-1} \xrightarrow{p} c \neq 0$, then we get:

$$N \rightarrow \infty : \quad \hat{\tau}^{idid} \xrightarrow{p} \tau$$

with the result following from the Law of Large Numbers and Slutsky's Theorem. \square